

PMDS603P Deep Learning Lab Assessment 2

August 2025

Note: Make a single PDF file of the work you are doing in a Jupyter notebook. Upload with the proper format. Please mention your name and roll no properly with the Experiment number on the first page of your submission.

Question1. Today, we will try to recall the work done in the previous lab first. The second problem attempted in the last lab was to use MNIST dataset which contains handwritten numbers (their images) from 0 to 9 digits. First try to fit a simple neural network model. Let us import the necessary modules required for this along with the dataset.

It contains 70000 handwritten images of digits from 0 to 9. So its a 10 class classification problem. Lets try to create a model that can do the classification task.

```
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import SGD
import matplotlib.pyplot as plt
batch_size = 128
num_classes = 10
epochs = 10
(x_train, y_train), (x_test, y_test) = mnist.load_data()
plt.imshow(x_train[3])
print(x_train[3])
```

Next let us reshape the whole images in to a single vector and normalize them.

```
x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
```

- Do the one-hot encoding for the labels with appropriate functions for both training and testing data.
- Now here I have given the number of iterations as 25 but it can be changed, and the results can be observed.
- Now create a simple neural network model using a sequential class with the following

architecture.

Two hidden layers with 512 neurons in each layer with a sigmoid activation function. An output layer with 10 neurons with a softmax activation function. (Since it's a multiclass classification problem)

- Compile and fit the model with an appropriate loss function and use the SGD optimizer.

You can change the learning rate in your optimizer SGD by defining

```
sgd1 = SGD(learning_rate = 0.01)
```

and then use sgd1 as your optimizer while compiling your model.

- You can print your model validation error and accuracy using You will get your validation

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

set loss and accuracy.

- Now if you want to check one case, how your model is predicting you can try any one image print it and use your model to predict the same and print. The argmax function gives us the index of the class with highest probability.

1 Regularization Techniques

Next, we will see how we can use different regularization techniques to get better accuracy for the models.

- **Dropout Technique.**

Now Already you have already created a model in the previous section. Now use this code to import Dropout layers.

```
from keras.layers import Dropout
```

Further, try to include a Dropout layer after your first and second hidden layers by using the code

```
model.add(Dropout(0.2))
```

Thus, in the layer before that, 20% of neurons become inactive during training in an epoch.

But it's important to understand that the dropout parameter value that you see as chosen as 0.2 is a hyperparameter that can be tuned and found out. An appropriate parameter value will give you better results. Try manually whether you can come up with an appropriate dropout parameter value for your model by giving different values for the parameter.

- **Early Stopping:** Next, we can see how the early stopping technique can be implemented to reduce overfitting of the model and help it to generalize well.

Now, for that, let's first divide the training data further into a train and a validation set. Do the one-hot encoding part for all three sets of data. Training, validation, and

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
X_subtrain, X_valid, y_subtrain, y_valid = train_test_split(x_train, y_train, test_size=0.10, random_state=1)
```

testing set.

Now create your model with appropriate activation functions and loss functions, etc, and print the model summary of your model.

Further import the class early stopping from keras using

from keras.callbacks import EarlyStopping

Define an object of the early stopping class

```
estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-5, mode = 'min', patience=4,
                      verbose=1, restore_best_weights=True)
```

Note:

monitor ='val_loss'

patience=4

min_delta: Minimum change in the monitored quantity to qualify as an improvement, i.e. an absolute change of less than min_delta, will count as no improvement.

mode: In min mode, training will stop when the quantity monitored has stopped decreasing;

in "max" mode it will stop when the quantity monitored has stopped increasing;

restore_best_weights=True

Further compile the model, and while fitting the model, use

```
model.fit(X_subtrain, y_subtrain,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(X_valid, y_valid), callbacks=[estop])
```

Here we have used the validation set for the validation part and then defined callbacks, as you see above, which will help perform early stopping.

Now you can evaluate your accuracy of the model and verify.

Challenging Question: Try for a scratch code for this case where you can create a custom neural network without using any inbuilt classes like sequential etc. Where you need to define a class neural network which has methods like forwardpass, backwardpass, and train. Figure out how we can do this.

This model has inputs as [0, 0, 1], [0, 1, 1], [1, 0, 1], [1, 1, 1] and the expected output as [0], [1], [1], [0] in each case. So there are three features in our dataset as you see above. The activation function is to be taken as sigmoid. The architecture is like we have only one hidden layer and an output layer with one neuron. Take the error function as $(1/2)(y - \hat{y})^2$.

Question 2. **cifar10** dataset is also an inbuilt dataset which contains 10 classes of images, mainly, 0-airplane, 1-automobile, 2-bird, 3-cat, 4-deer, 5-dog, 6-frog, 7-horse, 8-ship, 9-truck.

Load the inbuilt dataset cifar10 as you did in last lab by replacing `mnist.load_datae()` as `cifar10.load_data()`. First, try to import it from `keras.datasets` as you did for `mnist`.

Now, identify the size of the images you have first of all. You can now see $32*32*3$ images that is $32*32$ pixel images with 3 channels that give the RGB values since we have a color image.

Try to print the shape of each image and see. you will see it's stored like $32*32*3$ arrays. Now, try to visualize certain images using appropriate functions.

Check the size of `x_train` and `x_test` and reshape them into one-dimensional arrays as done in the case of `mnist` dataset.

Do necessary pre-processing and split the data into training, validation, and testing sets.

Create a new model using a sequential class with appropriate hidden layers and output layer neurons. Choose appropriate activation functions like sigmoid and relu, etc. And also an appropriate one in the output layer.

Choose the error function appropriately. Include early stopping technique in your model and run the model for 500 epochs. Try to come up with a better model with decent accuracy.

The choice we have taken in the model here may not be the appropriate one. But you can see the accuracy you are able to come up with without having overfitting happen there.

Question 3. Next from `keras.regularizers` import `l2`

Now we will try to include L2 regularization in our modeling part for performing regularization. You can add this weight decay in any of your hidden layers like.

```
model.add(Dense(256, activation='relu', kernel_regularizer=l2(0.001)))
```

The parameter 0.001 is an arbitrary choice for the regularization parameter α that we saw in the class.

Now try to fit your same model as above with this change and check for any improvements.

Question 4: Now, let's see how we can proceed to do perform some hyperparameter tuning and find out the appropriate parameter value. The following part is done for a very simple model with one hidden layer and an output layer. The number of neurons and the dropout parameter is being tuned to find appropriate ones.

What you see above is that we have loaded and pre-processed the data and created a model that would be re-used for the tuning part.

Next, we will set up the `keras_tuner` that can be used for the tuning part using `tuner.search()`. So our objective is to maximize `val_accuracy`. We are using `randomsearch` with a maximum 10 combinations of our hyperparameters with each `execution_per_trial` as 1. A directory is created on the disk with the name `mnist_tuning` which stores logs and information.

Now using `tuner.search` method, we can try to get the best model. Here I have only given 10 epochs. you can increase as well.

Now if you want to re-tune your model by adding few more layers etc then you have to first delete the already stored information in the directory `mnist_tuning`. For that you can use this

```

import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.utils import to_categorical
from keras.optimizers import SGD
import keras_tuner as kt

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(-1, 28 * 28).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28 * 28).astype('float32') / 255.0

print(x_train.shape)

y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

def build_model(hp):
    model = Sequential()
    model.add(Flatten(input_shape=(28*28,)))

    units = hp.Int('units', min_value=64, max_value=512, step=64)
    model.add(Dense(units, activation='relu'))

    dropout_rate = hp.Float('dropout', min_value=0.0, max_value=0.5, step=0.1)
    model.add(Dropout(dropout_rate))

    model.add(Dense(10, activation='softmax'))

    model.compile(
        optimizer=SGD(),
        loss='categorical_crossentropy',
        metrics=['accuracy']
    )
    return model

```

```

tuner = kt.RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=10,
    executions_per_trial=1,
    directory='mnist_tuning',
    project_name='dense_dropout_tune'
)

tuner.search(x_train, y_train,
            epochs=10,
            validation_split=0.2,
            batch_size=128,
            callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)])

best_model = tuner.get_best_models(num_models=1)[0]

test_loss, test_acc = best_model.evaluate(x_test, y_test)
print("Test accuracy:", test_acc)

best_hps = tuner.get_best_hyperparameters(1)[0]
print("Best units:", best_hps.get('units'))
print("Best dropout:", best_hps.get('dropout'))

```

import shutil

shutil.rmtree('mnist_tuning/dense_dropout_tune', ignore_errors=True)

Then again you can re-run.

So your work is to include L2 regularization and tune the above model to get better accuracy with mnist.

Then later go back to the cifar10 dataset problem and come up with your best model.