


## DA\_2

September 1, 2025

Name:	Tufan Kundu
Registration no:	24MDT0184
Course Name:	Deep Learning Lab
Course Code:	PMDS603P
Digital Assessment:	2

0.1 Question1. Today, we will try to recall the work done in the previous lab first. The second problem attempted in the last lab was to use MNIST dataset which contains handwritten numbers (their images) from 0 to 9 digits. First try to fit a simple neural network model. Let us import the necessary modules required for this along with the dataset. It contains 70000 handwritten images of digits from 0 to 9. So its a 10 class classification problem. Lets try to create a model that can do the classification task.

```
[4]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense,Dropout,Flatten
from keras.optimizers import SGD
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
batch_size = 128
num_classes = 10
epochs = 50
(x_train,y_train), (x_test,y_test) = mnist.load_data()
plt.imshow(x_train[3],cmap='gray')
plt.axis('off')
plt.show()
```



```
[5]: x_train = x_train.reshape(60000,784)
x_test = x_test.reshape(10000,784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train/=255
x_test/=255
print(x_train.shape[0],'train samples')
print(x_test.shape[0],'test samples')
y_train = keras.utils.to_categorical(y_train,num_classes)
y_test_ = keras.utils.to_categorical(y_test,num_classes)
```

60000 train samples

10000 test samples

### 0.1.1 Without dropout with ReLU activation

```
[6]: model = Sequential()
model.add(Dense(512, activation = 'relu',input_shape = (784,)))
model.add(Dense(512, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))
model.summary()
sgd1 = SGD(learning_rate=0.01)
model.compile(loss = 'categorical_crossentropy', optimizer = sgd1, metrics = _
↳ ['accuracy'])
```

```
history = model.  
    ↪fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test,y
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401,920
dense_4 (Dense)	(None, 512)	262,656
dense_5 (Dense)	(None, 10)	5,130

Total params: 669,706 (2.55 MB)

Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/50

469/469                      5s 9ms/step -  
accuracy: 0.6011 - loss: 1.6053 - val\_accuracy: 0.8828 - val\_loss: 0.5087

Epoch 2/50

469/469                      3s 6ms/step -  
accuracy: 0.8802 - loss: 0.4768 - val\_accuracy: 0.9051 - val\_loss: 0.3595

Epoch 3/50

469/469                      3s 7ms/step -  
accuracy: 0.8989 - loss: 0.3613 - val\_accuracy: 0.9145 - val\_loss: 0.3125

Epoch 4/50

469/469                      3s 7ms/step -  
accuracy: 0.9107 - loss: 0.3174 - val\_accuracy: 0.9214 - val\_loss: 0.2854

Epoch 5/50

469/469                      3s 6ms/step -  
accuracy: 0.9158 - loss: 0.2946 - val\_accuracy: 0.9269 - val\_loss: 0.2666

Epoch 6/50

469/469                      3s 6ms/step -  
accuracy: 0.9220 - loss: 0.2774 - val\_accuracy: 0.9301 - val\_loss: 0.2509

Epoch 7/50

469/469                      3s 7ms/step -  
accuracy: 0.9259 - loss: 0.2599 - val\_accuracy: 0.9323 - val\_loss: 0.2413

Epoch 8/50

469/469                      3s 7ms/step -  
accuracy: 0.9323 - loss: 0.2392 - val\_accuracy: 0.9373 - val\_loss: 0.2276

Epoch 9/50  
469/469                    3s 6ms/step -  
accuracy: 0.9350 - loss: 0.2308 - val\_accuracy: 0.9384 - val\_loss: 0.2186

Epoch 10/50  
469/469                    3s 7ms/step -  
accuracy: 0.9360 - loss: 0.2213 - val\_accuracy: 0.9418 - val\_loss: 0.2096

Epoch 11/50  
469/469                    3s 6ms/step -  
accuracy: 0.9404 - loss: 0.2095 - val\_accuracy: 0.9437 - val\_loss: 0.2027

Epoch 12/50  
469/469                    3s 6ms/step -  
accuracy: 0.9427 - loss: 0.2047 - val\_accuracy: 0.9450 - val\_loss: 0.1952

Epoch 13/50  
469/469                    3s 6ms/step -  
accuracy: 0.9445 - loss: 0.1950 - val\_accuracy: 0.9469 - val\_loss: 0.1887

Epoch 14/50  
469/469                    3s 7ms/step -  
accuracy: 0.9462 - loss: 0.1879 - val\_accuracy: 0.9481 - val\_loss: 0.1813

Epoch 15/50  
469/469                    3s 7ms/step -  
accuracy: 0.9484 - loss: 0.1783 - val\_accuracy: 0.9499 - val\_loss: 0.1750

Epoch 16/50  
469/469                    3s 7ms/step -  
accuracy: 0.9511 - loss: 0.1703 - val\_accuracy: 0.9510 - val\_loss: 0.1699

Epoch 17/50  
469/469                    3s 7ms/step -  
accuracy: 0.9543 - loss: 0.1621 - val\_accuracy: 0.9514 - val\_loss: 0.1651

Epoch 18/50  
469/469                    3s 7ms/step -  
accuracy: 0.9543 - loss: 0.1602 - val\_accuracy: 0.9532 - val\_loss: 0.1604

Epoch 19/50  
469/469                    3s 6ms/step -  
accuracy: 0.9560 - loss: 0.1575 - val\_accuracy: 0.9554 - val\_loss: 0.1554

Epoch 20/50  
469/469                    3s 6ms/step -  
accuracy: 0.9575 - loss: 0.1505 - val\_accuracy: 0.9546 - val\_loss: 0.1531

Epoch 21/50  
469/469                    3s 6ms/step -  
accuracy: 0.9602 - loss: 0.1442 - val\_accuracy: 0.9573 - val\_loss: 0.1477

Epoch 22/50  
469/469                    3s 7ms/step -  
accuracy: 0.9605 - loss: 0.1399 - val\_accuracy: 0.9573 - val\_loss: 0.1448

Epoch 23/50  
469/469                    3s 7ms/step -  
accuracy: 0.9627 - loss: 0.1343 - val\_accuracy: 0.9588 - val\_loss: 0.1402

Epoch 24/50  
469/469                    3s 6ms/step -  
accuracy: 0.9623 - loss: 0.1338 - val\_accuracy: 0.9608 - val\_loss: 0.1368

Epoch 25/50  
 469/469            3s 6ms/step -  
 accuracy: 0.9643 - loss: 0.1273 - val\_accuracy: 0.9612 - val\_loss: 0.1336

Epoch 26/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9655 - loss: 0.1246 - val\_accuracy: 0.9610 - val\_loss: 0.1322

Epoch 27/50  
 469/469            4s 8ms/step -  
 accuracy: 0.9650 - loss: 0.1225 - val\_accuracy: 0.9625 - val\_loss: 0.1279

Epoch 28/50  
 469/469            4s 7ms/step -  
 accuracy: 0.9673 - loss: 0.1169 - val\_accuracy: 0.9632 - val\_loss: 0.1269

Epoch 29/50  
 469/469            3s 6ms/step -  
 accuracy: 0.9692 - loss: 0.1129 - val\_accuracy: 0.9638 - val\_loss: 0.1238

Epoch 30/50  
 469/469            3s 6ms/step -  
 accuracy: 0.9683 - loss: 0.1120 - val\_accuracy: 0.9645 - val\_loss: 0.1216

Epoch 31/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9695 - loss: 0.1093 - val\_accuracy: 0.9649 - val\_loss: 0.1190

Epoch 32/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9714 - loss: 0.1053 - val\_accuracy: 0.9656 - val\_loss: 0.1170

Epoch 33/50  
 469/469            4s 7ms/step -  
 accuracy: 0.9704 - loss: 0.1042 - val\_accuracy: 0.9664 - val\_loss: 0.1159

Epoch 34/50  
 469/469            7s 14ms/step -  
 accuracy: 0.9714 - loss: 0.1033 - val\_accuracy: 0.9668 - val\_loss: 0.1133

Epoch 35/50  
 469/469            4s 8ms/step -  
 accuracy: 0.9734 - loss: 0.0968 - val\_accuracy: 0.9667 - val\_loss: 0.1122

Epoch 36/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9734 - loss: 0.0944 - val\_accuracy: 0.9677 - val\_loss: 0.1091

Epoch 37/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9739 - loss: 0.0940 - val\_accuracy: 0.9682 - val\_loss: 0.1076

Epoch 38/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9757 - loss: 0.0879 - val\_accuracy: 0.9684 - val\_loss: 0.1057

Epoch 39/50  
 469/469            3s 7ms/step -  
 accuracy: 0.9758 - loss: 0.0879 - val\_accuracy: 0.9689 - val\_loss: 0.1044

Epoch 40/50  
 469/469            5s 10ms/step -  
 accuracy: 0.9762 - loss: 0.0855 - val\_accuracy: 0.9693 - val\_loss: 0.1027

```

Epoch 41/50
469/469          4s 8ms/step -
accuracy: 0.9765 - loss: 0.0850 - val_accuracy: 0.9694 - val_loss: 0.1021
Epoch 42/50
469/469          4s 8ms/step -
accuracy: 0.9779 - loss: 0.0823 - val_accuracy: 0.9695 - val_loss: 0.1014
Epoch 43/50
469/469          5s 10ms/step -
accuracy: 0.9777 - loss: 0.0811 - val_accuracy: 0.9706 - val_loss: 0.0992
Epoch 44/50
469/469          7s 15ms/step -
accuracy: 0.9784 - loss: 0.0787 - val_accuracy: 0.9705 - val_loss: 0.0979
Epoch 45/50
469/469          8s 11ms/step -
accuracy: 0.9788 - loss: 0.0782 - val_accuracy: 0.9708 - val_loss: 0.0978
Epoch 46/50
469/469          4s 9ms/step -
accuracy: 0.9795 - loss: 0.0778 - val_accuracy: 0.9710 - val_loss: 0.0955
Epoch 47/50
469/469          5s 9ms/step -
accuracy: 0.9797 - loss: 0.0745 - val_accuracy: 0.9718 - val_loss: 0.0956
Epoch 48/50
469/469          3s 7ms/step -
accuracy: 0.9802 - loss: 0.0719 - val_accuracy: 0.9719 - val_loss: 0.0943
Epoch 49/50
469/469          4s 8ms/step -
accuracy: 0.9812 - loss: 0.0705 - val_accuracy: 0.9713 - val_loss: 0.0927
Epoch 50/50
469/469          3s 7ms/step -
accuracy: 0.9823 - loss: 0.0677 - val_accuracy: 0.9724 - val_loss: 0.0912

```

```

[7]: score = model.evaluate(x_test,y_test_, verbose = 1)
      print("Test loss:", score[0])
      print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          1s 3ms/step -
accuracy: 0.9668 - loss: 0.1081
Test loss: 0.09124796092510223
Test Accuracy:97.24%


```

```

[8]: plt.figure(figsize=(12,4))
      plt.subplot(1,2,1)
      plt.plot(history.history['accuracy'], label='Train Acc')
      plt.plot(history.history['val_accuracy'], label='Val Acc')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()
      plt.subplot(1,2,2)

```

```
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



## 0.2 Without dropout using sigmoid activation

```
[9]: model = Sequential()
model.add(Dense(512, activation = 'sigmoid',input_shape = (784,)))
model.add(Dense(512, activation = 'sigmoid'))
model.add(Dense(10, activation = 'softmax'))
model.summary()
sgd1 = SGD(learning_rate=0.01)
model.compile(loss = 'categorical_crossentropy', optimizer = sgd1, metrics = ['accuracy'])
history = model.
    fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test,y_test))
score = model.evaluate(x_test,y_test_, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 512)	401,920
dense_7 (Dense)	(None, 512)	262,656

dense\_8 (Dense)

(None, 10)

5,130

Total params: 669,706 (2.55 MB)

Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/50

469/469 5s 9ms/step -

accuracy: 0.1307 - loss: 2.2957 - val\_accuracy: 0.2895 - val\_loss: 2.2310

Epoch 2/50

469/469 4s 8ms/step -

accuracy: 0.3464 - loss: 2.2137 - val\_accuracy: 0.3752 - val\_loss: 2.1390

Epoch 3/50

469/469 3s 7ms/step -

accuracy: 0.4897 - loss: 2.1150 - val\_accuracy: 0.6444 - val\_loss: 2.0076

Epoch 4/50

469/469 3s 7ms/step -

accuracy: 0.5914 - loss: 1.9729 - val\_accuracy: 0.5940 - val\_loss: 1.8208

Epoch 5/50

469/469 3s 7ms/step -

accuracy: 0.6405 - loss: 1.7719 - val\_accuracy: 0.7126 - val\_loss: 1.5753

Epoch 6/50

469/469 4s 8ms/step -

accuracy: 0.6928 - loss: 1.5271 - val\_accuracy: 0.7186 - val\_loss: 1.3310

Epoch 7/50

469/469 5s 11ms/step -

accuracy: 0.7367 - loss: 1.2956 - val\_accuracy: 0.7665 - val\_loss: 1.1272

Epoch 8/50

469/469 4s 8ms/step -

accuracy: 0.7633 - loss: 1.1038 - val\_accuracy: 0.7860 - val\_loss: 0.9749

Epoch 9/50

469/469 3s 7ms/step -

accuracy: 0.7829 - loss: 0.9613 - val\_accuracy: 0.8059 - val\_loss: 0.8597

Epoch 10/50

469/469 3s 7ms/step -

accuracy: 0.8022 - loss: 0.8531 - val\_accuracy: 0.8145 - val\_loss: 0.7728

Epoch 11/50

469/469 4s 8ms/step -

accuracy: 0.8178 - loss: 0.7701 - val\_accuracy: 0.8305 - val\_loss: 0.7045

Epoch 12/50

469/469 3s 7ms/step -

accuracy: 0.8279 - loss: 0.7078 - val\_accuracy: 0.8367 - val\_loss: 0.6524

Epoch 13/50



469/469                    3s 7ms/step -  
 accuracy: 0.8368 - loss: 0.6597 - val\_accuracy: 0.8455 - val\_loss: 0.6090  
 Epoch 14/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8461 - loss: 0.6126 - val\_accuracy: 0.8567 - val\_loss: 0.5716  
 Epoch 15/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8530 - loss: 0.5838 - val\_accuracy: 0.8609 - val\_loss: 0.5424  
 Epoch 16/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8572 - loss: 0.5504 - val\_accuracy: 0.8656 - val\_loss: 0.5164  
 Epoch 17/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8638 - loss: 0.5206 - val\_accuracy: 0.8695 - val\_loss: 0.4945  
 Epoch 18/50  
 469/469                    3s 6ms/step -  
 accuracy: 0.8632 - loss: 0.5141 - val\_accuracy: 0.8738 - val\_loss: 0.4764  
 Epoch 19/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8701 - loss: 0.4913 - val\_accuracy: 0.8769 - val\_loss: 0.4600  
 Epoch 20/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8711 - loss: 0.4760 - val\_accuracy: 0.8813 - val\_loss: 0.4462  
 Epoch 21/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8785 - loss: 0.4564 - val\_accuracy: 0.8829 - val\_loss: 0.4341  
 Epoch 22/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8791 - loss: 0.4465 - val\_accuracy: 0.8860 - val\_loss: 0.4233  
 Epoch 23/50  
 469/469                    4s 9ms/step -  
 accuracy: 0.8815 - loss: 0.4374 - val\_accuracy: 0.8866 - val\_loss: 0.4147  
 Epoch 24/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8819 - loss: 0.4281 - val\_accuracy: 0.8899 - val\_loss: 0.4060  
 Epoch 25/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8859 - loss: 0.4187 - val\_accuracy: 0.8916 - val\_loss: 0.3981  
 Epoch 26/50  
 469/469                    4s 8ms/step -  
 accuracy: 0.8856 - loss: 0.4150 - val\_accuracy: 0.8921 - val\_loss: 0.3911  
 Epoch 27/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8889 - loss: 0.4053 - val\_accuracy: 0.8939 - val\_loss: 0.3838  
 Epoch 28/50  
 469/469                    3s 7ms/step -  
 accuracy: 0.8879 - loss: 0.3982 - val\_accuracy: 0.8959 - val\_loss: 0.3787  
 Epoch 29/50

469/469                    5s 10ms/step -  
accuracy: 0.8915 - loss: 0.3929 - val\_accuracy: 0.8949 - val\_loss: 0.3736  
Epoch 30/50

469/469                    3s 6ms/step -  
accuracy: 0.8921 - loss: 0.3888 - val\_accuracy: 0.8966 - val\_loss: 0.3689  
Epoch 31/50

469/469                    3s 7ms/step -  
accuracy: 0.8946 - loss: 0.3794 - val\_accuracy: 0.8988 - val\_loss: 0.3635  
Epoch 32/50

469/469                    4s 8ms/step -  
accuracy: 0.8930 - loss: 0.3759 - val\_accuracy: 0.8992 - val\_loss: 0.3599  
Epoch 33/50

469/469                    3s 7ms/step -  
accuracy: 0.8927 - loss: 0.3809 - val\_accuracy: 0.9008 - val\_loss: 0.3557  
Epoch 34/50

469/469                    3s 6ms/step -  
accuracy: 0.8979 - loss: 0.3624 - val\_accuracy: 0.9003 - val\_loss: 0.3516  
Epoch 35/50

469/469                    3s 6ms/step -  
accuracy: 0.8975 - loss: 0.3633 - val\_accuracy: 0.9008 - val\_loss: 0.3485  
Epoch 36/50

469/469                    3s 6ms/step -  
accuracy: 0.8955 - loss: 0.3685 - val\_accuracy: 0.9008 - val\_loss: 0.3465  
Epoch 37/50

469/469                    3s 7ms/step -  
accuracy: 0.9005 - loss: 0.3517 - val\_accuracy: 0.9016 - val\_loss: 0.3425  
Epoch 38/50

469/469                    4s 9ms/step -  
accuracy: 0.9001 - loss: 0.3515 - val\_accuracy: 0.9028 - val\_loss: 0.3403  
Epoch 39/50

469/469                    5s 10ms/step -  
accuracy: 0.9010 - loss: 0.3501 - val\_accuracy: 0.9034 - val\_loss: 0.3377  
Epoch 40/50

469/469                    4s 9ms/step -  
accuracy: 0.8989 - loss: 0.3556 - val\_accuracy: 0.9042 - val\_loss: 0.3346  
Epoch 41/50

469/469                    4s 8ms/step -  
accuracy: 0.9002 - loss: 0.3516 - val\_accuracy: 0.9042 - val\_loss: 0.3323  
Epoch 42/50

469/469                    3s 7ms/step -  
accuracy: 0.9012 - loss: 0.3486 - val\_accuracy: 0.9052 - val\_loss: 0.3320  
Epoch 43/50

469/469                    3s 6ms/step -  
accuracy: 0.9015 - loss: 0.3465 - val\_accuracy: 0.9056 - val\_loss: 0.3291  
Epoch 44/50

469/469                    3s 6ms/step -  
accuracy: 0.9021 - loss: 0.3456 - val\_accuracy: 0.9051 - val\_loss: 0.3261  
Epoch 45/50

```


469/469          3s 7ms/step -
accuracy: 0.9047 - loss: 0.3332 - val_accuracy: 0.9045 - val_loss: 0.3244
Epoch 46/50
469/469          4s 7ms/step -
accuracy: 0.9030 - loss: 0.3359 - val_accuracy: 0.9063 - val_loss: 0.3218
Epoch 47/50
469/469          3s 7ms/step -
accuracy: 0.9048 - loss: 0.3345 - val_accuracy: 0.9074 - val_loss: 0.3202
Epoch 48/50
469/469          3s 7ms/step -
accuracy: 0.9060 - loss: 0.3295 - val_accuracy: 0.9086 - val_loss: 0.3186
Epoch 49/50
469/469          3s 7ms/step -
accuracy: 0.9059 - loss: 0.3331 - val_accuracy: 0.9093 - val_loss: 0.3174
Epoch 50/50
469/469          4s 9ms/step -
accuracy: 0.9075 - loss: 0.3280 - val_accuracy: 0.9089 - val_loss: 0.3157
313/313          1s 5ms/step -
accuracy: 0.8966 - loss: 0.3592
Test loss: 0.3157140612602234
Test Accuracy:90.89%

```

```

[10]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```



## 0.3 Regularization Techniques

### 0.3.1 Using dropout(0.2)

```
[11]: model = Sequential()
model.add(Dense(512, activation = 'relu',input_shape = (784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation = 'softmax'))
model.summary()
sgd1 = SGD(learning_rate=0.01)
model.compile(loss = 'categorical_crossentropy', optimizer = sgd1, metrics = ['accuracy'])
history = model.fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test,y_test))
score = model.evaluate(x_test,y_test, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 512)	401,920
dropout (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 512)	262,656

dropout_1 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 10)	5,130

Total params: 669,706 (2.55 MB)

Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/50

469/469 8s 14ms/step -  
accuracy: 0.4983 - loss: 1.7143 - val\_accuracy: 0.8674 - val\_loss: 0.5604

Epoch 2/50

469/469 5s 10ms/step -  
accuracy: 0.8367 - loss: 0.5906 - val\_accuracy: 0.8942 - val\_loss: 0.3843

Epoch 3/50

469/469 4s 9ms/step -  
accuracy: 0.8694 - loss: 0.4529 - val\_accuracy: 0.9083 - val\_loss: 0.3272

Epoch 4/50

469/469 4s 9ms/step -  
accuracy: 0.8891 - loss: 0.3852 - val\_accuracy: 0.9164 - val\_loss: 0.2954

Epoch 5/50

469/469 4s 8ms/step -  
accuracy: 0.8969 - loss: 0.3535 - val\_accuracy: 0.9227 - val\_loss: 0.2733

Epoch 6/50

469/469 4s 8ms/step -  
accuracy: 0.9035 - loss: 0.3264 - val\_accuracy: 0.9254 - val\_loss: 0.2565

Epoch 7/50

469/469 4s 8ms/step -  
accuracy: 0.9127 - loss: 0.3057 - val\_accuracy: 0.9301 - val\_loss: 0.2421

Epoch 8/50

469/469 4s 8ms/step -  
accuracy: 0.9179 - loss: 0.2835 - val\_accuracy: 0.9342 - val\_loss: 0.2294

Epoch 9/50

469/469 4s 8ms/step -  
accuracy: 0.9208 - loss: 0.2781 - val\_accuracy: 0.9366 - val\_loss: 0.2189

Epoch 10/50

469/469 4s 8ms/step -  
accuracy: 0.9258 - loss: 0.2575 - val\_accuracy: 0.9385 - val\_loss: 0.2095

Epoch 11/50

469/469 4s 8ms/step -  
accuracy: 0.9290 - loss: 0.2486 - val\_accuracy: 0.9410 - val\_loss: 0.1999

Epoch 12/50

469/469 4s 7ms/step -

accuracy: 0.9304 - loss: 0.2378 - val\_accuracy: 0.9444 - val\_loss: 0.1907  
 Epoch 13/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9351 - loss: 0.2267 - val\_accuracy: 0.9455 - val\_loss: 0.1842  
 Epoch 14/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9370 - loss: 0.2153 - val\_accuracy: 0.9490 - val\_loss: 0.1772  
 Epoch 15/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9389 - loss: 0.2102 - val\_accuracy: 0.9508 - val\_loss: 0.1708  
 Epoch 16/50  
 469/469 4s 9ms/step -  
 accuracy: 0.9417 - loss: 0.2029 - val\_accuracy: 0.9513 - val\_loss: 0.1658  
 Epoch 17/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9449 - loss: 0.1921 - val\_accuracy: 0.9530 - val\_loss: 0.1594  
 Epoch 18/50  
 469/469 4s 9ms/step -  
 accuracy: 0.9464 - loss: 0.1857 - val\_accuracy: 0.9541 - val\_loss: 0.1545  
 Epoch 19/50  
 469/469 3s 7ms/step -  
 accuracy: 0.9484 - loss: 0.1794 - val\_accuracy: 0.9555 - val\_loss: 0.1495  
 Epoch 20/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9484 - loss: 0.1775 - val\_accuracy: 0.9562 - val\_loss: 0.1447  
 Epoch 21/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9509 - loss: 0.1704 - val\_accuracy: 0.9575 - val\_loss: 0.1408  
 Epoch 22/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9530 - loss: 0.1635 - val\_accuracy: 0.9584 - val\_loss: 0.1371  
 Epoch 23/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9534 - loss: 0.1632 - val\_accuracy: 0.9589 - val\_loss: 0.1351  
 Epoch 24/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9546 - loss: 0.1564 - val\_accuracy: 0.9605 - val\_loss: 0.1312  
 Epoch 25/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9562 - loss: 0.1542 - val\_accuracy: 0.9610 - val\_loss: 0.1277  
 Epoch 26/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9569 - loss: 0.1489 - val\_accuracy: 0.9619 - val\_loss: 0.1256  
 Epoch 27/50  
 469/469 5s 10ms/step -  
 accuracy: 0.9565 - loss: 0.1482 - val\_accuracy: 0.9622 - val\_loss: 0.1221  
 Epoch 28/50  
 469/469 4s 8ms/step -

accuracy: 0.9577 - loss: 0.1424 - val\_accuracy: 0.9636 - val\_loss: 0.1195  
 Epoch 29/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9592 - loss: 0.1391 - val\_accuracy: 0.9639 - val\_loss: 0.1168  
 Epoch 30/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9626 - loss: 0.1319 - val\_accuracy: 0.9648 - val\_loss: 0.1150  
 Epoch 31/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9612 - loss: 0.1313 - val\_accuracy: 0.9649 - val\_loss: 0.1137  
 Epoch 32/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9601 - loss: 0.1348 - val\_accuracy: 0.9658 - val\_loss: 0.1111  
 Epoch 33/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9637 - loss: 0.1251 - val\_accuracy: 0.9665 - val\_loss: 0.1092  
 Epoch 34/50  
 469/469 4s 7ms/step -  
 accuracy: 0.9633 - loss: 0.1246 - val\_accuracy: 0.9673 - val\_loss: 0.1071  
 Epoch 35/50  
 469/469 4s 7ms/step -  
 accuracy: 0.9646 - loss: 0.1237 - val\_accuracy: 0.9678 - val\_loss: 0.1054  
 Epoch 36/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9652 - loss: 0.1221 - val\_accuracy: 0.9681 - val\_loss: 0.1033  
 Epoch 37/50  
 469/469 4s 9ms/step -  
 accuracy: 0.9661 - loss: 0.1181 - val\_accuracy: 0.9690 - val\_loss: 0.1011  
 Epoch 38/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9662 - loss: 0.1148 - val\_accuracy: 0.9687 - val\_loss: 0.1002  
 Epoch 39/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9673 - loss: 0.1109 - val\_accuracy: 0.9694 - val\_loss: 0.0985  
 Epoch 40/50  
 469/469 4s 7ms/step -  
 accuracy: 0.9681 - loss: 0.1125 - val\_accuracy: 0.9702 - val\_loss: 0.0974  
 Epoch 41/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9679 - loss: 0.1084 - val\_accuracy: 0.9706 - val\_loss: 0.0969  
 Epoch 42/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9703 - loss: 0.1026 - val\_accuracy: 0.9704 - val\_loss: 0.0950  
 Epoch 43/50  
 469/469 4s 8ms/step -  
 accuracy: 0.9698 - loss: 0.1040 - val\_accuracy: 0.9708 - val\_loss: 0.0937  
 Epoch 44/50  
 469/469 4s 8ms/step -

```

accuracy: 0.9700 - loss: 0.1033 - val_accuracy: 0.9716 - val_loss: 0.0920
Epoch 45/50
469/469          4s 8ms/step -
accuracy: 0.9702 - loss: 0.1034 - val_accuracy: 0.9718 - val_loss: 0.0914
Epoch 46/50
469/469          3s 7ms/step -
accuracy: 0.9720 - loss: 0.0968 - val_accuracy: 0.9719 - val_loss: 0.0907
Epoch 47/50
469/469          4s 8ms/step -
accuracy: 0.9726 - loss: 0.0963 - val_accuracy: 0.9730 - val_loss: 0.0890
Epoch 48/50
469/469          4s 8ms/step -
accuracy: 0.9720 - loss: 0.0955 - val_accuracy: 0.9728 - val_loss: 0.0881
Epoch 49/50
469/469          4s 8ms/step -
accuracy: 0.9723 - loss: 0.0944 - val_accuracy: 0.9734 - val_loss: 0.0869
Epoch 50/50
469/469          4s 8ms/step -
accuracy: 0.9739 - loss: 0.0916 - val_accuracy: 0.9736 - val_loss: 0.0864
313/313          1s 3ms/step -
accuracy: 0.9689 - loss: 0.1017
Test loss: 0.08635895699262619
Test Accuracy:97.36%


```

```

[12]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```





## 0.4 Early stopping

```
[14]: from sklearn.model_selection import train_test_split
(x_train,y_train),(x_test,y_test) = mnist.load_data()
x_subtrain,x_valid,y_subtrain,y_valid = \
    train_test_split(x_train,y_train,test_size = 0.10, random_state = 1)
x_train = x_train/255
x_test = x_test/255
x_subtrain = x_subtrain/255
x_valid=x_valid/255
```

### 0.4.1 making the ANN

```
[15]: model = Sequential()
model.add(Flatten(input_shape = (28,28)))
model.add(Dense(512, activation = 'relu'))
model.add(Dense(512, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))
model.summary()
sgd1 = SGD(learning_rate=0.01)
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense_12 (Dense)	(None, 512)	401,920
dense_13 (Dense)	(None, 512)	262,656

dense\_14 (Dense) (None, 10) 5,130

Total params: 669,706 (2.55 MB)

Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)

```
[16]: from keras.callbacks import EarlyStopping
model.compile(loss= 'sparse_categorical_crossentropy',optimizer = sgd1,metrics_
↳ ['accuracy'])
estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-3, mode = 'min',_
↳ patience = 4, verbose = 1, restore_best_weights= True)
history = model.fit(x_subtrain,y_subtrain, batch_size=batch_size, epochs = 100,_
↳ verbose = 1, validation_data=(x_valid,y_valid))
```

Epoch 1/100

422/422 4s 8ms/step -

accuracy: 0.5949 - loss: 1.6476 - val\_accuracy: 0.8600 - val\_loss: 0.5826

Epoch 2/100

422/422 3s 7ms/step -

accuracy: 0.8701 - loss: 0.5185 - val\_accuracy: 0.8885 - val\_loss: 0.4104

Epoch 3/100

422/422 3s 7ms/step -

accuracy: 0.8935 - loss: 0.3859 - val\_accuracy: 0.9037 - val\_loss: 0.3525

Epoch 4/100

422/422 3s 6ms/step -

accuracy: 0.9050 - loss: 0.3394 - val\_accuracy: 0.9103 - val\_loss: 0.3219

Epoch 5/100

422/422 3s 7ms/step -

accuracy: 0.9129 - loss: 0.3069 - val\_accuracy: 0.9160 - val\_loss: 0.3012

Epoch 6/100

422/422 3s 6ms/step -

accuracy: 0.9180 - loss: 0.2882 - val\_accuracy: 0.9188 - val\_loss: 0.2832

Epoch 7/100

422/422 3s 6ms/step -

accuracy: 0.9247 - loss: 0.2650 - val\_accuracy: 0.9240 - val\_loss: 0.2690

Epoch 8/100

422/422 3s 6ms/step -

accuracy: 0.9286 - loss: 0.2531 - val\_accuracy: 0.9263 - val\_loss: 0.2579

Epoch 9/100

422/422 3s 7ms/step -

accuracy: 0.9307 - loss: 0.2442 - val\_accuracy: 0.9292 - val\_loss: 0.2479

Epoch 10/100

422/422                    3s 6ms/step -  
 accuracy: 0.9343 - loss: 0.2264 - val\_accuracy: 0.9320 - val\_loss: 0.2387  
 Epoch 11/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9356 - loss: 0.2237 - val\_accuracy: 0.9342 - val\_loss: 0.2303  
 Epoch 12/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9414 - loss: 0.2073 - val\_accuracy: 0.9352 - val\_loss: 0.2252  
 Epoch 13/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9425 - loss: 0.2027 - val\_accuracy: 0.9365 - val\_loss: 0.2181  
 Epoch 14/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9439 - loss: 0.1951 - val\_accuracy: 0.9407 - val\_loss: 0.2099  
 Epoch 15/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9447 - loss: 0.1906 - val\_accuracy: 0.9418 - val\_loss: 0.2042  
 Epoch 16/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9497 - loss: 0.1794 - val\_accuracy: 0.9432 - val\_loss: 0.2003  
 Epoch 17/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9509 - loss: 0.1740 - val\_accuracy: 0.9430 - val\_loss: 0.1951  
 Epoch 18/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9502 - loss: 0.1740 - val\_accuracy: 0.9448 - val\_loss: 0.1894  
 Epoch 19/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9535 - loss: 0.1615 - val\_accuracy: 0.9470 - val\_loss: 0.1845  
 Epoch 20/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9550 - loss: 0.1623 - val\_accuracy: 0.9467 - val\_loss: 0.1825  
 Epoch 21/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9559 - loss: 0.1563 - val\_accuracy: 0.9488 - val\_loss: 0.1756  
 Epoch 22/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9576 - loss: 0.1476 - val\_accuracy: 0.9500 - val\_loss: 0.1722  
 Epoch 23/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9594 - loss: 0.1424 - val\_accuracy: 0.9500 - val\_loss: 0.1687  
 Epoch 24/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9604 - loss: 0.1405 - val\_accuracy: 0.9517 - val\_loss: 0.1661  
 Epoch 25/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9612 - loss: 0.1374 - val\_accuracy: 0.9513 - val\_loss: 0.1638  
 Epoch 26/100

422/422                    3s 6ms/step -  
 accuracy: 0.9627 - loss: 0.1315 - val\_accuracy: 0.9520 - val\_loss: 0.1595  
 Epoch 27/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9639 - loss: 0.1278 - val\_accuracy: 0.9532 - val\_loss: 0.1580  
 Epoch 28/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9633 - loss: 0.1277 - val\_accuracy: 0.9548 - val\_loss: 0.1562  
 Epoch 29/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9652 - loss: 0.1215 - val\_accuracy: 0.9567 - val\_loss: 0.1515  
 Epoch 30/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9666 - loss: 0.1195 - val\_accuracy: 0.9565 - val\_loss: 0.1488  
 Epoch 31/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9669 - loss: 0.1175 - val\_accuracy: 0.9577 - val\_loss: 0.1477  
 Epoch 32/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9680 - loss: 0.1117 - val\_accuracy: 0.9577 - val\_loss: 0.1450  
 Epoch 33/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9684 - loss: 0.1121 - val\_accuracy: 0.9582 - val\_loss: 0.1429  
 Epoch 34/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9695 - loss: 0.1081 - val\_accuracy: 0.9590 - val\_loss: 0.1402  
 Epoch 35/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9698 - loss: 0.1077 - val\_accuracy: 0.9597 - val\_loss: 0.1391  
 Epoch 36/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9696 - loss: 0.1061 - val\_accuracy: 0.9615 - val\_loss: 0.1364  
 Epoch 37/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9712 - loss: 0.1042 - val\_accuracy: 0.9607 - val\_loss: 0.1367  
 Epoch 38/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9718 - loss: 0.0989 - val\_accuracy: 0.9625 - val\_loss: 0.1335  
 Epoch 39/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9735 - loss: 0.0974 - val\_accuracy: 0.9628 - val\_loss: 0.1310  
 Epoch 40/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9733 - loss: 0.0936 - val\_accuracy: 0.9630 - val\_loss: 0.1294  
 Epoch 41/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9742 - loss: 0.0947 - val\_accuracy: 0.9622 - val\_loss: 0.1285  
 Epoch 42/100

422/422                    3s 7ms/step -  
 accuracy: 0.9758 - loss: 0.0882 - val\_accuracy: 0.9623 - val\_loss: 0.1285  
 Epoch 43/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9755 - loss: 0.0873 - val\_accuracy: 0.9628 - val\_loss: 0.1256  
 Epoch 44/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9759 - loss: 0.0872 - val\_accuracy: 0.9643 - val\_loss: 0.1240  
 Epoch 45/100  
 422/422                    3s 8ms/step -  
 accuracy: 0.9764 - loss: 0.0849 - val\_accuracy: 0.9650 - val\_loss: 0.1223  
 Epoch 46/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9770 - loss: 0.0832 - val\_accuracy: 0.9645 - val\_loss: 0.1214  
 Epoch 47/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9771 - loss: 0.0827 - val\_accuracy: 0.9653 - val\_loss: 0.1198  
 Epoch 48/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9782 - loss: 0.0794 - val\_accuracy: 0.9660 - val\_loss: 0.1192  
 Epoch 49/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9795 - loss: 0.0747 - val\_accuracy: 0.9672 - val\_loss: 0.1174  
 Epoch 50/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9788 - loss: 0.0787 - val\_accuracy: 0.9663 - val\_loss: 0.1168  
 Epoch 51/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9802 - loss: 0.0742 - val\_accuracy: 0.9668 - val\_loss: 0.1151  
 Epoch 52/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9799 - loss: 0.0743 - val\_accuracy: 0.9677 - val\_loss: 0.1139  
 Epoch 53/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9800 - loss: 0.0716 - val\_accuracy: 0.9670 - val\_loss: 0.1134  
 Epoch 54/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9814 - loss: 0.0704 - val\_accuracy: 0.9673 - val\_loss: 0.1128  
 Epoch 55/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9820 - loss: 0.0676 - val\_accuracy: 0.9683 - val\_loss: 0.1110  
 Epoch 56/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9808 - loss: 0.0704 - val\_accuracy: 0.9687 - val\_loss: 0.1108  
 Epoch 57/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9821 - loss: 0.0655 - val\_accuracy: 0.9680 - val\_loss: 0.1100  
 Epoch 58/100

422/422                    3s 7ms/step -  
 accuracy: 0.9832 - loss: 0.0634 - val\_accuracy: 0.9695 - val\_loss: 0.1086  
 Epoch 59/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9826 - loss: 0.0642 - val\_accuracy: 0.9683 - val\_loss: 0.1074  
 Epoch 60/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9844 - loss: 0.0609 - val\_accuracy: 0.9692 - val\_loss: 0.1070  
 Epoch 61/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9846 - loss: 0.0594 - val\_accuracy: 0.9697 - val\_loss: 0.1060  
 Epoch 62/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9844 - loss: 0.0618 - val\_accuracy: 0.9695 - val\_loss: 0.1049  
 Epoch 63/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9851 - loss: 0.0592 - val\_accuracy: 0.9700 - val\_loss: 0.1050  
 Epoch 64/100  
 422/422                    3s 8ms/step -  
 accuracy: 0.9853 - loss: 0.0576 - val\_accuracy: 0.9698 - val\_loss: 0.1045  
 Epoch 65/100  
 422/422                    5s 8ms/step -  
 accuracy: 0.9847 - loss: 0.0588 - val\_accuracy: 0.9712 - val\_loss: 0.1029  
 Epoch 66/100  
 422/422                    3s 6ms/step -  
 accuracy: 0.9863 - loss: 0.0546 - val\_accuracy: 0.9703 - val\_loss: 0.1029  
 Epoch 67/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9866 - loss: 0.0537 - val\_accuracy: 0.9705 - val\_loss: 0.1019  
 Epoch 68/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9861 - loss: 0.0528 - val\_accuracy: 0.9712 - val\_loss: 0.1018  
 Epoch 69/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9863 - loss: 0.0525 - val\_accuracy: 0.9700 - val\_loss: 0.1011  
 Epoch 70/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9869 - loss: 0.0524 - val\_accuracy: 0.9708 - val\_loss: 0.1009  
 Epoch 71/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9861 - loss: 0.0531 - val\_accuracy: 0.9715 - val\_loss: 0.1000  
 Epoch 72/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9871 - loss: 0.0513 - val\_accuracy: 0.9722 - val\_loss: 0.0985  
 Epoch 73/100  
 422/422                    3s 7ms/step -  
 accuracy: 0.9874 - loss: 0.0490 - val\_accuracy: 0.9725 - val\_loss: 0.0987  
 Epoch 74/100

422/422            3s 7ms/step -  
accuracy: 0.9882 - loss: 0.0483 - val\_accuracy: 0.9715 - val\_loss: 0.0998  
Epoch 75/100

422/422            3s 7ms/step -  
accuracy: 0.9867 - loss: 0.0514 - val\_accuracy: 0.9723 - val\_loss: 0.0991  
Epoch 76/100

422/422            3s 7ms/step -  
accuracy: 0.9879 - loss: 0.0476 - val\_accuracy: 0.9707 - val\_loss: 0.0981  
Epoch 77/100

422/422            3s 7ms/step -  
accuracy: 0.9888 - loss: 0.0449 - val\_accuracy: 0.9720 - val\_loss: 0.0979  
Epoch 78/100

422/422            3s 7ms/step -  
accuracy: 0.9896 - loss: 0.0443 - val\_accuracy: 0.9728 - val\_loss: 0.0965  
Epoch 79/100

422/422            3s 7ms/step -  
accuracy: 0.9895 - loss: 0.0434 - val\_accuracy: 0.9725 - val\_loss: 0.0954  
Epoch 80/100

422/422            3s 7ms/step -  
accuracy: 0.9896 - loss: 0.0426 - val\_accuracy: 0.9722 - val\_loss: 0.0957  
Epoch 81/100

422/422            3s 7ms/step -  
accuracy: 0.9887 - loss: 0.0450 - val\_accuracy: 0.9727 - val\_loss: 0.0960  
Epoch 82/100

422/422            3s 7ms/step -  
accuracy: 0.9895 - loss: 0.0442 - val\_accuracy: 0.9728 - val\_loss: 0.0947  
Epoch 83/100

422/422            3s 7ms/step -  
accuracy: 0.9900 - loss: 0.0411 - val\_accuracy: 0.9735 - val\_loss: 0.0941  
Epoch 84/100

422/422            3s 6ms/step -  
accuracy: 0.9903 - loss: 0.0421 - val\_accuracy: 0.9730 - val\_loss: 0.0944  
Epoch 85/100

422/422            3s 7ms/step -  
accuracy: 0.9893 - loss: 0.0421 - val\_accuracy: 0.9727 - val\_loss: 0.0933  
Epoch 86/100

422/422            4s 8ms/step -  
accuracy: 0.9906 - loss: 0.0402 - val\_accuracy: 0.9725 - val\_loss: 0.0942  
Epoch 87/100

422/422            3s 7ms/step -  
accuracy: 0.9906 - loss: 0.0391 - val\_accuracy: 0.9728 - val\_loss: 0.0937  
Epoch 88/100

422/422            3s 7ms/step -  
accuracy: 0.9909 - loss: 0.0397 - val\_accuracy: 0.9725 - val\_loss: 0.0935  
Epoch 89/100

422/422            4s 9ms/step -  
accuracy: 0.9912 - loss: 0.0378 - val\_accuracy: 0.9733 - val\_loss: 0.0932  
Epoch 90/100

```

422/422          3s 8ms/step -
accuracy: 0.9915 - loss: 0.0368 - val_accuracy: 0.9727 - val_loss: 0.0923
Epoch 91/100
422/422          3s 8ms/step -
accuracy: 0.9909 - loss: 0.0367 - val_accuracy: 0.9730 - val_loss: 0.0925
Epoch 92/100
422/422          4s 9ms/step -
accuracy: 0.9918 - loss: 0.0351 - val_accuracy: 0.9742 - val_loss: 0.0919
Epoch 93/100
422/422          3s 8ms/step -
accuracy: 0.9915 - loss: 0.0359 - val_accuracy: 0.9733 - val_loss: 0.0913
Epoch 94/100
422/422          4s 8ms/step -
accuracy: 0.9917 - loss: 0.0354 - val_accuracy: 0.9738 - val_loss: 0.0914
Epoch 95/100
422/422          4s 10ms/step -
accuracy: 0.9921 - loss: 0.0340 - val_accuracy: 0.9735 - val_loss: 0.0910
Epoch 96/100
422/422          4s 9ms/step -
accuracy: 0.9931 - loss: 0.0326 - val_accuracy: 0.9737 - val_loss: 0.0909
Epoch 97/100
422/422          3s 7ms/step -
accuracy: 0.9933 - loss: 0.0315 - val_accuracy: 0.9733 - val_loss: 0.0904
Epoch 98/100
422/422          3s 7ms/step -
accuracy: 0.9928 - loss: 0.0329 - val_accuracy: 0.9738 - val_loss: 0.0915
Epoch 99/100
422/422          3s 7ms/step -
accuracy: 0.9925 - loss: 0.0334 - val_accuracy: 0.9738 - val_loss: 0.0900
Epoch 100/100
422/422          4s 8ms/step -
accuracy: 0.9934 - loss: 0.0309 - val_accuracy: 0.9743 - val_loss: 0.0898

```

```

[17]: score = model.evaluate(x_test,y_test, verbose = 1)
      print("Test loss:", score[0])
      print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          1s 3ms/step -
accuracy: 0.9736 - loss: 0.0866
Test loss: 0.07452523708343506
Test Accuracy:97.74%

```


```

[18]: plt.figure(figsize=(12,4))
      plt.subplot(1,2,1)
      plt.plot(history.history['accuracy'], label='Train Acc')
      plt.plot(history.history['val_accuracy'], label='Val Acc')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')

```



```
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



**0.5 Challenging Question:** Try for a scratch code for this case where you can create a custom neural network without using any inbuilt classes like sequential etc. Where you need to define a class neural network which has methods like forwardpass, backwardpass, and train. Figure out how we can do this. This model has inputs as [0, 0, 1], [0, 1, 1], [1, 0, 1], [1, 1, 1] and the expected output as [0], [1], [1], [0] in each case. So there are three features in our dataset as you see above. The activation function is to be taken as sigmoid. The architecture is like we have only one hidden layer and an output layer with one neuron. Take the error function as  $(1/2)(y - \hat{y})^2$

```
[19]: import numpy as np

# sigmoid activation function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# derivative of sigmoid function
def sigmoid_derivative(x):
    return x*(1-x)

# mean squared error loss
def mse_loss(y_true,y_pred):
    return 0.5*np.mean((y_true-y_pred)**2)
```

```

# Input dataset (XOR gate inputs with bias term)
x = np.array([[0,0,1],
[0,1,1],
[1,0,1],
[1,1,1]])

# output labels
y = np.array([[0],
[1],
[1],
[0]])

# seed for reproducibility
np.random.seed(1)

# Initialize weights randomly with mean 0
input_size = 3 # 3 input features
hidden_size = 2 # 2 hidden layers
output_size = 1 # 1 output neuron
# Weights
w1 = 2 * np.random.random((input_size, hidden_size))-1
w2 = 2 * np.random.random((hidden_size, output_size))-1
# Biases
b1 = np.zeros((1, hidden_size))
b2 = np.zeros((1, output_size))
# Learning rate
lr = 0.1
# Training loop

for epoch in range(10000):
    ##----- Forward pass -----
    a1 = np.dot(x,w1) + b1
    h1 = sigmoid(a1) # activation of hidden layer
    a2 = np.dot(h1,w2) + b2
    output = sigmoid(a2) # final prediction
    # loss calculation
    loss = mse_loss(y,output)
    ##----- Back propagation -----
    # output layer error
    output_error = output - y
    output_delta = output_error * sigmoid_derivative(output)
    ## hidden layer error
    hidden_error = np.dot(output_delta, w2.T)
    hidden_delta = hidden_error * sigmoid_derivative(h1)
    ##-----Updating weights and biases -----
    w2 -= lr * np.dot(h1.T,output_delta)

```

```

b2 -= lr * np.sum(output_delta, axis = 0, keepdims = True)
w1 -= lr * np.dot(x.T, hidden_delta)
b1 -= lr * np.sum(hidden_delta, axis = 0, keepdims = True)
# Print loss every 1000 epochs
if epoch % 1000 == 0:
    print(f"Epoch {epoch}, Loss: {loss:.4f}")
# ----- Final Output -----
print("\nFinal predictions after training:")
print(output.round(3))

```

```

Epoch 0, Loss: 0.1267
Epoch 1000, Loss: 0.1215
Epoch 2000, Loss: 0.1029
Epoch 3000, Loss: 0.0905
Epoch 4000, Loss: 0.0828
Epoch 5000, Loss: 0.0433
Epoch 6000, Loss: 0.0105
Epoch 7000, Loss: 0.0049
Epoch 8000, Loss: 0.0031
Epoch 9000, Loss: 0.0022

```

```

Final predictions after training:
[[0.049]
 [0.945]
 [0.945]
 [0.071]]

```

## 0.6 Question2.

0.6.1 cifar10 dataset is also an inbuilt dataset which contains 10 classes of images, mainly, 0-airplane, 1-automobile, 2-bird, 3-cat, 4-deer, 5-dog, 6-frog, 7-horse, 8-ship, 9-truck. Load the inbuilt dataset cifar10 as you did in last lab by replacing `mnist.load_data()` as `cifar10.load_data()`. First, try to import it from `keras.datasets` as you did for `mnist`. Now, identify the size of the images you have first of all. You can now see  $32 * 32 * 3$  images that is  $32 * 32$  pixel images with 3 channels that give the RGB values since we have a color image. Try to print the shape of each image and see. you will see it's stored like  $32 * 32 * 3$  arrays. Now, try to visualize certain images using appropriate functions. Check the size of `x_train` and `x_test` and reshape them into one-dimensional arrays as done in the case of `mnist` dataset. Do necessary pre-processing and split the data into training, validation, and testing sets. Create a new model using a sequential class with appropriate hidden layers and output layer neurons. Choose appropriate activation functions like sigmoid and relu, etc. And also an appropriate one in the output layer. Choose the error function appropriately. Include early stopping technique in your model and run the model for 500 epochs. Try to come up with a better model with decent accuracy. The choice we have taken in the model here may not be the appropriate one. But you can see the accuracy you are able to come up with without having overfitting happen there.

### 0.6.2 Importing the necessary libraries

```
[20]: import keras
      from keras.datasets import cifar10
      from keras.models import Sequential
      from keras.layers import Dense, Dropout, Flatten
      from keras.optimizers import Adam, SGD
      import matplotlib.pyplot as plt
      from keras.callbacks import EarlyStopping
      from sklearn.model_selection import train_test_split
      from keras.utils import to_categorical
      from keras.regularizers import l2
      import warnings
      warnings.filterwarnings('ignore')
```

### 0.6.3 Loading the dataset and validation split

```
[21]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
      x_subtrain, x_valid, y_subtrain, y_valid = \
      ↪ train_test_split(x_train, y_train, test_size = 0.10, random_state = 1)
```

```
[22]: x_subtrain.shape
```

```
[22]: (45000, 32, 32, 3)
```

```
[23]: x_valid.shape
```

```
[23]: (5000, 32, 32, 3)
```

```
[24]: x_train = x_train/255  
x_test = x_test/255  
x_subtrain = x_subtrain/255  
x_valid=x_valid/255
```

```
[25]: # Flattening images  
x_subtrain_flat = x_subtrain.reshape(x_subtrain.shape[0], -1)  
x_valid_flat = x_valid.reshape(x_valid.shape[0], -1)  
x_test_flat = x_test.reshape(x_test.shape[0], -1)
```

#### 0.6.4 Some random images from the dataset along with their labels

```
[26]: import numpy as np  
  
# Mapping of label numbers to class names  
label_names = {  
    0: "airplane",  
    1: "automobile",  
    2: "bird",  
    3: "cat",  
    4: "deer",  
    5: "dog",  
    6: "frog",  
    7: "horse",  
    8: "ship",  
    9: "truck"  
}  
  
# Picking 5 random indexes  
random_indices = np.random.choice(len(x_train), size=5, replace=False)  
  
# Plotting the images with labels  
plt.figure(figsize=(2, 10))  
for i, idx in enumerate(random_indices):  
    plt.subplot(5, 1, i+1)  
    plt.imshow(x_train[idx])  
    plt.title(label_names[int(y_train[idx])])  
    plt.axis('off')  
  
plt.tight_layout()  
plt.show()
```

cat



horse



bird



airplane



frog



### 0.6.5 One hot encoding the target labels

```
[27]: y_subtrain_cat = to_categorical(y_subtrain, 10)
      y_valid_cat = to_categorical(y_valid, 10)
      y_test_cat = to_categorical(y_test, 10)
```

### 0.6.6 Building the model architecture

- without regularization or dropout layer

```
[28]: model = Sequential()
      model.add(Dense(512, activation = 'relu', input_shape = (3072,)))
      model.add(Dense(256, activation = 'relu'))
      model.add(Dense(128, activation = 'relu'))
      model.add(Dense(10, activation = 'softmax'))
      model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 512)	1,573,376
dense_16 (Dense)	(None, 256)	131,328
dense_17 (Dense)	(None, 128)	32,896
dense_18 (Dense)	(None, 10)	1,290

Total params: 1,738,890 (6.63 MB)

Trainable params: 1,738,890 (6.63 MB)

Non-trainable params: 0 (0.00 B)

### 0.6.7 Compiling and running the model

```
[29]: from keras.callbacks import EarlyStopping
```

```

model.compile(loss= 'categorical_crossentropy',optimizer = Adam(learning_rate=0.
↳001),metrics = ['accuracy'])

estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-4, mode = 'min',
↳patience = 5, verbose = 1, restore_best_weights=True)
history = model.fit(x_subtrain_flat,y_subtrain_cat, batch_size=128, epochs =
↳500, verbose = 1, validation_data=(x_valid_flat,y_valid_cat),
↳callbacks=[estop])

```

Epoch 1/500

352/352 9s 22ms/step -

accuracy: 0.2554 - loss: 2.0817 - val\_accuracy: 0.3670 - val\_loss: 1.7583

Epoch 2/500

352/352 8s 23ms/step -

accuracy: 0.3826 - loss: 1.7192 - val\_accuracy: 0.4096 - val\_loss: 1.6545

Epoch 3/500

352/352 7s 20ms/step -

accuracy: 0.4096 - loss: 1.6392 - val\_accuracy: 0.4168 - val\_loss: 1.6265

Epoch 4/500

352/352 7s 20ms/step -

accuracy: 0.4449 - loss: 1.5593 - val\_accuracy: 0.4224 - val\_loss: 1.6146

Epoch 5/500

352/352 7s 21ms/step -

accuracy: 0.4554 - loss: 1.5212 - val\_accuracy: 0.4634 - val\_loss: 1.5215

Epoch 6/500

352/352 7s 20ms/step -

accuracy: 0.4682 - loss: 1.4804 - val\_accuracy: 0.4612 - val\_loss: 1.5085

Epoch 7/500

352/352 7s 20ms/step -

accuracy: 0.4845 - loss: 1.4437 - val\_accuracy: 0.4658 - val\_loss: 1.4851

Epoch 8/500

352/352 7s 21ms/step -

accuracy: 0.4944 - loss: 1.4155 - val\_accuracy: 0.4784 - val\_loss: 1.4744

Epoch 9/500

352/352 9s 25ms/step -

accuracy: 0.5056 - loss: 1.3740 - val\_accuracy: 0.4386 - val\_loss: 1.5614

Epoch 10/500

352/352 8s 22ms/step -

accuracy: 0.5130 - loss: 1.3540 - val\_accuracy: 0.4904 - val\_loss: 1.4252

Epoch 11/500

352/352 7s 20ms/step -

accuracy: 0.5279 - loss: 1.3258 - val\_accuracy: 0.4926 - val\_loss: 1.4139

Epoch 12/500

352/352 7s 21ms/step -

accuracy: 0.5347 - loss: 1.3026 - val\_accuracy: 0.4996 - val\_loss: 1.4141

Epoch 13/500

352/352 7s 21ms/step -



```

accuracy: 0.5466 - loss: 1.2611 - val_accuracy: 0.5052 - val_loss: 1.4028
Epoch 14/500
352/352          9s 25ms/step -
accuracy: 0.5610 - loss: 1.2322 - val_accuracy: 0.5026 - val_loss: 1.3980
Epoch 15/500
352/352          8s 23ms/step -
accuracy: 0.5694 - loss: 1.2097 - val_accuracy: 0.5058 - val_loss: 1.4046
Epoch 16/500
352/352          7s 21ms/step -
accuracy: 0.5741 - loss: 1.1996 - val_accuracy: 0.5158 - val_loss: 1.3886
Epoch 17/500
352/352          8s 23ms/step -
accuracy: 0.5836 - loss: 1.1684 - val_accuracy: 0.5126 - val_loss: 1.3907
Epoch 18/500
352/352          8s 23ms/step -
accuracy: 0.5900 - loss: 1.1437 - val_accuracy: 0.5034 - val_loss: 1.4234
Epoch 19/500
352/352          7s 21ms/step -
accuracy: 0.5990 - loss: 1.1282 - val_accuracy: 0.5032 - val_loss: 1.4273
Epoch 20/500
352/352          7s 20ms/step -
accuracy: 0.6108 - loss: 1.0971 - val_accuracy: 0.5144 - val_loss: 1.4316
Epoch 21/500
352/352          7s 19ms/step -
accuracy: 0.6178 - loss: 1.0706 - val_accuracy: 0.5052 - val_loss: 1.4354
Epoch 21: early stopping
Restoring model weights from the end of the best epoch: 16.

```

```

[30]: score = model.evaluate(x_test_flat,y_test_cat, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          1s 3ms/step -
accuracy: 0.5154 - loss: 1.3739
Test loss: 1.3852113485336304
Test Accuracy:51.97%

```

```

[31]: plt.figure(figsize=(12,4))


plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')

```

```
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



### 0.6.8 Model with dropout layers

```
[32]: model = Sequential()
model.add(Dense(512, activation = 'relu', input_shape = (3072,)))
model.add(Dropout(0.3))
model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.3))
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation = 'softmax'))
model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 512)	1,573,376
dropout_2 (Dropout)	(None, 512)	0
dense_20 (Dense)	(None, 256)	131,328
dropout_3 (Dropout)	(None, 256)	0

dense_21 (Dense)	(None, 128)	32,896
dropout_4 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 10)	1,290

Total params: 1,738,890 (6.63 MB)

Trainable params: 1,738,890 (6.63 MB)

Non-trainable params: 0 (0.00 B)

```
[33]: from keras.callbacks import EarlyStopping

model.compile(loss= 'categorical_crossentropy',optimizer = Adam(learning_rate=0.
↳001),metrics = ['accuracy'])

estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-4, mode = 'min',
↳patience = 5, verbose = 1, restore_best_weights=True)
history = model.fit(x_subtrain_flat,y_subtrain_cat, batch_size=128, epochs =
↳500, verbose = 1, validation_data=(x_valid_flat,y_valid_cat),
↳callbacks=[estop])
```

Epoch 1/500

352/352 10s 25ms/step -

accuracy: 0.1807 - loss: 2.2071 - val\_accuracy: 0.3238 - val\_loss: 1.8934

Epoch 2/500

352/352 8s 23ms/step -

accuracy: 0.2852 - loss: 1.9350 - val\_accuracy: 0.3202 - val\_loss: 1.8745

Epoch 3/500

352/352 7s 21ms/step -

accuracy: 0.3015 - loss: 1.8916 - val\_accuracy: 0.3638 - val\_loss: 1.7975

Epoch 4/500

352/352 7s 20ms/step -

accuracy: 0.3119 - loss: 1.8670 - val\_accuracy: 0.3634 - val\_loss: 1.7696

Epoch 5/500

352/352 8s 23ms/step -

accuracy: 0.3225 - loss: 1.8450 - val\_accuracy: 0.3570 - val\_loss: 1.7812

Epoch 6/500

352/352 9s 25ms/step -

accuracy: 0.3302 - loss: 1.8304 - val\_accuracy: 0.3758 - val\_loss: 1.7633

Epoch 7/500

352/352 10s 28ms/step -

accuracy: 0.3340 - loss: 1.8129 - val\_accuracy: 0.3810 - val\_loss: 1.7565

Epoch 8/500  
352/352 9s 25ms/step -  
accuracy: 0.3451 - loss: 1.8041 - val\_accuracy: 0.3864 - val\_loss: 1.7250

Epoch 9/500  
352/352 9s 26ms/step -  
accuracy: 0.3491 - loss: 1.7873 - val\_accuracy: 0.4022 - val\_loss: 1.7131

Epoch 10/500  
352/352 8s 22ms/step -  
accuracy: 0.3491 - loss: 1.7870 - val\_accuracy: 0.3996 - val\_loss: 1.7129

Epoch 11/500  
352/352 7s 21ms/step -  
accuracy: 0.3579 - loss: 1.7663 - val\_accuracy: 0.3984 - val\_loss: 1.7022

Epoch 12/500  
352/352 7s 21ms/step -  
accuracy: 0.3630 - loss: 1.7551 - val\_accuracy: 0.3994 - val\_loss: 1.6952

Epoch 13/500  
352/352 7s 21ms/step -  
accuracy: 0.3633 - loss: 1.7485 - val\_accuracy: 0.3978 - val\_loss: 1.6983

Epoch 14/500  
352/352 7s 21ms/step -  
accuracy: 0.3591 - loss: 1.7584 - val\_accuracy: 0.4076 - val\_loss: 1.6914

Epoch 15/500  
352/352 8s 21ms/step -  
accuracy: 0.3694 - loss: 1.7258 - val\_accuracy: 0.4110 - val\_loss: 1.6813

Epoch 16/500  
352/352 8s 21ms/step -  
accuracy: 0.3750 - loss: 1.7254 - val\_accuracy: 0.4200 - val\_loss: 1.6518

Epoch 17/500  
352/352 8s 22ms/step -  
accuracy: 0.3788 - loss: 1.7118 - val\_accuracy: 0.4244 - val\_loss: 1.6719

Epoch 18/500  
352/352 7s 20ms/step -  
accuracy: 0.3794 - loss: 1.7138 - val\_accuracy: 0.4250 - val\_loss: 1.6546

Epoch 19/500  
352/352 7s 21ms/step -  
accuracy: 0.3791 - loss: 1.7157 - val\_accuracy: 0.4228 - val\_loss: 1.6645

Epoch 20/500  
352/352 8s 23ms/step -  
accuracy: 0.3806 - loss: 1.7142 - val\_accuracy: 0.4260 - val\_loss: 1.6423

Epoch 21/500  
352/352 7s 19ms/step -  
accuracy: 0.3810 - loss: 1.7043 - val\_accuracy: 0.4218 - val\_loss: 1.6621

Epoch 22/500  
352/352 7s 19ms/step -  
accuracy: 0.3816 - loss: 1.7004 - val\_accuracy: 0.4054 - val\_loss: 1.6683

Epoch 23/500  
352/352 6s 18ms/step -  
accuracy: 0.3839 - loss: 1.7022 - val\_accuracy: 0.4262 - val\_loss: 1.6465

Epoch 24/500  
 352/352 6s 18ms/step -  
 accuracy: 0.3908 - loss: 1.6887 - val\_accuracy: 0.4112 - val\_loss: 1.6733  
 Epoch 25/500  
 352/352 7s 19ms/step -  
 accuracy: 0.3828 - loss: 1.6960 - val\_accuracy: 0.4102 - val\_loss: 1.6692  
 Epoch 25: early stopping  
 Restoring model weights from the end of the best epoch: 20.

```
[34]: score = model.evaluate(x_test_flat,y_test_cat, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")
```


313/313 1s 3ms/step -  
 accuracy: 0.4278 - loss: 1.6321  
 Test loss: 1.6369125843048096  
 Test Accuracy:42.79%

```
[35]: plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



## 0.7 Question 3. Next from keras.regularizers import l2

### 0.7.1 Model building with l2 regularizer

```
[36]: model = Sequential()
model.add(Dense(512, activation = 'relu', kernel_regularizer=l2(0.
    ↳0001),input_shape = (3072,)))
model.add(Dense(256, activation = 'relu',kernel_regularizer=l2(0.0001)))
model.add(Dense(128, activation = 'relu',kernel_regularizer=l2(0.005)))
model.add(Dense(10, activation = 'softmax'))
model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 512)	1,573,376
dense_24 (Dense)	(None, 256)	131,328
dense_25 (Dense)	(None, 128)	32,896
dense_26 (Dense)	(None, 10)	1,290

Total params: 1,738,890 (6.63 MB)

Trainable params: 1,738,890 (6.63 MB)

Non-trainable params: 0 (0.00 B)

### 0.7.2 Without using early stopping

```
[37]: model.compile(loss= 'categorical_crossentropy',optimizer = Adam(learning_rate=0.
    ↳001),metrics = ['accuracy'])

estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-4, mode = 'min',
    ↳patience = 5, verbose = 1, restore_best_weights=True)
history = model.fit(x_subtrain_flat,y_subtrain_cat, batch_size=128, epochs =
    ↳500, verbose = 1, validation_data=(x_valid_flat,y_valid_cat),
    ↳callbacks=[estop])
```

Epoch 1/500  
352/352 10s 23ms/step -  
accuracy: 0.2553 - loss: 2.6025 - val\_accuracy: 0.3716 - val\_loss: 1.9682

Epoch 2/500  
352/352 7s 20ms/step -  
accuracy: 0.3771 - loss: 1.9061 - val\_accuracy: 0.4190 - val\_loss: 1.7790

Epoch 3/500  
352/352 7s 20ms/step -  
accuracy: 0.4159 - loss: 1.7583 - val\_accuracy: 0.4314 - val\_loss: 1.7220

Epoch 4/500  
352/352 7s 20ms/step -  
accuracy: 0.4350 - loss: 1.6784 - val\_accuracy: 0.4676 - val\_loss: 1.6170

Epoch 5/500  
352/352 7s 20ms/step -  
accuracy: 0.4565 - loss: 1.6113 - val\_accuracy: 0.4622 - val\_loss: 1.6180

Epoch 6/500  
352/352 7s 20ms/step -  
accuracy: 0.4679 - loss: 1.5756 - val\_accuracy: 0.4544 - val\_loss: 1.6137

Epoch 7/500  
352/352 7s 20ms/step -  
accuracy: 0.4787 - loss: 1.5499 - val\_accuracy: 0.4600 - val\_loss: 1.5985

Epoch 8/500  
352/352 7s 21ms/step -  
accuracy: 0.4832 - loss: 1.5293 - val\_accuracy: 0.4656 - val\_loss: 1.5763

Epoch 9/500  
352/352 7s 20ms/step -  
accuracy: 0.4918 - loss: 1.5201 - val\_accuracy: 0.4738 - val\_loss: 1.5532

Epoch 10/500  
352/352 7s 20ms/step -  
accuracy: 0.5001 - loss: 1.4908 - val\_accuracy: 0.4738 - val\_loss: 1.5866

Epoch 11/500  
352/352 7s 20ms/step -  
accuracy: 0.5093 - loss: 1.4653 - val\_accuracy: 0.4852 - val\_loss: 1.5406

Epoch 12/500  
352/352 7s 20ms/step -  
accuracy: 0.5160 - loss: 1.4462 - val\_accuracy: 0.4908 - val\_loss: 1.5158

Epoch 13/500  
352/352 7s 20ms/step -  
accuracy: 0.5194 - loss: 1.4355 - val\_accuracy: 0.4954 - val\_loss: 1.5275

Epoch 14/500  
352/352 7s 21ms/step -  
accuracy: 0.5283 - loss: 1.4164 - val\_accuracy: 0.5062 - val\_loss: 1.4899

Epoch 15/500  
352/352 7s 20ms/step -  
accuracy: 0.5300 - loss: 1.4028 - val\_accuracy: 0.4844 - val\_loss: 1.5423

Epoch 16/500  
352/352 7s 20ms/step -  
accuracy: 0.5437 - loss: 1.3788 - val\_accuracy: 0.5064 - val\_loss: 1.4921

```

Epoch 17/500
352/352          7s 20ms/step -
accuracy: 0.5483 - loss: 1.3751 - val_accuracy: 0.5058 - val_loss: 1.4971
Epoch 18/500
352/352          7s 20ms/step -
accuracy: 0.5472 - loss: 1.3703 - val_accuracy: 0.5154 - val_loss: 1.4749
Epoch 19/500
352/352          8s 21ms/step -
accuracy: 0.5584 - loss: 1.3410 - val_accuracy: 0.4946 - val_loss: 1.5050
Epoch 20/500
352/352          7s 21ms/step -
accuracy: 0.5597 - loss: 1.3463 - val_accuracy: 0.5168 - val_loss: 1.4709
Epoch 21/500
352/352          7s 20ms/step -
accuracy: 0.5629 - loss: 1.3367 - val_accuracy: 0.5042 - val_loss: 1.4976
Epoch 22/500
352/352          7s 20ms/step -
accuracy: 0.5648 - loss: 1.3345 - val_accuracy: 0.5042 - val_loss: 1.4977
Epoch 23/500
352/352          7s 21ms/step -
accuracy: 0.5659 - loss: 1.3266 - val_accuracy: 0.5152 - val_loss: 1.4802
Epoch 24/500
352/352          7s 21ms/step -
accuracy: 0.5714 - loss: 1.3183 - val_accuracy: 0.5066 - val_loss: 1.5034
Epoch 25/500
352/352          7s 20ms/step -
accuracy: 0.5798 - loss: 1.2953 - val_accuracy: 0.5174 - val_loss: 1.5022
Epoch 25: early stopping
Restoring model weights from the end of the best epoch: 20.

```

```

[38]: score = model.evaluate(x_test_flat,y_test_cat, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          1s 4ms/step -
accuracy: 0.5204 - loss: 1.4501
Test loss: 1.4567960500717163
Test Accuracy:51.61%

```

```

[39]: plt.figure(figsize=(12,4))


plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

```



```
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



### 0.7.3 using early stopping along with l2 regularization

```
[40]: estop = EarlyStopping(monitor = 'val_loss', min_delta = 1e-4, mode = 'min',
    ↪patience = 5, verbose = 1, restore_best_weights=True)
history = model.fit(x_subtrain_flat,y_subtrain_cat, batch_size=128, epochs =
    ↪500, verbose = 1, validation_data=(x_valid_flat,y_valid_cat),
    ↪callbacks=[estop])
```

```
Epoch 1/500
352/352          9s 24ms/step -
accuracy: 0.5652 - loss: 1.3301 - val_accuracy: 0.4984 - val_loss: 1.5603
Epoch 2/500
352/352          7s 21ms/step -
accuracy: 0.5633 - loss: 1.3336 - val_accuracy: 0.5030 - val_loss: 1.5436
Epoch 3/500
352/352          7s 20ms/step -
accuracy: 0.5672 - loss: 1.3255 - val_accuracy: 0.5118 - val_loss: 1.4963
Epoch 4/500
352/352          7s 21ms/step -
accuracy: 0.5742 - loss: 1.3156 - val_accuracy: 0.4996 - val_loss: 1.5386
Epoch 5/500
352/352          7s 20ms/step -
accuracy: 0.5724 - loss: 1.3093 - val_accuracy: 0.5086 - val_loss: 1.5058
Epoch 6/500
```

```

352/352          8s 21ms/step -
accuracy: 0.5819 - loss: 1.2856 - val_accuracy: 0.5030 - val_loss: 1.5206
Epoch 7/500
352/352          7s 20ms/step -
accuracy: 0.5764 - loss: 1.2972 - val_accuracy: 0.5168 - val_loss: 1.5215
Epoch 8/500
352/352          7s 20ms/step -
accuracy: 0.5907 - loss: 1.2719 - val_accuracy: 0.5230 - val_loss: 1.4814
Epoch 9/500
352/352          7s 20ms/step -
accuracy: 0.5892 - loss: 1.2758 - val_accuracy: 0.5266 - val_loss: 1.4918
Epoch 10/500
352/352          7s 21ms/step -
accuracy: 0.5846 - loss: 1.2767 - val_accuracy: 0.5162 - val_loss: 1.4942
Epoch 11/500
352/352          7s 20ms/step -
accuracy: 0.5923 - loss: 1.2511 - val_accuracy: 0.5014 - val_loss: 1.5157
Epoch 12/500
352/352          7s 20ms/step -
accuracy: 0.5935 - loss: 1.2714 - val_accuracy: 0.5276 - val_loss: 1.5037
Epoch 13/500
352/352          7s 20ms/step -
accuracy: 0.5961 - loss: 1.2632 - val_accuracy: 0.5250 - val_loss: 1.4913
Epoch 13: early stopping
Restoring model weights from the end of the best epoch: 8.

```

```

[41]: score = model.evaluate(x_test_flat,y_test_cat, verbose = 1)
print("Test loss:", score[0])
print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          1s 3ms/step -
accuracy: 0.5239 - loss: 1.4569
Test loss: 1.465133547782898
Test Accuracy:52.07%

```

```

[42]: plt.figure(figsize=(12,4))


plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')

```

```
plt.ylabel('Loss')
plt.legend()

plt.show()
```



**0.8 Question 4:** Now, let's see how we can proceed to do perform some hyperparameter tuning and find out the appropriate parameter value. The following part is done for a very simple model with one hidden layer and an output layer. The number of neurons and the dropout parameter is being tuned to find appropriate ones.

```
[43]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.utils import to_categorical
from keras.optimizers import SGD, Adam
import keras_tuner as kt
```

```
[44]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
print(x_train.shape)

x_train = x_train.reshape(-1, 28*28).astype('float32')/255.0
x_test = x_test.reshape(-1, 28*28).astype('float32')/255.0

print(x_train.shape)
print(x_test.shape)
```

```
(60000, 28, 28)
(60000, 784)
(10000, 784)
```

```
[45]: y_train_ = to_categorical(y_train,10)
y_test_ = to_categorical(y_test,10)

def build_model(hp):
    model = Sequential()
    model.add(Flatten(input_shape= (28*28,)))

    units = hp.Int('units', min_value = 64, max_value = 512, step = 64)
    model.add(Dense(units,activation = 'relu'))
    dropout_rate = hp.Float('dropout',min_value = 0.0, max_value = 0.5, step = 0.1)
    model.add(Dropout(dropout_rate))
    model.add(Dense(10, activation = 'softmax'))

    model.compile(
        optimizer = SGD(),
        loss = 'categorical_crossentropy',
        metrics = ['accuracy']
    )
    return model
```

```
[46]: tuner = kt.RandomSearch(
    build_model,
    objective = 'val_accuracy',
    max_trials = 10,
    executions_per_trial = 1,
    directory = 'mnist_tuning',
    project_name = 'dense_dropout_tune',
    overwrite=True
)

tuner.search(x_train,y_train_, epochs = 10, validation_split = 0.2,batch_size = 128,
callbacks = [keras.callbacks.EarlyStopping(monitor = 'val_loss',patience = 5)])
best_model = tuner.get_best_models(num_models = 1)[0]

test_loss, test_acc = best_model.evaluate(x_test,y_test_)
print("Test Accuracy:",test_acc)

best_hps = tuner.get_best_hyperparameters(1)[0]
print("Best Units:", best_hps.get('units'))
print("Best dropout:",best_hps.get('dropout'))
```

Trial 10 Complete [00h 00m 15s]  
val\_accuracy: 0.9211666584014893

Best val\_accuracy So Far: 0.9235833287239075  
Total elapsed time: 00h 02m 53s

```

313/313          1s 2ms/step -
accuracy: 0.9105 - loss: 0.3241
Test Accuracy: 0.9232000112533569
Best Units: 320
Best dropout: 0.2

```

```
[47]: best_model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 320)	251,200
dropout (Dropout)	(None, 320)	0
dense_1 (Dense)	(None, 10)	3,210

```
Total params: 254,410 (993.79 KB)
```

```
Trainable params: 254,410 (993.79 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[48]: history = best_model.fit(x_train,y_train_, batch_size=128, epochs = 50, verbose=
    ↪ 1, validation_split=0.2,callbacks = [keras.callbacks.EarlyStopping(monitor=
    ↪ 'val_loss', patience = 5)])
```

```
Epoch 1/50
```

```
375/375          2s 4ms/step -
accuracy: 0.9090 - loss: 0.3180 - val_accuracy: 0.9255 - val_loss: 0.2722
```

```
Epoch 2/50
```

```
375/375          2s 4ms/step -
accuracy: 0.9152 - loss: 0.3035 - val_accuracy: 0.9270 - val_loss: 0.2646
```

```
Epoch 3/50
```

```
375/375          2s 5ms/step -
accuracy: 0.9161 - loss: 0.2941 - val_accuracy: 0.9298 - val_loss: 0.2577
```

```
Epoch 4/50
```

```
375/375          2s 4ms/step -
accuracy: 0.9172 - loss: 0.2921 - val_accuracy: 0.9303 - val_loss: 0.2517
```

```
Epoch 5/50
```

```
375/375          2s 5ms/step -
```

accuracy: 0.9197 - loss: 0.2824 - val\_accuracy: 0.9316 - val\_loss: 0.2458  
 Epoch 6/50  
 375/375                    2s 5ms/step -  
 accuracy: 0.9223 - loss: 0.2770 - val\_accuracy: 0.9337 - val\_loss: 0.2405  
 Epoch 7/50  
 375/375                    3s 7ms/step -  
 accuracy: 0.9268 - loss: 0.2656 - val\_accuracy: 0.9352 - val\_loss: 0.2350  
 Epoch 8/50  
 375/375                    2s 5ms/step -  
 accuracy: 0.9239 - loss: 0.2645 - val\_accuracy: 0.9362 - val\_loss: 0.2305  
 Epoch 9/50  
 375/375                    2s 4ms/step -  
 accuracy: 0.9281 - loss: 0.2586 - val\_accuracy: 0.9377 - val\_loss: 0.2257  
 Epoch 10/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9309 - loss: 0.2508 - val\_accuracy: 0.9388 - val\_loss: 0.2217  
 Epoch 11/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9322 - loss: 0.2432 - val\_accuracy: 0.9400 - val\_loss: 0.2173  
 Epoch 12/50  
 375/375                    2s 4ms/step -  
 accuracy: 0.9310 - loss: 0.2428 - val\_accuracy: 0.9411 - val\_loss: 0.2133  
 Epoch 13/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9351 - loss: 0.2337 - val\_accuracy: 0.9423 - val\_loss: 0.2098  
 Epoch 14/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9338 - loss: 0.2364 - val\_accuracy: 0.9433 - val\_loss: 0.2059  
 Epoch 15/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9344 - loss: 0.2321 - val\_accuracy: 0.9444 - val\_loss: 0.2024  
 Epoch 16/50  
 375/375                    2s 5ms/step -  
 accuracy: 0.9367 - loss: 0.2207 - val\_accuracy: 0.9447 - val\_loss: 0.1991  
 Epoch 17/50  
 375/375                    1s 4ms/step -  
 accuracy: 0.9358 - loss: 0.2241 - val\_accuracy: 0.9456 - val\_loss: 0.1965  
 Epoch 18/50  
 375/375                    2s 4ms/step -  
 accuracy: 0.9386 - loss: 0.2182 - val\_accuracy: 0.9463 - val\_loss: 0.1933  
 Epoch 19/50  
 375/375                    2s 4ms/step -  
 accuracy: 0.9408 - loss: 0.2071 - val\_accuracy: 0.9473 - val\_loss: 0.1903  
 Epoch 20/50  
 375/375                    2s 4ms/step -  
 accuracy: 0.9416 - loss: 0.2077 - val\_accuracy: 0.9477 - val\_loss: 0.1876  
 Epoch 21/50  
 375/375                    2s 4ms/step -

accuracy: 0.9411 - loss: 0.2061 - val\_accuracy: 0.9483 - val\_loss: 0.1847  
 Epoch 22/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9429 - loss: 0.2023 - val\_accuracy: 0.9499 - val\_loss: 0.1823  
 Epoch 23/50  
 375/375 2s 4ms/step -  
 accuracy: 0.9432 - loss: 0.2012 - val\_accuracy: 0.9496 - val\_loss: 0.1801  
 Epoch 24/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9447 - loss: 0.1935 - val\_accuracy: 0.9504 - val\_loss: 0.1776  
 Epoch 25/50  
 375/375 2s 6ms/step -  
 accuracy: 0.9476 - loss: 0.1912 - val\_accuracy: 0.9525 - val\_loss: 0.1749  
 Epoch 26/50  
 375/375 3s 8ms/step -  
 accuracy: 0.9465 - loss: 0.1877 - val\_accuracy: 0.9523 - val\_loss: 0.1728  
 Epoch 27/50  
 375/375 3s 8ms/step -  
 accuracy: 0.9474 - loss: 0.1862 - val\_accuracy: 0.9531 - val\_loss: 0.1706  
 Epoch 28/50  
 375/375 3s 7ms/step -  
 accuracy: 0.9468 - loss: 0.1876 - val\_accuracy: 0.9534 - val\_loss: 0.1686  
 Epoch 29/50  
 375/375 2s 6ms/step -  
 accuracy: 0.9485 - loss: 0.1836 - val\_accuracy: 0.9541 - val\_loss: 0.1666  
 Epoch 30/50  
 375/375 2s 6ms/step -  
 accuracy: 0.9483 - loss: 0.1779 - val\_accuracy: 0.9544 - val\_loss: 0.1647  
 Epoch 31/50  
 375/375 3s 7ms/step -  
 accuracy: 0.9482 - loss: 0.1792 - val\_accuracy: 0.9548 - val\_loss: 0.1627  
 Epoch 32/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9515 - loss: 0.1737 - val\_accuracy: 0.9553 - val\_loss: 0.1609  
 Epoch 33/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9491 - loss: 0.1751 - val\_accuracy: 0.9557 - val\_loss: 0.1591  
 Epoch 34/50  
 375/375 2s 4ms/step -  
 accuracy: 0.9526 - loss: 0.1663 - val\_accuracy: 0.9565 - val\_loss: 0.1576  
 Epoch 35/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9514 - loss: 0.1672 - val\_accuracy: 0.9573 - val\_loss: 0.1557  
 Epoch 36/50  
 375/375 2s 5ms/step -  
 accuracy: 0.9518 - loss: 0.1689 - val\_accuracy: 0.9575 - val\_loss: 0.1543  
 Epoch 37/50  
 375/375 2s 5ms/step -

```

accuracy: 0.9562 - loss: 0.1581 - val_accuracy: 0.9582 - val_loss: 0.1526
Epoch 38/50
375/375          2s 5ms/step -
accuracy: 0.9558 - loss: 0.1584 - val_accuracy: 0.9587 - val_loss: 0.1512
Epoch 39/50
375/375          2s 4ms/step -
accuracy: 0.9537 - loss: 0.1617 - val_accuracy: 0.9587 - val_loss: 0.1498
Epoch 40/50
375/375          2s 5ms/step -
accuracy: 0.9546 - loss: 0.1581 - val_accuracy: 0.9592 - val_loss: 0.1482
Epoch 41/50
375/375          2s 4ms/step -
accuracy: 0.9554 - loss: 0.1565 - val_accuracy: 0.9603 - val_loss: 0.1468
Epoch 42/50
375/375          2s 4ms/step -
accuracy: 0.9557 - loss: 0.1566 - val_accuracy: 0.9601 - val_loss: 0.1455
Epoch 43/50
375/375          2s 4ms/step -
accuracy: 0.9568 - loss: 0.1529 - val_accuracy: 0.9599 - val_loss: 0.1441
Epoch 44/50
375/375          2s 4ms/step -
accuracy: 0.9573 - loss: 0.1486 - val_accuracy: 0.9609 - val_loss: 0.1430
Epoch 45/50
375/375          2s 5ms/step -
accuracy: 0.9573 - loss: 0.1520 - val_accuracy: 0.9606 - val_loss: 0.1419
Epoch 46/50
375/375          2s 4ms/step -
accuracy: 0.9574 - loss: 0.1528 - val_accuracy: 0.9610 - val_loss: 0.1405
Epoch 47/50
375/375          2s 5ms/step -
accuracy: 0.9596 - loss: 0.1436 - val_accuracy: 0.9613 - val_loss: 0.1396
Epoch 48/50
375/375          2s 4ms/step -
accuracy: 0.9585 - loss: 0.1478 - val_accuracy: 0.9613 - val_loss: 0.1382
Epoch 49/50
375/375          2s 4ms/step -
accuracy: 0.9593 - loss: 0.1450 - val_accuracy: 0.9617 - val_loss: 0.1371
Epoch 50/50
375/375          2s 6ms/step -
accuracy: 0.9593 - loss: 0.1439 - val_accuracy: 0.9618 - val_loss: 0.1360

```

```

[49]: score = best_model.evaluate(x_test,y_test_, verbose = 1)
      print("Test loss:", score[0])
      print(f"Test Accuracy:{score[1]*100:.2f}%")

```

```

313/313          2s 6ms/step -
accuracy: 0.9556 - loss: 0.1555
Test loss: 0.13252055644989014

```




Test Accuracy:96.27%

```
[50]: plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



### 0.8.1 Tuning further parameters like

- Number of hidden layers
- units per layers
- dropout per layer
- optimizer type
- learning rate

```
[51]: from tensorflow import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout
from keras.utils import to_categorical
```

```
import keras_tuner as kt
from keras.regularizers import l2
import numpy as np
```

## 0.8.2 Loading and preprocessing the dataset


```
[52]: (x_train,y_train),(x_test,y_test) = mnist.load_data()

## Plotting some random images from the dataset
random_indices = np.random.choice(len(x_train), size=5, replace=False)
plt.figure(figsize=(2, 10))
for i, idx in enumerate(random_indices):
    plt.subplot(5, 1, i+1)
    plt.imshow(x_train[idx],cmap = 'gray')
    plt.title(int(y_train[idx]))
    plt.axis('off')

plt.tight_layout()
plt.show()

x_train = x_train.reshape(-1,28*28).astype('float32')/255.0
x_test = x_test.reshape(-1,28*28).astype('float32')/255.0
y_train_ = to_categorical(y_train,10)
y_test_ = to_categorical(y_test, 10)
```

4



9



2



4



5



### 0.8.3 Building model for tuner

- with L2 regularization

```
[53]: def build_model(hp):
    model = Sequential()
    model.add(Flatten(input_shape = (784,)))

    ## Tuning the number of layers
    for i in range(hp.Int('num_layers',1,3)):
        model.add(Dense(
            ## no of nodes in each layer
            units = hp.Int(f'units_{i}', min_value = 64, max_value = 512, step_u
↪= 64),
            activation = 'relu',
            kernel_regularizer=l2(
                hp.Choice(f'l2_{i}', values = [0.0,1e-4,1e-3,1e-2])
            )
        ))
        ## Tuning the dropout rate
        model.add(Dropout(
            rate = hp.Float(f'dropout_{i}', min_value = 0.0,max_value = 0.5,u
↪step = 0.05)
        ))

    ## output layer
    model.add(Dense(10, activation = 'softmax'))

    ## Tuning the optimizer type and the learning rate
    optimizer_choice = hp.Choice('optimizer', values = ['adam','sgd'])
    learning_rate = hp.Choice('learning_rate', values=[1e-2,1e-3,1e-4])

    if optimizer_choice == 'adam':
        optimizer = keras.optimizers.Adam(learning_rate=learning_rate)
    else:
        optimizer = keras.optimizers.SGD(learning_rate=learning_rate)

    model.compile(
        optimizer = optimizer,
        loss = 'categorical_crossentropy',
        metrics = ['accuracy']
    )
    return model
```

```
[54]: tuner = kt.RandomSearch(
    build_model,
    objective = 'val_accuracy',
    max_trials=20,
    executions_per_trial = 2,
    overwrite = True,
    directory = 'mnist_tuning',
    project_name = 'advanced_dense_tune'
)

tuner.search(
    x_train,y_train_,
    epochs = 10,
    validation_split = 0.2,
    batch_size = 128,
    callbacks = [keras.callbacks.EarlyStopping(monitor = 'val_loss', patience = 3, min_delta=1e-4, restore_best_weights = True)]
)
```

Trial 20 Complete [00h 00m 30s]  
val\_accuracy: 0.28724999725818634

Best val\_accuracy So Far: 0.9781250059604645  
Total elapsed time: 00h 16m 35s

```
[55]: # best model and hyperparameters
best_model = tuner.get_best_models(num_models=1)[0]
best_hps = tuner.get_best_hyperparameters(1)[0]
print("Best Hyperparameters:")
for key in best_hps.values.keys():
    print(f"{key}: {best_hps.get(key)}")

# Final training
history = best_model.fit(
    x_train, y_train_,
    batch_size=128,
    epochs=50,
    validation_split=0.2,
    callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience = 3, min_delta=1e-4, restore_best_weights = True)]
)

# Evaluate
score = best_model.evaluate(x_test, y_test_, verbose=1)
print("Test loss:", score[0])
print(f"Test Accuracy: {score[1]*100:.2f}%")
```

Best Hyperparameters:

```

num_layers: 3
units_0: 256
l2_0: 0.0
dropout_0: 0.30000000000000004
optimizer: adam
learning_rate: 0.001
units_1: 512
l2_1: 0.0
dropout_1: 0.1
units_2: 448
l2_2: 0.0001
dropout_2: 0.0
Epoch 1/50
375/375          5s 10ms/step -
accuracy: 0.9851 - loss: 0.0591 - val_accuracy: 0.9788 - val_loss: 0.0916
Epoch 2/50
375/375          3s 9ms/step -
accuracy: 0.9868 - loss: 0.0522 - val_accuracy: 0.9793 - val_loss: 0.0876
Epoch 3/50
375/375          3s 9ms/step -
accuracy: 0.9856 - loss: 0.0532 - val_accuracy: 0.9787 - val_loss: 0.0891
Epoch 4/50
375/375          4s 10ms/step -
accuracy: 0.9869 - loss: 0.0471 - val_accuracy: 0.9801 - val_loss: 0.0915
Epoch 5/50
375/375          3s 9ms/step -
accuracy: 0.9880 - loss: 0.0465 - val_accuracy: 0.9793 - val_loss: 0.0898
313/313          1s 3ms/step -
accuracy: 0.9778 - loss: 0.0878
Test loss: 0.07718071341514587
Test Accuracy: 98.14%

```

```


[56]: plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```
plt.show()
```



0.9 Then later go back to the cifar10 dataset problem and come up with your best model

0.9.1 hyper parameter tuning for cifar 10 dataset

0.9.2 importing the necessary libraries

```
[57]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Input
from keras.regularizers import l2
from keras.utils import to_categorical
import keras_tuner as kt
```

0.9.3 loading the cifar10 dataset

```
[58]: (x_train,y_train),(x_test,y_test) = cifar10.load_data()

# class label mapping
label_map = {
    0: 'airplane', 1: 'automobile', 2: 'bird', 3: 'cat', 4: 'deer',
    5: 'dog', 6: 'frog', 7: 'horse', 8: 'ship', 9: 'truck'
}
```

#### 0.9.4 Showing some sample images

```
[59]: idx = np.random.choice(len(x_train), 5, replace = False)
plt.figure(figsize = (2,10))
for i, id in enumerate(idx):
    plt.subplot(5,1,i+1)
    plt.imshow(x_train[id])
    plt.title(label_map[int(y_train[id])])
    plt.axis('off')
plt.show()
```



dog




cat



truck



truck



horse



### 0.9.5 preprocessing the dataset

```
[60]: ## normalizing
x_train = x_train.astype('float32')/255.0
x_test = x_test.astype('float32')/255.0
x_train = x_train.reshape(-1,32*32*3)
x_test = x_test.reshape(-1,32*32*3)

y_train_ = to_categorical(y_train,10)
y_test_ = to_categorical(y_test,10)
```

### 0.9.6 model building for hyperparameter tuning

```
[61]: def build_model(hp):
    model = Sequential()
    model.add(Input(shape=(32*32*3,)))
    # Tune number of hidden layers (1-3)
    for i in range(hp.Int('num_layers', 1, 3)):
        model.add(Dense(
            units=hp.Int(f'units_{i}', min_value=128, max_value=512, step=64),
            activation='relu',
            kernel_regularizer=l2(hp.Choice(f'l2_{i}', values=[0.0, 1e-4, ↵
↵1e-3])))
        ))
        model.add(Dropout(
            rate=hp.Float(f'dropout_{i}', min_value=0.0, max_value=0.5, step=0.
↵1)
        ))

    # Output layer
    model.add(Dense(10, activation='softmax'))

    # Tune optimizer type & learning rate
    optimizer_choice = hp.Choice('optimizer', values=['adam', 'sgd'])
    learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

    if optimizer_choice == 'adam':
        optimizer = keras.optimizers.Adam(learning_rate=learning_rate)
    else:
        optimizer = keras.optimizers.SGD(learning_rate=learning_rate, ↵
↵momentum=0.9)

    model.compile(
        optimizer=optimizer,
        loss='categorical_crossentropy',
        metrics=['accuracy']
    )
```

```
return model
```

```
[62]: tuner = kt.RandomSearch(
        build_model,
        objective='val_accuracy',
        max_trials=10,
        executions_per_trial=1,
        overwrite=True,
        directory='cifar10_tuning',
        project_name='dense_dropout_l2_tune'
    )

    tuner.search(
        x_train, y_train_,
        epochs=20,
        validation_split=0.2,
        batch_size=128,
        callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience = 3,
        ↪min_delta=1e-4, restore_best_weights = True)]
    )
```

Trial 10 Complete [00h 01m 22s]  
val\_accuracy: 0.20059999823570251

Best val\_accuracy So Far: 0.5169000029563904  
Total elapsed time: 00h 12m 32s

### 0.9.7 Best model and hyperparameters

```
[63]: best_model = tuner.get_best_models(num_models=1)[0]
        best_hps = tuner.get_best_hyperparameters(1)[0]
        print("\nBest Hyperparameters:")
        for key in best_hps.values.keys():
            print(f"{key}: {best_hps.get(key)}")
```

Best Hyperparameters:  
num\_layers: 3  
units\_0: 384  
l2\_0: 0.0001  
dropout\_0: 0.1  
optimizer: sgd  
learning\_rate: 0.01  
units\_1: 320  
l2\_1: 0.001  
dropout\_1: 0.1  
units\_2: 128  
l2\_2: 0.0

dropout\_2: 0.0

### 0.9.8 Training the best model

```
[64]: history = best_model.fit(  
    x_train, y_train_,  
    batch_size=128,  
    epochs=50,  
    validation_split=0.2,  
    callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience =  
↪3,min_delta=1e-4, restore_best_weights = True)]  
)
```

Epoch 1/50

313/313 8s 20ms/step -

accuracy: 0.5608 - loss: 1.3661 - val\_accuracy: 0.4981 - val\_loss: 1.5471

Epoch 2/50

313/313 5s 16ms/step -

accuracy: 0.5558 - loss: 1.3684 - val\_accuracy: 0.5067 - val\_loss: 1.5259

Epoch 3/50

313/313 6s 18ms/step -

accuracy: 0.5577 - loss: 1.3547 - val\_accuracy: 0.5179 - val\_loss: 1.4900

Epoch 4/50

313/313 5s 17ms/step -

accuracy: 0.5693 - loss: 1.3283 - val\_accuracy: 0.5103 - val\_loss: 1.5069

Epoch 5/50

313/313 6s 18ms/step -

accuracy: 0.5712 - loss: 1.3270 - val\_accuracy: 0.5167 - val\_loss: 1.4978

Epoch 6/50

313/313 5s 17ms/step -

accuracy: 0.5786 - loss: 1.3023 - val\_accuracy: 0.5223 - val\_loss: 1.4850

Epoch 7/50

313/313 6s 18ms/step -

accuracy: 0.5784 - loss: 1.3036 - val\_accuracy: 0.5225 - val\_loss: 1.5013

Epoch 8/50

313/313 5s 16ms/step -

accuracy: 0.5848 - loss: 1.2859 - val\_accuracy: 0.5355 - val\_loss: 1.4660

Epoch 9/50

313/313 5s 15ms/step -

accuracy: 0.5829 - loss: 1.2819 - val\_accuracy: 0.5328 - val\_loss: 1.4694

Epoch 10/50

313/313 4s 14ms/step -

accuracy: 0.5939 - loss: 1.2719 - val\_accuracy: 0.5330 - val\_loss: 1.4760

Epoch 11/50

313/313 4s 14ms/step -

accuracy: 0.5964 - loss: 1.2596 - val\_accuracy: 0.5281 - val\_loss: 1.4908

### 0.9.9 Evaluating the model

```
[65]: score = best_model.evaluate(x_test, y_test_, verbose=1)
print("Test loss:", score[0])
print(f"Test Accuracy: {score[1]*100:.2f}%")
```

```
313/313          1s 3ms/step -
accuracy: 0.5398 - loss: 1.4360
Test loss: 1.4406636953353882
Test Accuracy: 53.62%
```

### 0.9.10 Accuracy and loss plot

```
[66]: plt.figure(figsize=(12, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy Over Epochs')

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Loss Over Epochs')

plt.show()
```

