

DA_3

September 15, 2025

Name:	Tufan Kundu
Registration no:	24MDT0184
Course Name:	Deep Learning Lab
Course Code:	PMDS603P
Digital Assessment:	3

0.1 Question1. Use the MNIST dataset and do necessary pre-processing, and split the data into training, validation, and testing sets. Create a new ANN model with appropriate hidden layers and output layer neurons. Choose appropriate activation functions. Choose the error function appropriately and use SGD as the optimizer. Include early stopping technique in your model and run the model for 500 epochs and report the Performance.

0.1.1 Importing the necessary libraries

```
[38]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout
from keras.optimizers import SGD
from keras.callbacks import EarlyStopping
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import numpy as np

import warnings
warnings.filterwarnings('ignore')
```


0.1.2 Loading the MNIST Dataset

```
[39]: from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
```

0.1.3 Displaying Random Images

```
[40]: class_names = [str(i) for i in range(10)]
indices = np.random.choice(len(x_train), size = 5 , replace = False)

plt.figure(figsize = (2*5,3))
for i, idx in enumerate(indices):
    ax = plt.subplot(1,5,i+1)
    img = x_train[idx]
    plt.imshow(img, cmap = 'gray')
    label = class_names[y_train[idx]]
    plt.title(label)
    plt.axis('off')
plt.show()
```



```
[41]: x_train, x_val, y_train, y_val = train_test_split(
    x_train, y_train,
    test_size=0.2, random_state=42, stratify=y_train
)

print(f"X_train shape: {x_train.shape}")
print(f"X_valid shape: {x_val.shape}")
print(f"X_test shape: {x_test.shape}")

print(f"y_train shape: {y_train.shape}")
print(f"y_valid shape: {y_val.shape}")
print(f"y_test shape: {y_test.shape}")
```

```
X_train shape: (48000, 28, 28)
X_valid shape: (12000, 28, 28)
X_test shape: (10000, 28, 28)
y_train shape: (48000,)
y_valid shape: (12000,)
y_test shape: (10000,)
```

0.1.4 Building the ANN Model

```
[42]: model = Sequential()
model.add(Flatten(input_shape = (28,28)))
model.add(Dense(512, activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.3))
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation = 'softmax'))
model.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
flatten_9 (Flatten)	(None, 784)	0
dense_36 (Dense)	(None, 512)	401,920
dropout_27 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 256)	131,328
dropout_28 (Dropout)	(None, 256)	0
dense_38 (Dense)	(None, 128)	32,896
dropout_29 (Dropout)	(None, 128)	0
dense_39 (Dense)	(None, 10)	1,290

Total params: 567,434 (2.16 MB)

Trainable params: 567,434 (2.16 MB)

Non-trainable params: 0 (0.00 B)

```
[43]: sgd = SGD(learning_rate= 0.01, momentum=0.9)
model.compile(optimizer = sgd, loss = 'sparse_categorical_crossentropy',
metrics = ['accuracy'])
```

```
[44]: estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
    ↪ verbose = 1, restore_best_weights=True)
    history = model.fit(x_train,y_train,batch_size=128, epochs = 500, verbose = 1,
    ↪ validation_data=(x_val,y_val), callbacks=[estop])
```

Epoch 1/500

375/375 3s 6ms/step -

accuracy: 0.5860 - loss: 1.2496 - val_accuracy: 0.9233 - val_loss: 0.2554

Epoch 2/500

375/375 2s 6ms/step -

accuracy: 0.8955 - loss: 0.3440 - val_accuracy: 0.9447 - val_loss: 0.1780

Epoch 3/500

375/375 2s 6ms/step -

accuracy: 0.9253 - loss: 0.2481 - val_accuracy: 0.9555 - val_loss: 0.1487

Epoch 4/500

375/375 2s 6ms/step -

accuracy: 0.9389 - loss: 0.2023 - val_accuracy: 0.9625 - val_loss: 0.1249

Epoch 5/500

375/375 2s 6ms/step -

accuracy: 0.9488 - loss: 0.1700 - val_accuracy: 0.9670 - val_loss: 0.1131

Epoch 6/500

375/375 2s 6ms/step -

accuracy: 0.9540 - loss: 0.1538 - val_accuracy: 0.9707 - val_loss: 0.1025

Epoch 7/500

375/375 2s 6ms/step -

accuracy: 0.9582 - loss: 0.1383 - val_accuracy: 0.9699 - val_loss: 0.1021

Epoch 8/500

375/375 2s 6ms/step -

accuracy: 0.9599 - loss: 0.1319 - val_accuracy: 0.9734 - val_loss: 0.0913

Epoch 9/500

375/375 2s 6ms/step -

accuracy: 0.9651 - loss: 0.1150 - val_accuracy: 0.9741 - val_loss: 0.0900

Epoch 10/500

375/375 3s 7ms/step -

accuracy: 0.9662 - loss: 0.1120 - val_accuracy: 0.9757 - val_loss: 0.0849

Epoch 11/500

375/375 3s 7ms/step -

accuracy: 0.9700 - loss: 0.1007 - val_accuracy: 0.9768 - val_loss: 0.0815

Epoch 12/500

375/375 3s 7ms/step -

accuracy: 0.9709 - loss: 0.0953 - val_accuracy: 0.9755 - val_loss: 0.0851

Epoch 13/500

375/375 3s 7ms/step -

accuracy: 0.9729 - loss: 0.0859 - val_accuracy: 0.9783 - val_loss: 0.0781

Epoch 14/500

375/375 3s 7ms/step -

accuracy: 0.9745 - loss: 0.0814 - val_accuracy: 0.9792 - val_loss: 0.0751

Epoch 15/500

```

375/375          3s 7ms/step -
accuracy: 0.9742 - loss: 0.0817 - val_accuracy: 0.9789 - val_loss: 0.0751
Epoch 16/500
375/375          3s 8ms/step -
accuracy: 0.9764 - loss: 0.0755 - val_accuracy: 0.9801 - val_loss: 0.0745
Epoch 17/500
375/375          3s 8ms/step -
accuracy: 0.9776 - loss: 0.0748 - val_accuracy: 0.9791 - val_loss: 0.0740
Epoch 18/500
375/375          3s 9ms/step -
accuracy: 0.9775 - loss: 0.0723 - val_accuracy: 0.9803 - val_loss: 0.0742
Epoch 19/500
375/375          4s 10ms/step -
accuracy: 0.9783 - loss: 0.0672 - val_accuracy: 0.9798 - val_loss: 0.0748
Epoch 20/500
375/375          3s 7ms/step -
accuracy: 0.9804 - loss: 0.0633 - val_accuracy: 0.9803 - val_loss: 0.0740
Epoch 21/500
375/375          3s 7ms/step -
accuracy: 0.9804 - loss: 0.0623 - val_accuracy: 0.9810 - val_loss: 0.0722
Epoch 22/500
375/375          3s 7ms/step -
accuracy: 0.9829 - loss: 0.0543 - val_accuracy: 0.9814 - val_loss: 0.0695
Epoch 23/500
375/375          3s 7ms/step -
accuracy: 0.9826 - loss: 0.0557 - val_accuracy: 0.9811 - val_loss: 0.0713
Epoch 24/500
375/375          3s 8ms/step -
accuracy: 0.9836 - loss: 0.0539 - val_accuracy: 0.9803 - val_loss: 0.0698
Epoch 25/500
375/375          3s 8ms/step -
accuracy: 0.9837 - loss: 0.0511 - val_accuracy: 0.9810 - val_loss: 0.0721
Epoch 26/500
375/375          3s 7ms/step -
accuracy: 0.9860 - loss: 0.0452 - val_accuracy: 0.9816 - val_loss: 0.0695
Epoch 27/500
375/375          3s 7ms/step -
accuracy: 0.9838 - loss: 0.0498 - val_accuracy: 0.9817 - val_loss: 0.0736
Epoch 27: early stopping
Restoring model weights from the end of the best epoch: 22.

```

```
[45]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```


313/313          1s 2ms/step -
accuracy: 0.9762 - loss: 0.0817

```

```
[46]: print(f"Test Accuracy with ANN model: {val_accuracy*100:.4f}%")
```

```
Test Accuracy with ANN model: 98.0700%
```

```
[47]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epochs")
plt.ylabel('Loss')
plt.legend()
plt.show()
```



0.1.5 Question 2: Now refit the model with three learning rate schedulers, linear, polynomialdecay and exponentialdecay and report the answers. The fitting should be done with early stopping on.

```
[48]: def build_model():
    model = Sequential()
    model.add(Flatten(input_shape = (28,28)))
    model.add(Dense(512, activation = 'relu'))
    model.add(Dropout(0.5))
    model.add(Dense(256, activation = 'relu'))
    model.add(Dropout(0.3))
    model.add(Dense(128, activation = 'relu'))
    model.add(Dropout(0.2))
    model.add(Dense(10, activation = 'softmax'))
    return model
```

```
[49]: import keras.optimizers
```

```
initial_lr = 0.01
```

```
## Linear decay
```

```
linear_decay = keras.optimizers.schedules.PolynomialDecay(  
    initial_learning_rate= initial_lr,  
    decay_steps= 10000,  
    end_learning_rate=0.0,  
    power = 1.0  
)
```

```
## Polynomial Decay
```

```
polynomial_decay = keras.optimizers.schedules.PolynomialDecay(  
    initial_learning_rate=initial_lr,  
    decay_steps=10000,  
    end_learning_rate=0.0001,  
    power = 2.0  
)
```

```
## Exponential Decay
```

```
exponential_decay = keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=initial_lr,  
    decay_steps=1000,  
    decay_rate=0.9  
)
```

```
schedulers = {  
    "Linear Decay": linear_decay,  
    "Polynomial Decay": polynomial_decay,  
    "Exponential Decay": exponential_decay  
}
```

```
[50]: results = {}
```

```
for name, schedule in schedulers.items():  
    sgd = SGD(learning_rate=schedule, momentum=0.9)  
    print(f"\nTraining with {name} learning rate scheduler:\n")  
    model = build_model()  
    model.compile(optimizer=sgd, loss = 'sparse_categorical_crossentropy',  
↪ metrics = ['accuracy'])  
  
    history = model.fit(x_train,y_train, validation_data=(x_val,y_val),  
↪ batch_size=128, epochs = 500, callbacks=[estop], verbose = 1)  
    test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)  
    print(f"{name} - Test Accuracy: {test_acc*100:.4f}%, Test Loss: {test_loss:.  
↪ 4f}")
```

```

# Store results
results[name] = {
    "history": history.history,
    "test_acc": test_acc,
    "test_loss": test_loss
}

```

Training with Linear Decay learning rate scheduler:

```

Epoch 1/500
375/375      3s 7ms/step -
accuracy: 0.5811 - loss: 1.2293 - val_accuracy: 0.9234 - val_loss: 0.2504
Epoch 2/500
375/375      3s 7ms/step -
accuracy: 0.8992 - loss: 0.3363 - val_accuracy: 0.9457 - val_loss: 0.1758
Epoch 3/500
375/375      3s 7ms/step -
accuracy: 0.9253 - loss: 0.2468 - val_accuracy: 0.9551 - val_loss: 0.1467
Epoch 4/500
375/375      3s 7ms/step -
accuracy: 0.9383 - loss: 0.2063 - val_accuracy: 0.9619 - val_loss: 0.1266
Epoch 5/500
375/375      3s 7ms/step -
accuracy: 0.9471 - loss: 0.1791 - val_accuracy: 0.9657 - val_loss: 0.1171
Epoch 6/500
375/375      3s 8ms/step -
accuracy: 0.9522 - loss: 0.1595 - val_accuracy: 0.9703 - val_loss: 0.1044
Epoch 7/500
375/375      3s 7ms/step -
accuracy: 0.9584 - loss: 0.1340 - val_accuracy: 0.9722 - val_loss: 0.0980
Epoch 8/500
375/375      3s 7ms/step -
accuracy: 0.9620 - loss: 0.1286 - val_accuracy: 0.9721 - val_loss: 0.0954
Epoch 9/500
375/375      3s 7ms/step -
accuracy: 0.9664 - loss: 0.1109 - val_accuracy: 0.9739 - val_loss: 0.0913
Epoch 10/500
375/375      3s 7ms/step -
accuracy: 0.9656 - loss: 0.1142 - val_accuracy: 0.9747 - val_loss: 0.0883
Epoch 11/500
375/375      3s 7ms/step -
accuracy: 0.9675 - loss: 0.1061 - val_accuracy: 0.9768 - val_loss: 0.0846
Epoch 12/500
375/375      3s 7ms/step -
accuracy: 0.9700 - loss: 0.0993 - val_accuracy: 0.9785 - val_loss: 0.0829
Epoch 13/500
375/375      3s 7ms/step -

```


accuracy: 0.9713 - loss: 0.0942 - val_accuracy: 0.9784 - val_loss: 0.0807
 Epoch 14/500
 375/375 3s 7ms/step -
 accuracy: 0.9720 - loss: 0.0895 - val_accuracy: 0.9769 - val_loss: 0.0820
 Epoch 15/500
 375/375 3s 7ms/step -
 accuracy: 0.9735 - loss: 0.0886 - val_accuracy: 0.9777 - val_loss: 0.0799
 Epoch 16/500
 375/375 3s 7ms/step -
 accuracy: 0.9751 - loss: 0.0807 - val_accuracy: 0.9794 - val_loss: 0.0768
 Epoch 17/500
 375/375 3s 7ms/step -
 accuracy: 0.9750 - loss: 0.0797 - val_accuracy: 0.9798 - val_loss: 0.0761
 Epoch 18/500
 375/375 3s 7ms/step -
 accuracy: 0.9762 - loss: 0.0726 - val_accuracy: 0.9792 - val_loss: 0.0759
 Epoch 19/500
 375/375 3s 7ms/step -
 accuracy: 0.9777 - loss: 0.0767 - val_accuracy: 0.9801 - val_loss: 0.0761
 Epoch 20/500
 375/375 3s 7ms/step -
 accuracy: 0.9784 - loss: 0.0722 - val_accuracy: 0.9803 - val_loss: 0.0756
 Epoch 21/500
 375/375 3s 7ms/step -
 accuracy: 0.9776 - loss: 0.0738 - val_accuracy: 0.9804 - val_loss: 0.0754
 Epoch 22/500
 375/375 3s 7ms/step -
 accuracy: 0.9798 - loss: 0.0660 - val_accuracy: 0.9808 - val_loss: 0.0745
 Epoch 23/500
 375/375 3s 7ms/step -
 accuracy: 0.9810 - loss: 0.0631 - val_accuracy: 0.9809 - val_loss: 0.0748
 Epoch 24/500
 375/375 3s 7ms/step -
 accuracy: 0.9788 - loss: 0.0666 - val_accuracy: 0.9812 - val_loss: 0.0740
 Epoch 25/500
 375/375 3s 7ms/step -
 accuracy: 0.9815 - loss: 0.0599 - val_accuracy: 0.9813 - val_loss: 0.0740
 Epoch 26/500
 375/375 3s 7ms/step -
 accuracy: 0.9813 - loss: 0.0630 - val_accuracy: 0.9814 - val_loss: 0.0739
 Epoch 27/500
 375/375 3s 7ms/step -
 accuracy: 0.9814 - loss: 0.0601 - val_accuracy: 0.9816 - val_loss: 0.0738
 Epoch 28/500
 375/375 3s 7ms/step -
 accuracy: 0.9801 - loss: 0.0611 - val_accuracy: 0.9816 - val_loss: 0.0738
 Epoch 29/500
 375/375 3s 7ms/step -

accuracy: 0.9809 - loss: 0.0612 - val_accuracy: 0.9816 - val_loss: 0.0738
Epoch 30/500
375/375 3s 7ms/step -
accuracy: 0.9804 - loss: 0.0612 - val_accuracy: 0.9816 - val_loss: 0.0738
Epoch 31/500
375/375 3s 7ms/step -
accuracy: 0.9805 - loss: 0.0636 - val_accuracy: 0.9816 - val_loss: 0.0738
Epoch 32/500
375/375 3s 7ms/step -
accuracy: 0.9819 - loss: 0.0603 - val_accuracy: 0.9816 - val_loss: 0.0738
Epoch 32: early stopping
Restoring model weights from the end of the best epoch: 27.
Linear Decay - Test Accuracy: 98.0200%, Test Loss: 0.0670

Training with Polynomial Decay learning rate scheduler:

Epoch 1/500
375/375 3s 7ms/step -
accuracy: 0.5726 - loss: 1.2735 - val_accuracy: 0.9208 - val_loss: 0.2617
Epoch 2/500
375/375 3s 7ms/step -
accuracy: 0.8925 - loss: 0.3502 - val_accuracy: 0.9428 - val_loss: 0.1863
Epoch 3/500
375/375 3s 8ms/step -
accuracy: 0.9234 - loss: 0.2588 - val_accuracy: 0.9534 - val_loss: 0.1523
Epoch 4/500
375/375 3s 7ms/step -
accuracy: 0.9338 - loss: 0.2138 - val_accuracy: 0.9595 - val_loss: 0.1323
Epoch 5/500
375/375 3s 7ms/step -
accuracy: 0.9455 - loss: 0.1875 - val_accuracy: 0.9646 - val_loss: 0.1191
Epoch 6/500
375/375 3s 7ms/step -
accuracy: 0.9502 - loss: 0.1661 - val_accuracy: 0.9653 - val_loss: 0.1143
Epoch 7/500
375/375 3s 7ms/step -
accuracy: 0.9546 - loss: 0.1517 - val_accuracy: 0.9692 - val_loss: 0.1052
Epoch 8/500
375/375 3s 7ms/step -
accuracy: 0.9570 - loss: 0.1411 - val_accuracy: 0.9704 - val_loss: 0.1003
Epoch 9/500
375/375 3s 7ms/step -
accuracy: 0.9620 - loss: 0.1267 - val_accuracy: 0.9712 - val_loss: 0.0983
Epoch 10/500
375/375 3s 7ms/step -
accuracy: 0.9633 - loss: 0.1232 - val_accuracy: 0.9722 - val_loss: 0.0958
Epoch 11/500
375/375 3s 7ms/step -

accuracy: 0.9638 - loss: 0.1166 - val_accuracy: 0.9739 - val_loss: 0.0921
 Epoch 12/500
 375/375 3s 7ms/step -
 accuracy: 0.9653 - loss: 0.1145 - val_accuracy: 0.9737 - val_loss: 0.0904
 Epoch 13/500
 375/375 3s 7ms/step -
 accuracy: 0.9685 - loss: 0.1056 - val_accuracy: 0.9752 - val_loss: 0.0886
 Epoch 14/500
 375/375 3s 7ms/step -
 accuracy: 0.9686 - loss: 0.1071 - val_accuracy: 0.9756 - val_loss: 0.0864
 Epoch 15/500
 375/375 3s 7ms/step -
 accuracy: 0.9685 - loss: 0.1014 - val_accuracy: 0.9753 - val_loss: 0.0870
 Epoch 16/500
 375/375 3s 7ms/step -
 accuracy: 0.9706 - loss: 0.0972 - val_accuracy: 0.9761 - val_loss: 0.0854
 Epoch 17/500
 375/375 3s 7ms/step -
 accuracy: 0.9724 - loss: 0.0963 - val_accuracy: 0.9758 - val_loss: 0.0850
 Epoch 18/500
 375/375 3s 7ms/step -
 accuracy: 0.9713 - loss: 0.0943 - val_accuracy: 0.9761 - val_loss: 0.0845
 Epoch 19/500
 375/375 3s 7ms/step -
 accuracy: 0.9704 - loss: 0.0950 - val_accuracy: 0.9759 - val_loss: 0.0842
 Epoch 20/500
 375/375 3s 7ms/step -
 accuracy: 0.9728 - loss: 0.0908 - val_accuracy: 0.9762 - val_loss: 0.0837
 Epoch 21/500
 375/375 3s 7ms/step -
 accuracy: 0.9733 - loss: 0.0882 - val_accuracy: 0.9766 - val_loss: 0.0831
 Epoch 22/500
 375/375 3s 7ms/step -
 accuracy: 0.9735 - loss: 0.0893 - val_accuracy: 0.9766 - val_loss: 0.0831
 Epoch 23/500
 375/375 3s 7ms/step -
 accuracy: 0.9741 - loss: 0.0860 - val_accuracy: 0.9769 - val_loss: 0.0831
 Epoch 24/500
 375/375 3s 7ms/step -
 accuracy: 0.9727 - loss: 0.0910 - val_accuracy: 0.9770 - val_loss: 0.0828
 Epoch 25/500
 375/375 3s 8ms/step -
 accuracy: 0.9752 - loss: 0.0863 - val_accuracy: 0.9770 - val_loss: 0.0827
 Epoch 26/500
 375/375 3s 7ms/step -
 accuracy: 0.9730 - loss: 0.0864 - val_accuracy: 0.9770 - val_loss: 0.0826
 Epoch 27/500
 375/375 3s 7ms/step -

accuracy: 0.9725 - loss: 0.0869 - val_accuracy: 0.9771 - val_loss: 0.0826
 Epoch 28/500
 375/375 3s 7ms/step -
 accuracy: 0.9727 - loss: 0.0907 - val_accuracy: 0.9772 - val_loss: 0.0826
 Epoch 29/500
 375/375 3s 7ms/step -
 accuracy: 0.9727 - loss: 0.0871 - val_accuracy: 0.9771 - val_loss: 0.0825
 Epoch 30/500
 375/375 3s 7ms/step -
 accuracy: 0.9746 - loss: 0.0814 - val_accuracy: 0.9771 - val_loss: 0.0825
 Epoch 31/500
 375/375 3s 7ms/step -
 accuracy: 0.9735 - loss: 0.0866 - val_accuracy: 0.9770 - val_loss: 0.0824
 Epoch 32/500
 375/375 3s 7ms/step -
 accuracy: 0.9737 - loss: 0.0844 - val_accuracy: 0.9770 - val_loss: 0.0824
 Epoch 33/500
 375/375 3s 7ms/step -
 accuracy: 0.9742 - loss: 0.0840 - val_accuracy: 0.9770 - val_loss: 0.0824
 Epoch 34/500
 375/375 3s 7ms/step -
 accuracy: 0.9759 - loss: 0.0845 - val_accuracy: 0.9769 - val_loss: 0.0824
 Epoch 35/500
 375/375 3s 7ms/step -
 accuracy: 0.9756 - loss: 0.0806 - val_accuracy: 0.9770 - val_loss: 0.0823
 Epoch 36/500
 375/375 3s 7ms/step -
 accuracy: 0.9737 - loss: 0.0869 - val_accuracy: 0.9770 - val_loss: 0.0822
 Epoch 37/500
 375/375 3s 8ms/step -
 accuracy: 0.9749 - loss: 0.0829 - val_accuracy: 0.9770 - val_loss: 0.0823
 Epoch 38/500
 375/375 3s 7ms/step -
 accuracy: 0.9707 - loss: 0.0904 - val_accuracy: 0.9769 - val_loss: 0.0822
 Epoch 39/500
 375/375 3s 8ms/step -
 accuracy: 0.9737 - loss: 0.0853 - val_accuracy: 0.9769 - val_loss: 0.0821
 Epoch 40/500
 375/375 3s 7ms/step -
 accuracy: 0.9722 - loss: 0.0859 - val_accuracy: 0.9769 - val_loss: 0.0821
 Epoch 41/500
 375/375 3s 8ms/step -
 accuracy: 0.9739 - loss: 0.0842 - val_accuracy: 0.9770 - val_loss: 0.0820
 Epoch 42/500
 375/375 3s 8ms/step -
 accuracy: 0.9740 - loss: 0.0854 - val_accuracy: 0.9772 - val_loss: 0.0819
 Epoch 43/500
 375/375 3s 8ms/step -

accuracy: 0.9736 - loss: 0.0861 - val_accuracy: 0.9772 - val_loss: 0.0820
 Epoch 44/500
 375/375 3s 7ms/step -
 accuracy: 0.9732 - loss: 0.0876 - val_accuracy: 0.9772 - val_loss: 0.0819
 Epoch 45/500
 375/375 3s 8ms/step -
 accuracy: 0.9724 - loss: 0.0887 - val_accuracy: 0.9772 - val_loss: 0.0820
 Epoch 46/500
 375/375 3s 7ms/step -
 accuracy: 0.9735 - loss: 0.0855 - val_accuracy: 0.9772 - val_loss: 0.0821
 Epoch 47/500
 375/375 3s 8ms/step -
 accuracy: 0.9734 - loss: 0.0835 - val_accuracy: 0.9770 - val_loss: 0.0820
 Epoch 47: early stopping
 Restoring model weights from the end of the best epoch: 42.
 Polynomial Decay - Test Accuracy: 97.6400%, Test Loss: 0.0770

Training with Exponential Decay learning rate scheduler:

Epoch 1/500
 375/375 4s 8ms/step -
 accuracy: 0.5878 - loss: 1.2296 - val_accuracy: 0.9227 - val_loss: 0.2570
 Epoch 2/500
 375/375 3s 8ms/step -
 accuracy: 0.8972 - loss: 0.3423 - val_accuracy: 0.9445 - val_loss: 0.1753
 Epoch 3/500
 375/375 3s 7ms/step -
 accuracy: 0.9231 - loss: 0.2518 - val_accuracy: 0.9544 - val_loss: 0.1486
 Epoch 4/500
 375/375 3s 8ms/step -
 accuracy: 0.9385 - loss: 0.2098 - val_accuracy: 0.9609 - val_loss: 0.1307
 Epoch 5/500
 375/375 3s 7ms/step -
 accuracy: 0.9468 - loss: 0.1741 - val_accuracy: 0.9627 - val_loss: 0.1203
 Epoch 6/500
 375/375 3s 8ms/step -
 accuracy: 0.9504 - loss: 0.1632 - val_accuracy: 0.9673 - val_loss: 0.1070
 Epoch 7/500
 375/375 3s 8ms/step -
 accuracy: 0.9569 - loss: 0.1420 - val_accuracy: 0.9695 - val_loss: 0.1030
 Epoch 8/500
 375/375 3s 8ms/step -
 accuracy: 0.9593 - loss: 0.1339 - val_accuracy: 0.9712 - val_loss: 0.0987
 Epoch 9/500
 375/375 3s 8ms/step -
 accuracy: 0.9660 - loss: 0.1180 - val_accuracy: 0.9712 - val_loss: 0.0958
 Epoch 10/500
 375/375 3s 9ms/step -

accuracy: 0.9681 - loss: 0.1067 - val_accuracy: 0.9743 - val_loss: 0.0898
 Epoch 11/500
 375/375 3s 8ms/step -
 accuracy: 0.9672 - loss: 0.1050 - val_accuracy: 0.9738 - val_loss: 0.0901
 Epoch 12/500
 375/375 3s 8ms/step -
 accuracy: 0.9708 - loss: 0.0956 - val_accuracy: 0.9751 - val_loss: 0.0874
 Epoch 13/500
 375/375 3s 8ms/step -
 accuracy: 0.9703 - loss: 0.0973 - val_accuracy: 0.9749 - val_loss: 0.0832
 Epoch 14/500
 375/375 3s 8ms/step -
 accuracy: 0.9701 - loss: 0.0938 - val_accuracy: 0.9759 - val_loss: 0.0814
 Epoch 15/500
 375/375 3s 7ms/step -
 accuracy: 0.9720 - loss: 0.0888 - val_accuracy: 0.9767 - val_loss: 0.0817
 Epoch 16/500
 375/375 3s 8ms/step -
 accuracy: 0.9748 - loss: 0.0820 - val_accuracy: 0.9779 - val_loss: 0.0771
 Epoch 17/500
 375/375 3s 8ms/step -
 accuracy: 0.9759 - loss: 0.0791 - val_accuracy: 0.9778 - val_loss: 0.0789
 Epoch 18/500
 375/375 3s 8ms/step -
 accuracy: 0.9766 - loss: 0.0781 - val_accuracy: 0.9781 - val_loss: 0.0772
 Epoch 19/500
 375/375 3s 8ms/step -
 accuracy: 0.9792 - loss: 0.0681 - val_accuracy: 0.9785 - val_loss: 0.0770
 Epoch 20/500
 375/375 3s 8ms/step -
 accuracy: 0.9777 - loss: 0.0699 - val_accuracy: 0.9787 - val_loss: 0.0767
 Epoch 21/500
 375/375 3s 8ms/step -
 accuracy: 0.9786 - loss: 0.0696 - val_accuracy: 0.9787 - val_loss: 0.0769
 Epoch 22/500
 375/375 3s 8ms/step -
 accuracy: 0.9794 - loss: 0.0651 - val_accuracy: 0.9783 - val_loss: 0.0756
 Epoch 23/500
 375/375 3s 8ms/step -
 accuracy: 0.9805 - loss: 0.0642 - val_accuracy: 0.9786 - val_loss: 0.0758
 Epoch 24/500
 375/375 3s 8ms/step -
 accuracy: 0.9817 - loss: 0.0606 - val_accuracy: 0.9791 - val_loss: 0.0747
 Epoch 25/500
 375/375 3s 9ms/step -
 accuracy: 0.9815 - loss: 0.0600 - val_accuracy: 0.9795 - val_loss: 0.0731
 Epoch 26/500
 375/375 3s 8ms/step -

```

accuracy: 0.9815 - loss: 0.0587 - val_accuracy: 0.9791 - val_loss: 0.0737
Epoch 27/500
375/375          3s 9ms/step -
accuracy: 0.9822 - loss: 0.0547 - val_accuracy: 0.9793 - val_loss: 0.0737
Epoch 28/500
375/375          3s 8ms/step -
accuracy: 0.9820 - loss: 0.0570 - val_accuracy: 0.9803 - val_loss: 0.0728
Epoch 29/500
375/375          3s 8ms/step -
accuracy: 0.9832 - loss: 0.0537 - val_accuracy: 0.9803 - val_loss: 0.0724
Epoch 30/500
375/375          3s 8ms/step -
accuracy: 0.9841 - loss: 0.0523 - val_accuracy: 0.9797 - val_loss: 0.0735
Epoch 31/500
375/375          3s 8ms/step -
accuracy: 0.9837 - loss: 0.0515 - val_accuracy: 0.9801 - val_loss: 0.0729
Epoch 32/500
375/375          3s 8ms/step -
accuracy: 0.9838 - loss: 0.0526 - val_accuracy: 0.9793 - val_loss: 0.0740
Epoch 33/500
375/375          3s 8ms/step -
accuracy: 0.9852 - loss: 0.0485 - val_accuracy: 0.9805 - val_loss: 0.0738
Epoch 34/500
375/375          3s 8ms/step -
accuracy: 0.9847 - loss: 0.0504 - val_accuracy: 0.9803 - val_loss: 0.0719
Epoch 35/500
375/375          3s 9ms/step -
accuracy: 0.9847 - loss: 0.0465 - val_accuracy: 0.9803 - val_loss: 0.0726
Epoch 36/500
375/375          3s 8ms/step -
accuracy: 0.9852 - loss: 0.0476 - val_accuracy: 0.9803 - val_loss: 0.0722
Epoch 37/500
375/375          3s 8ms/step -
accuracy: 0.9862 - loss: 0.0455 - val_accuracy: 0.9800 - val_loss: 0.0730
Epoch 38/500
375/375          3s 8ms/step -
accuracy: 0.9854 - loss: 0.0464 - val_accuracy: 0.9808 - val_loss: 0.0733
Epoch 39/500
375/375          3s 8ms/step -
accuracy: 0.9864 - loss: 0.0421 - val_accuracy: 0.9803 - val_loss: 0.0738
Epoch 39: early stopping
Restoring model weights from the end of the best epoch: 34.
Exponential Decay - Test Accuracy: 98.1600%, Test Loss: 0.0633

```

0.1.6 Plotting loss and accuracy curves

```
[51]: plt.figure(figsize=(14,6))


# Accuracy comparison
for name, res in results.items():
    plt.plot(res["history"]["val_accuracy"], label=f"{name} Val Acc")

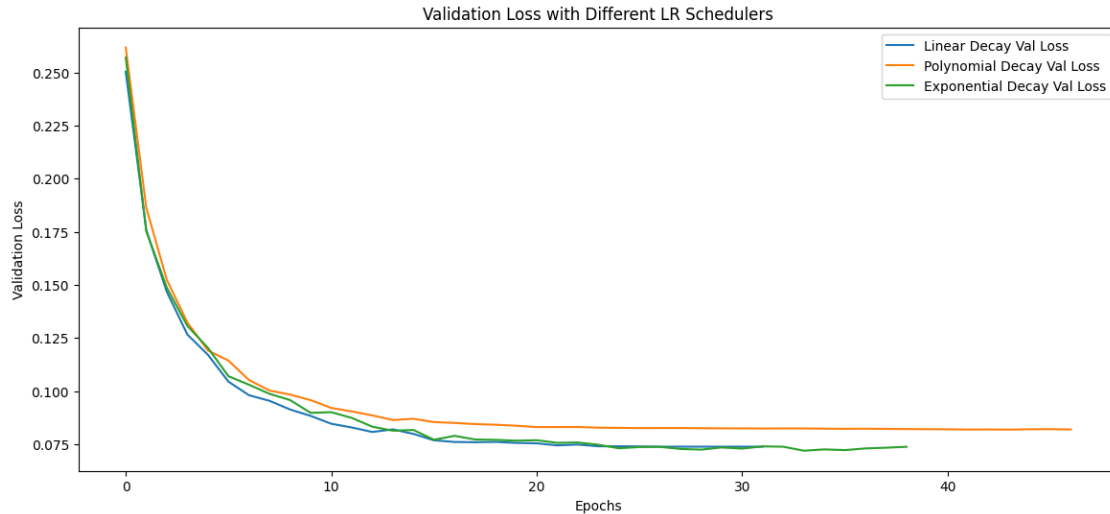
plt.xlabel("Epochs")
plt.ylabel("Validation Accuracy")
plt.title("Validation Accuracy with Different LR Schedulers")
plt.legend()
plt.show()

plt.figure(figsize=(14,6))

# Loss comparison
for name, res in results.items():
    plt.plot(res["history"]["val_loss"], label=f"{name} Val Loss")

plt.xlabel("Epochs")
plt.ylabel("Validation Loss")
plt.title("Validation Loss with Different LR Schedulers")
plt.legend()
plt.show()
```





```
[52]: for name, res in results.items():
      print(f"\n{name}:")
      print(f"  Final Test Accuracy = {res['test_acc']*100:.4f}%")
      print(f"  Final Test Loss      = {res['test_loss']:.4f}")
```

Linear Decay:

```
Final Test Accuracy = 98.0200%
Final Test Loss      = 0.0670
```

Polynomial Decay:

```
Final Test Accuracy = 97.6400%
Final Test Loss      = 0.0770
```

Exponential Decay:

```
Final Test Accuracy = 98.1600%
Final Test Loss      = 0.0633
```

0.2 Question 3: Optimizer Comparison

0.2.1 Report the best optimizer that would result in the best performance for the above models. Try at least three to four optimizers (e.g., SGD with momentum, RMSprop, Adam). Train the same model architecture with each optimizer. Compare their performances. Report which optimizer gives the best results

```
[53]: def build_model():
      model = Sequential()
      model.add(Flatten(input_shape=(28,28)))
      model.add(Dense(512, activation='relu'))
      model.add(Dropout(0.5))
```

```

model.add(Dense(256, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
return model

```

```

[54]: from keras.optimizers import SGD, Adam, RMSprop, Adamax
optimizers = {
    "SGD": SGD(learning_rate=0.01, momentum=0.9),
    "RMSprop": RMSprop(learning_rate=0.001),
    "Adam": Adam(learning_rate=0.001),
    "Adamax": Adamax(learning_rate=0.002)
}

estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
↳ verbose = 1, restore_best_weights=True)

```

```

[55]: results = {}

for name, optimizer in optimizers.items():
    print(f"\nTraining with {name} optimizer:\n")

    model = build_model()
    model.compile(optimizer=optimizer,
                  loss = 'sparse_categorical_crossentropy',
                  metrics = ['accuracy'])
    history = model.fit(
        x_train, y_train,
        validation_data=(x_val, y_val),
        batch_size=128,
        epochs=500,
        callbacks=[estop],
        verbose=1
    )
    # Evaluating on test set
    test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
    print(f"{name} - Test Accuracy: {test_acc*100:.4f}%, Test Loss: {test_loss:.
↳ 4f}")

    results[name] = {
        "history": history.history,
        "test_acc": test_acc,
        "test_loss": test_loss
    }

```

Training with SGD optimizer:

Epoch 1/500
 375/375 3s 7ms/step -
 accuracy: 0.5820 - loss: 1.2470 - val_accuracy: 0.9237 - val_loss: 0.2473

Epoch 2/500
 375/375 2s 7ms/step -
 accuracy: 0.8986 - loss: 0.3402 - val_accuracy: 0.9463 - val_loss: 0.1743

Epoch 3/500
 375/375 3s 7ms/step -
 accuracy: 0.9257 - loss: 0.2489 - val_accuracy: 0.9560 - val_loss: 0.1439

Epoch 4/500
 375/375 3s 7ms/step -
 accuracy: 0.9385 - loss: 0.2056 - val_accuracy: 0.9617 - val_loss: 0.1243

Epoch 5/500
 375/375 2s 6ms/step -
 accuracy: 0.9465 - loss: 0.1781 - val_accuracy: 0.9647 - val_loss: 0.1170

Epoch 6/500
 375/375 3s 7ms/step -
 accuracy: 0.9536 - loss: 0.1527 - val_accuracy: 0.9693 - val_loss: 0.1026

Epoch 7/500
 375/375 3s 7ms/step -
 accuracy: 0.9600 - loss: 0.1366 - val_accuracy: 0.9712 - val_loss: 0.1001

Epoch 8/500
 375/375 3s 7ms/step -
 accuracy: 0.9620 - loss: 0.1259 - val_accuracy: 0.9736 - val_loss: 0.0910

Epoch 9/500
 375/375 2s 7ms/step -
 accuracy: 0.9637 - loss: 0.1175 - val_accuracy: 0.9737 - val_loss: 0.0884

Epoch 10/500
 375/375 3s 7ms/step -
 accuracy: 0.9648 - loss: 0.1152 - val_accuracy: 0.9756 - val_loss: 0.0857

Epoch 11/500
 375/375 3s 8ms/step -
 accuracy: 0.9697 - loss: 0.0987 - val_accuracy: 0.9758 - val_loss: 0.0806

Epoch 12/500
 375/375 3s 7ms/step -
 accuracy: 0.9706 - loss: 0.0984 - val_accuracy: 0.9755 - val_loss: 0.0858

Epoch 13/500
 375/375 3s 7ms/step -
 accuracy: 0.9732 - loss: 0.0891 - val_accuracy: 0.9777 - val_loss: 0.0785

Epoch 14/500
 375/375 3s 7ms/step -
 accuracy: 0.9754 - loss: 0.0826 - val_accuracy: 0.9787 - val_loss: 0.0788

Epoch 15/500
 375/375 3s 7ms/step -
 accuracy: 0.9752 - loss: 0.0798 - val_accuracy: 0.9778 - val_loss: 0.0791

Epoch 16/500
 375/375 3s 7ms/step -

accuracy: 0.9755 - loss: 0.0787 - val_accuracy: 0.9784 - val_loss: 0.0756
 Epoch 17/500
 375/375 3s 8ms/step -
 accuracy: 0.9766 - loss: 0.0728 - val_accuracy: 0.9799 - val_loss: 0.0755
 Epoch 18/500
 375/375 3s 7ms/step -
 accuracy: 0.9793 - loss: 0.0698 - val_accuracy: 0.9797 - val_loss: 0.0717
 Epoch 19/500
 375/375 3s 8ms/step -
 accuracy: 0.9793 - loss: 0.0666 - val_accuracy: 0.9791 - val_loss: 0.0771
 Epoch 20/500
 375/375 3s 7ms/step -
 accuracy: 0.9798 - loss: 0.0632 - val_accuracy: 0.9794 - val_loss: 0.0724
 Epoch 21/500
 375/375 3s 7ms/step -
 accuracy: 0.9812 - loss: 0.0603 - val_accuracy: 0.9804 - val_loss: 0.0726
 Epoch 22/500
 375/375 2s 7ms/step -
 accuracy: 0.9800 - loss: 0.0613 - val_accuracy: 0.9806 - val_loss: 0.0715
 Epoch 23/500
 375/375 3s 7ms/step -
 accuracy: 0.9819 - loss: 0.0580 - val_accuracy: 0.9797 - val_loss: 0.0739
 Epoch 24/500
 375/375 3s 8ms/step -
 accuracy: 0.9832 - loss: 0.0528 - val_accuracy: 0.9805 - val_loss: 0.0760
 Epoch 25/500
 375/375 3s 8ms/step -
 accuracy: 0.9833 - loss: 0.0527 - val_accuracy: 0.9798 - val_loss: 0.0737
 Epoch 26/500
 375/375 3s 7ms/step -
 accuracy: 0.9846 - loss: 0.0495 - val_accuracy: 0.9799 - val_loss: 0.0753
 Epoch 27/500
 375/375 3s 7ms/step -
 accuracy: 0.9843 - loss: 0.0496 - val_accuracy: 0.9807 - val_loss: 0.0721
 Epoch 27: early stopping
 Restoring model weights from the end of the best epoch: 22.
 SGD - Test Accuracy: 98.1900%, Test Loss: 0.0638

Training with RMSprop optimizer:

Epoch 1/500
 375/375 4s 9ms/step -
 accuracy: 0.7769 - loss: 0.6982 - val_accuracy: 0.9516 - val_loss: 0.1627
 Epoch 2/500
 375/375 3s 9ms/step -
 accuracy: 0.9400 - loss: 0.2034 - val_accuracy: 0.9666 - val_loss: 0.1122
 Epoch 3/500
 375/375 3s 8ms/step -

accuracy: 0.9534 - loss: 0.1557 - val_accuracy: 0.9717 - val_loss: 0.1035
 Epoch 4/500
 375/375 3s 8ms/step -
 accuracy: 0.9612 - loss: 0.1275 - val_accuracy: 0.9719 - val_loss: 0.1004
 Epoch 5/500
 375/375 3s 8ms/step -
 accuracy: 0.9663 - loss: 0.1158 - val_accuracy: 0.9747 - val_loss: 0.0919
 Epoch 6/500
 375/375 3s 8ms/step -
 accuracy: 0.9707 - loss: 0.1004 - val_accuracy: 0.9783 - val_loss: 0.0824
 Epoch 7/500
 375/375 3s 9ms/step -
 accuracy: 0.9714 - loss: 0.0934 - val_accuracy: 0.9789 - val_loss: 0.0837
 Epoch 8/500
 375/375 3s 8ms/step -
 accuracy: 0.9742 - loss: 0.0911 - val_accuracy: 0.9780 - val_loss: 0.0829
 Epoch 9/500
 375/375 3s 8ms/step -
 accuracy: 0.9773 - loss: 0.0777 - val_accuracy: 0.9766 - val_loss: 0.0935
 Epoch 10/500
 375/375 3s 8ms/step -
 accuracy: 0.9766 - loss: 0.0793 - val_accuracy: 0.9797 - val_loss: 0.0841
 Epoch 11/500
 375/375 3s 8ms/step -
 accuracy: 0.9782 - loss: 0.0741 - val_accuracy: 0.9794 - val_loss: 0.0841
 Epoch 11: early stopping
 Restoring model weights from the end of the best epoch: 6.
 RMSprop - Test Accuracy: 97.8600%, Test Loss: 0.0766

Training with Adam optimizer:

Epoch 1/500
 375/375 4s 9ms/step -
 accuracy: 0.7487 - loss: 0.7641 - val_accuracy: 0.9537 - val_loss: 0.1502
 Epoch 2/500
 375/375 3s 9ms/step -
 accuracy: 0.9392 - loss: 0.2043 - val_accuracy: 0.9664 - val_loss: 0.1108
 Epoch 3/500
 375/375 3s 9ms/step -
 accuracy: 0.9536 - loss: 0.1570 - val_accuracy: 0.9714 - val_loss: 0.1003
 Epoch 4/500
 375/375 3s 8ms/step -
 accuracy: 0.9619 - loss: 0.1261 - val_accuracy: 0.9747 - val_loss: 0.0907
 Epoch 5/500
 375/375 3s 9ms/step -
 accuracy: 0.9674 - loss: 0.1077 - val_accuracy: 0.9771 - val_loss: 0.0819
 Epoch 6/500
 375/375 3s 8ms/step -

accuracy: 0.9695 - loss: 0.1004 - val_accuracy: 0.9772 - val_loss: 0.0796
 Epoch 7/500
 375/375 3s 9ms/step -
 accuracy: 0.9722 - loss: 0.0909 - val_accuracy: 0.9789 - val_loss: 0.0753
 Epoch 8/500
 375/375 3s 8ms/step -
 accuracy: 0.9748 - loss: 0.0800 - val_accuracy: 0.9803 - val_loss: 0.0717
 Epoch 9/500
 375/375 3s 8ms/step -
 accuracy: 0.9760 - loss: 0.0789 - val_accuracy: 0.9811 - val_loss: 0.0714
 Epoch 10/500
 375/375 3s 9ms/step -
 accuracy: 0.9781 - loss: 0.0735 - val_accuracy: 0.9809 - val_loss: 0.0748
 Epoch 11/500
 375/375 4s 10ms/step -
 accuracy: 0.9796 - loss: 0.0658 - val_accuracy: 0.9804 - val_loss: 0.0751
 Epoch 12/500
 375/375 3s 9ms/step -
 accuracy: 0.9809 - loss: 0.0637 - val_accuracy: 0.9799 - val_loss: 0.0748
 Epoch 13/500
 375/375 3s 8ms/step -
 accuracy: 0.9815 - loss: 0.0609 - val_accuracy: 0.9795 - val_loss: 0.0766
 Epoch 14/500
 375/375 3s 8ms/step -
 accuracy: 0.9808 - loss: 0.0585 - val_accuracy: 0.9812 - val_loss: 0.0715
 Epoch 14: early stopping
 Restoring model weights from the end of the best epoch: 9.
 Adam - Test Accuracy: 97.9700%, Test Loss: 0.0671

Training with Adamax optimizer:

Epoch 1/500
 375/375 5s 10ms/step -
 accuracy: 0.7543 - loss: 0.7553 - val_accuracy: 0.9427 - val_loss: 0.1813
 Epoch 2/500
 375/375 3s 9ms/step -
 accuracy: 0.9311 - loss: 0.2269 - val_accuracy: 0.9601 - val_loss: 0.1302
 Epoch 3/500
 375/375 3s 9ms/step -
 accuracy: 0.9504 - loss: 0.1704 - val_accuracy: 0.9684 - val_loss: 0.1073
 Epoch 4/500
 375/375 3s 9ms/step -
 accuracy: 0.9587 - loss: 0.1397 - val_accuracy: 0.9710 - val_loss: 0.0986
 Epoch 5/500
 375/375 3s 9ms/step -
 accuracy: 0.9640 - loss: 0.1200 - val_accuracy: 0.9742 - val_loss: 0.0933
 Epoch 6/500
 375/375 4s 9ms/step -

accuracy: 0.9672 - loss: 0.1025 - val_accuracy: 0.9765 - val_loss: 0.0828
 Epoch 7/500
 375/375 3s 9ms/step -
 accuracy: 0.9707 - loss: 0.0938 - val_accuracy: 0.9764 - val_loss: 0.0835
 Epoch 8/500
 375/375 3s 9ms/step -
 accuracy: 0.9736 - loss: 0.0821 - val_accuracy: 0.9772 - val_loss: 0.0828
 Epoch 9/500
 375/375 3s 9ms/step -
 accuracy: 0.9750 - loss: 0.0798 - val_accuracy: 0.9781 - val_loss: 0.0785
 Epoch 10/500
 375/375 3s 9ms/step -
 accuracy: 0.9767 - loss: 0.0736 - val_accuracy: 0.9787 - val_loss: 0.0776
 Epoch 11/500
 375/375 3s 9ms/step -
 accuracy: 0.9794 - loss: 0.0674 - val_accuracy: 0.9786 - val_loss: 0.0790
 Epoch 12/500
 375/375 3s 9ms/step -
 accuracy: 0.9795 - loss: 0.0653 - val_accuracy: 0.9811 - val_loss: 0.0737
 Epoch 13/500
 375/375 3s 9ms/step -
 accuracy: 0.9812 - loss: 0.0587 - val_accuracy: 0.9809 - val_loss: 0.0729
 Epoch 14/500
 375/375 3s 9ms/step -
 accuracy: 0.9840 - loss: 0.0534 - val_accuracy: 0.9793 - val_loss: 0.0754
 Epoch 15/500
 375/375 4s 9ms/step -
 accuracy: 0.9824 - loss: 0.0540 - val_accuracy: 0.9816 - val_loss: 0.0738
 Epoch 16/500
 375/375 3s 9ms/step -
 accuracy: 0.9837 - loss: 0.0489 - val_accuracy: 0.9807 - val_loss: 0.0720
 Epoch 17/500
 375/375 3s 9ms/step -
 accuracy: 0.9849 - loss: 0.0463 - val_accuracy: 0.9822 - val_loss: 0.0726
 Epoch 18/500
 375/375 3s 9ms/step -
 accuracy: 0.9879 - loss: 0.0395 - val_accuracy: 0.9827 - val_loss: 0.0711
 Epoch 19/500
 375/375 3s 9ms/step -
 accuracy: 0.9858 - loss: 0.0421 - val_accuracy: 0.9827 - val_loss: 0.0688
 Epoch 20/500
 375/375 3s 9ms/step -
 accuracy: 0.9879 - loss: 0.0392 - val_accuracy: 0.9824 - val_loss: 0.0718
 Epoch 21/500
 375/375 3s 9ms/step -
 accuracy: 0.9872 - loss: 0.0384 - val_accuracy: 0.9824 - val_loss: 0.0698
 Epoch 22/500
 375/375 3s 9ms/step -

```

accuracy: 0.9869 - loss: 0.0408 - val_accuracy: 0.9817 - val_loss: 0.0732
Epoch 23/500
375/375          3s 9ms/step -
accuracy: 0.9885 - loss: 0.0338 - val_accuracy: 0.9817 - val_loss: 0.0705
Epoch 24/500
375/375          4s 10ms/step -
accuracy: 0.9900 - loss: 0.0300 - val_accuracy: 0.9814 - val_loss: 0.0756
Epoch 24: early stopping
Restoring model weights from the end of the best epoch: 19.
Adamax - Test Accuracy: 98.1300%, Test Loss: 0.0676


```


```

[56]: plt.figure(figsize=(14,6))
      for name, res in results.items():
          plt.plot(res["history"]["val_accuracy"], label=f"{name}")
      plt.xlabel("Epochs")
      plt.ylabel("Validation Accuracy")
      plt.title("Validation Accuracy with Different Optimizers")
      plt.legend()
      plt.show()

      plt.figure(figsize=(14,6))
      for name, res in results.items():
          plt.plot(res["history"]["val_loss"], label=f"{name}")
      plt.xlabel("Epochs")
      plt.ylabel("Validation Loss")
      plt.title("Validation Loss with Different Optimizers")
      plt.legend()
      plt.show()

```





```
[57]: for name, res in results.items():
      print(f"\n{name}:")
      print(f"\tFinal Test Accuracy = {res['test_acc']*100:.4f}%")
      print(f"\tFinal Test Loss      = {res['test_loss']:.4f}")
```

SGD:

```
Final Test Accuracy = 98.1900%
Final Test Loss      = 0.0638
```

RMSprop:

```
Final Test Accuracy = 97.8600%
Final Test Loss      = 0.0766
```

Adam:

```
Final Test Accuracy = 97.9700%
Final Test Loss      = 0.0671
```

Adamax:

```
Final Test Accuracy = 98.1300%
Final Test Loss      = 0.0676
```

- Adamax Optimizer gives the best test accuracy