

DA_4

November 1, 2025

Name:	Tufan Kundu
Registration no:	24MDT0184
Course Name:	Deep Learning Lab
Course Code:	PMDS603P
Digital Assignment:	4

- 0.1 Question 1: Collect the dataset regarding ECG (Electrocardiogram) signals of different subjects given in the moodle platform and prepare an autoencoder that can perform anomaly detection. The dataset you see has two classes, normal (0) and abnormal (1) ECG signals. The labels are also provided in the same csv file. The task is to construct an autoencoder that can detect an abnormal ECG signal. Train an autoencoder on the normal ECG signals. Then use a suitable method using reconstruction error to flag abnormal ecg's. And you can check with your test set data the outputs.

0.1.1 Importing necessary libraries

```
[80]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, roc_auc_score, confusion_matrix
import seaborn as sns
import tensorflow as tf
from tensorflow.keras import regularizers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
```

0.1.2 Loading the dataset

```
[81]: df = pd.read_csv("as.csv")
df
```

```
[81]:      Unnamed: 0         0         1         2         3         4         5 \
0          0 -0.112522 -2.827204 -3.773897 -4.349751 -4.376041 -3.474986
1          1 -1.100878 -3.996840 -4.285843 -4.506579 -4.022377 -3.234368
2          2 -0.567088 -2.593450 -3.874230 -4.584095 -4.187449 -3.151462
3          3  0.490473 -1.914407 -3.616364 -4.318823 -4.268016 -3.881110
4          4  0.800232 -0.874252 -2.384761 -3.973292 -4.338224 -3.802422
...
4993     ...   ...   ...   ...   ...   ...
4993     4993  0.608558 -0.335651 -0.990948 -1.784153 -2.626145 -2.957065
4994     4994 -2.060402 -2.860116 -3.405074 -3.748719 -3.513561 -3.006545
4995     4995 -1.122969 -2.252925 -2.867628 -3.358605 -3.167849 -2.638360
4996     4996 -0.547705 -1.889545 -2.839779 -3.457912 -3.929149 -3.966026
4997     4997 -1.351779 -2.209006 -2.520225 -3.061475 -3.065141 -3.030739

          6         7         8   ...    131        132        133 \
0 -2.181408 -1.818286 -1.250522   ...  0.792168  0.933541  0.796958
1 -1.566126 -0.992258 -0.754680   ...  0.538356  0.656881  0.787490
2 -1.742940 -1.490659 -1.183580   ...  0.886073  0.531452  0.311377
3 -2.993280 -1.671131 -1.333884   ...  0.350816  0.499111  0.600345
4 -2.534510 -1.783423 -1.594450   ...  1.148884  0.958434  1.059025
...
4993 ...   ...   ...   ...   ...
4993 -2.931897 -2.664816 -2.090137   ...  1.757705  2.291923  2.704595
4994 -2.234850 -1.593270 -1.075279   ...  1.388947  2.079675  2.433375
4995 -1.664162 -0.935655 -0.866953   ... -0.472419 -1.310147 -2.029521
4996 -3.492560 -2.695270 -1.849691   ...  1.258419  1.907530  2.280888
4997 -2.622720 -2.044092 -1.295874   ... -1.512234 -2.076075 -2.586042

          134        135        136        137        138        139        140
0  0.578621  0.257740  0.228077  0.123431  0.925286  0.193137  1.0
1  0.724046  0.555784  0.476333  0.773820  1.119621 -1.436250  1.0
2 -0.021919 -0.713683 -0.532197  0.321097  0.904227 -0.421797  1.0
3  0.842069  0.952074  0.990133  1.086798  1.403011 -0.383564  1.0
4  1.371682  1.277392  0.960304  0.971020  1.614392  1.421456  1.0
...
4993 2.451519  2.017396  1.704358  1.688542  1.629593  1.342651  0.0
4994 2.159484  1.819747  1.534767  1.696818  1.483832  1.047612  0.0
4995 -3.221294 -4.176790 -4.009720 -2.874136 -2.008369 -1.808334  0.0
4996 1.895242  1.437702  1.193433  1.261335  1.150449  0.804932  0.0
4997 -3.322799 -3.627311 -3.437038 -2.260023 -1.577823 -0.684531  0.0
```

[4998 rows x 142 columns]

```
[82]: ## dropping the first column containing the serial no
df = df.drop(df.columns[0], axis=1)
df
```

```
[82]:      0         1         2         3         4         5         6 \
0 -0.112522 -2.827204 -3.773897 -4.349751 -4.376041 -3.474986 -2.181408
```

```

1   -1.100878 -3.996840 -4.285843 -4.506579 -4.022377 -3.234368 -1.566126
2   -0.567088 -2.593450 -3.874230 -4.584095 -4.187449 -3.151462 -1.742940
3    0.490473 -1.914407 -3.616364 -4.318823 -4.268016 -3.881110 -2.993280
4    0.800232 -0.874252 -2.384761 -3.973292 -4.338224 -3.802422 -2.534510
...
...   ...   ...   ...   ...   ...
4993  0.608558 -0.335651 -0.990948 -1.784153 -2.626145 -2.957065 -2.931897
4994 -2.060402 -2.860116 -3.405074 -3.748719 -3.513561 -3.006545 -2.234850
4995 -1.122969 -2.252925 -2.867628 -3.358605 -3.167849 -2.638360 -1.664162
4996 -0.547705 -1.889545 -2.839779 -3.457912 -3.929149 -3.966026 -3.492560
4997 -1.351779 -2.209006 -2.520225 -3.061475 -3.065141 -3.030739 -2.622720

          7      8      9   ...    131     132     133  \
0   -1.818286 -1.250522 -0.477492   ...  0.792168  0.933541  0.796958
1   -0.992258 -0.754680  0.042321   ...  0.538356  0.656881  0.787490
2   -1.490659 -1.183580 -0.394229   ...  0.886073  0.531452  0.311377
3   -1.671131 -1.333884 -0.965629   ...  0.350816  0.499111  0.600345
4   -1.783423 -1.594450 -0.753199   ...  1.148884  0.958434  1.059025
...
...   ...   ...   ...   ...
4993 -2.664816 -2.090137 -1.461841   ...  1.757705  2.291923  2.704595
4994 -1.593270 -1.075279 -0.976047   ...  1.388947  2.079675  2.433375
4995 -0.935655 -0.866953 -0.645363   ... -0.472419 -1.310147 -2.029521
4996 -2.695270 -1.849691 -1.374321   ...  1.258419  1.907530  2.280888
4997 -2.044092 -1.295874 -0.733839   ... -1.512234 -2.076075 -2.586042

         134     135     136     137     138     139    140
0    0.578621  0.257740  0.228077  0.123431  0.925286  0.193137  1.0
1    0.724046  0.555784  0.476333  0.773820  1.119621 -1.436250  1.0
2   -0.021919 -0.713683 -0.532197  0.321097  0.904227 -0.421797  1.0
3    0.842069  0.952074  0.990133  1.086798  1.403011 -0.383564  1.0
4    1.371682  1.277392  0.960304  0.971020  1.614392  1.421456  1.0
...
...   ...   ...   ...   ...
4993  2.451519  2.017396  1.704358  1.688542  1.629593  1.342651  0.0
4994  2.159484  1.819747  1.534767  1.696818  1.483832  1.047612  0.0
4995 -3.221294 -4.176790 -4.009720 -2.874136 -2.008369 -1.808334  0.0
4996  1.895242  1.437702  1.193433  1.261335  1.150449  0.804932  0.0
4997 -3.322799 -3.627311 -3.437038 -2.260023 -1.577823 -0.684531  0.0

```

[4998 rows x 141 columns]

0.1.3 Separating the feature and target column

[83]:

```
x = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

[84]:

```
x
```

```
[84]: array([[-0.11252183, -2.8272038 , -3.7738969 , ... , 0.12343082,
   0.92528624,  0.19313742],
[-1.1008778 , -3.9968398 , -4.2858426 , ... , 0.77381971,
  1.1196209 , -1.4362499 ],
[-0.56708802, -2.5934502 , -3.8742297 , ... , 0.32109663,
  0.90422673, -0.42179659],
...,
[-1.1229693 , -2.2529248 , -2.8676281 , ... , -2.874136 ,
 -2.0083694 , -1.8083338 ],
[-0.54770461, -1.8895451 , -2.8397786 , ... , 1.261335 ,
  1.1504486 ,  0.80493225],
[-1.3517791 , -2.2090058 , -2.5202247 , ... , -2.2600228 ,
 -1.577823 , -0.68453092]])
```

```
[85]: y
```

```
[85]: array([1., 1., 1., ... , 0., 0., 0.])
```

```
[86]: print(f"Normal samples: {(y==0).sum()}\nAbnormal samples: {(y==1).sum()}")
```

```
Normal samples: 2079,
Abnormal samples: 2919
```

0.1.4 Splitting normal and abnormal ECG samples

```
[87]: x_normal = x[y == 0]
x_abnormal = x[y == 1]
```

0.1.5 Splitting normal data into train and test

```
[88]: x_train,x_test_normal = train_test_split(x_normal, test_size=0.
                                         ↪2,random_state=42)
```

```
[89]: x_test = np.vstack([x_test_normal, x_abnormal])
y_test = np.hstack([
    np.zeros(len(x_test_normal)),
    np.ones(len(x_abnormal))
])
```

```
[90]: scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

```
[91]: print(f"Train shape: {x_train_scaled.shape}, Test shape: {x_test_scaled.shape}")
```

```
Train shape: (1663, 140), Test shape: (3335, 140)
```

0.1.6 Building the model

```
[92]: input_dim = x.shape[1]

autoencoder = Sequential([
    Dense(256, activation = 'relu', input_shape = (input_dim,)),
    Dense(128, activation = 'relu'),
    Dense(64, activation = 'relu'),
    Dense(32, activation = 'relu'),
    Dense(64, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(256, activation = 'relu'),
    Dense(input_dim, activation='linear')
])

autoencoder.summary()
```

```
c:\Users\TUFAN\.conda\envs\tf_env\lib\site-
packages\keras\src\layers\core\dense.py:93: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 256)	36,096
dense_23 (Dense)	(None, 128)	32,896
dense_24 (Dense)	(None, 64)	8,256
dense_25 (Dense)	(None, 32)	2,080
dense_26 (Dense)	(None, 64)	2,112
dense_27 (Dense)	(None, 128)	8,320
dense_28 (Dense)	(None, 256)	33,024
dense_29 (Dense)	(None, 140)	35,980

Total params: 158,764 (620.17 KB)

```
Trainable params: 158,764 (620.17 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[93]: autoencoder.compile(optimizer='adam', loss = 'mse', metrics=['mae'])
```


```
[94]: estop = EarlyStopping(monitor='val_loss', patience = 5, verbose = 1, u
      ↪restore_best_weights = True )
history = autoencoder.fit(
    x_train_scaled, x_train_scaled,
    epochs=50,
    batch_size=32,
    validation_split=0.1,
    callbacks = [estop],
    verbose=1
)
```

```
Epoch 1/50
47/47          2s 8ms/step - loss:
0.8639 - mae: 0.6321 - val_loss: 0.3930 - val_mae: 0.4205
Epoch 2/50
47/47          0s 4ms/step - loss:
0.4076 - mae: 0.4054 - val_loss: 0.2963 - val_mae: 0.3455
Epoch 3/50
47/47          0s 4ms/step - loss:
0.3001 - mae: 0.3370 - val_loss: 0.2339 - val_mae: 0.2954
Epoch 4/50
47/47          0s 5ms/step - loss:
0.2178 - mae: 0.2869 - val_loss: 0.2152 - val_mae: 0.2796
Epoch 5/50
47/47          0s 4ms/step - loss:
0.1815 - mae: 0.2631 - val_loss: 0.2054 - val_mae: 0.2663
Epoch 6/50
47/47          0s 3ms/step - loss:
0.1630 - mae: 0.2553 - val_loss: 0.1935 - val_mae: 0.2568
Epoch 7/50
47/47          0s 4ms/step - loss:
0.1763 - mae: 0.2542 - val_loss: 0.1825 - val_mae: 0.2538
Epoch 8/50
47/47          0s 4ms/step - loss:
0.1555 - mae: 0.2420 - val_loss: 0.1754 - val_mae: 0.2411
Epoch 9/50
47/47          0s 4ms/step - loss:
0.1339 - mae: 0.2251 - val_loss: 0.1718 - val_mae: 0.2370
Epoch 10/50
47/47          0s 4ms/step - loss:
0.1085 - mae: 0.2112 - val_loss: 0.1634 - val_mae: 0.2314
```

```
Epoch 11/50
47/47          0s 4ms/step - loss:
0.1132 - mae: 0.2127 - val_loss: 0.1733 - val_mae: 0.2378
Epoch 12/50
47/47          0s 4ms/step - loss:
0.1235 - mae: 0.2226 - val_loss: 0.1634 - val_mae: 0.2322
Epoch 13/50
47/47          0s 5ms/step - loss:
0.1017 - mae: 0.2078 - val_loss: 0.1535 - val_mae: 0.2197
Epoch 14/50
47/47          0s 6ms/step - loss:
0.1033 - mae: 0.2069 - val_loss: 0.1497 - val_mae: 0.2175
Epoch 15/50
47/47          0s 3ms/step - loss:
0.0908 - mae: 0.1975 - val_loss: 0.1524 - val_mae: 0.2219
Epoch 16/50
47/47          0s 4ms/step - loss:
0.0932 - mae: 0.2037 - val_loss: 0.1504 - val_mae: 0.2183
Epoch 17/50
47/47          0s 4ms/step - loss:
0.0817 - mae: 0.1921 - val_loss: 0.1484 - val_mae: 0.2172
Epoch 18/50
47/47          0s 4ms/step - loss:
0.0759 - mae: 0.1862 - val_loss: 0.1418 - val_mae: 0.2103
Epoch 19/50
47/47          0s 4ms/step - loss:
0.0825 - mae: 0.1926 - val_loss: 0.1404 - val_mae: 0.2089
Epoch 20/50
47/47          0s 4ms/step - loss:
0.0759 - mae: 0.1827 - val_loss: 0.1389 - val_mae: 0.2111
Epoch 21/50
47/47          0s 4ms/step - loss:
0.0771 - mae: 0.1878 - val_loss: 0.1433 - val_mae: 0.2133
Epoch 22/50
47/47          0s 3ms/step - loss:
0.0739 - mae: 0.1843 - val_loss: 0.1362 - val_mae: 0.2033
Epoch 23/50
47/47          0s 4ms/step - loss:
0.0649 - mae: 0.1744 - val_loss: 0.1413 - val_mae: 0.2097
Epoch 24/50
47/47          0s 3ms/step - loss:
0.0711 - mae: 0.1797 - val_loss: 0.1423 - val_mae: 0.2126
Epoch 25/50
47/47          0s 4ms/step - loss:
0.0695 - mae: 0.1794 - val_loss: 0.1420 - val_mae: 0.2087
Epoch 26/50
47/47          0s 4ms/step - loss:
0.0721 - mae: 0.1825 - val_loss: 0.1478 - val_mae: 0.2207
```

```
Epoch 27/50
47/47          0s 4ms/step - loss:
0.0702 - mae: 0.1814 - val_loss: 0.1435 - val_mae: 0.2097
Epoch 27: early stopping
Restoring model weights from the end of the best epoch: 22.
```

```
[95]: # Plot training history
plt.figure(figsize=(6,4))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Autoencoder Training Loss (Without Sparsity)')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.legend()
plt.show()
```



0.1.7 Anomaly detection using reconstruction error

```
[96]: def reconstruction_errors(model, data):
    reconstructed = model.predict(data, verbose=0)
    errors = np.mean((data - reconstructed) ** 2, axis=1)
    return errors
```

```
[97]: ## error on test test
errors = reconstruction_errors(autoencoder, x_test_scaled)

# Computing threshold using normal data
normal_errors = reconstruction_errors(autoencoder, x_test_scaled[y_test == 0])
threshold = np.percentile(normal_errors, 90)
print(f"Detection threshold: {threshold:.6f}")

## Flagging the anomalies
y_pred = (errors > threshold).astype(int)
```

Detection threshold: 0.104929


```
[98]: print("\nClassification Report (Without Sparsity):")
print(classification_report(y_test, y_pred))

print(f"ROC-AUC Score: {roc_auc_score(y_test, errors):.4f}")
```

	precision	recall	f1-score	support
0.0	0.98	0.90	0.94	416
1.0	0.99	1.00	0.99	2919
accuracy			0.98	3335
macro avg	0.98	0.95	0.96	3335
weighted avg	0.98	0.98	0.98	3335

ROC-AUC Score: 0.9480

```
[99]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix (Without Sparsity)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



0.2 Question 2: Include sparsity regularization and compare your results.

0.2.1 Building the sparse autoencoder model

```
[100]: sparse_autoencoder = Sequential([
    Dense(256, activation = 'relu', input_shape = (input_dim,), 
          activity_regularizer = regularizers.l1(1e-4)),
    Dense(128, activation = 'relu', activity_regularizer = regularizers.
          l1(1e-4)),
    Dense(64, activation = 'relu', activity_regularizer = regularizers.l1(1e-4)),
    Dense(32, activation = 'relu', activity_regularizer = regularizers.l1(1e-4)),
    Dense(64, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(256, activation = 'relu'),
    Dense(input_dim, activation='linear')
])

sparse_autoencoder.summary()
sparse_autoencoder.compile(optimizer='adam', loss='mse', metrics=['mae'])
```

```
c:\Users\TUFAN\.conda\envs\tf_env\lib\site-packages\keras\src\layers\core\dense.py:93: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 256)	36,096
dense_31 (Dense)	(None, 128)	32,896
dense_32 (Dense)	(None, 64)	8,256
dense_33 (Dense)	(None, 32)	2,080
dense_34 (Dense)	(None, 64)	2,112
dense_35 (Dense)	(None, 128)	8,320
dense_36 (Dense)	(None, 256)	33,024
dense_37 (Dense)	(None, 140)	35,980

```
Total params: 158,764 (620.17 KB)
```

```
Trainable params: 158,764 (620.17 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[101]: history_sparse = sparse_autoencoder.fit(  
        x_train_scaled, x_train_scaled,  
        epochs=50,  
        batch_size=32,  
        validation_split=0.1,  
        callbacks = [estop],  
        verbose=1  
)  
  
# Plot training history comparison  
plt.figure(figsize=(6,4))
```

```

plt.plot(history_sparse.history['loss'], label='Sparse AE Train Loss')
plt.plot(history_sparse.history['val_loss'], label='Sparse AE Val Loss')
plt.title('Sparse Autoencoder Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.legend()
plt.show()

```

Epoch 1/50
 47/47 2s 7ms/step - loss:
 1.2393 - mae: 0.6909 - val_loss: 0.7497 - val_mae: 0.5672

Epoch 2/50
 47/47 0s 4ms/step - loss:
 0.7218 - mae: 0.5470 - val_loss: 0.6087 - val_mae: 0.4791

Epoch 3/50
 47/47 0s 3ms/step - loss:
 0.5971 - mae: 0.4668 - val_loss: 0.5281 - val_mae: 0.4544

Epoch 4/50
 47/47 0s 4ms/step - loss:
 0.5515 - mae: 0.4494 - val_loss: 0.4991 - val_mae: 0.4394

Epoch 5/50
 47/47 0s 3ms/step - loss:
 0.5420 - mae: 0.4375 - val_loss: 0.4740 - val_mae: 0.4176

Epoch 6/50
 47/47 0s 5ms/step - loss:
 0.4928 - mae: 0.4206 - val_loss: 0.4433 - val_mae: 0.4008

Epoch 7/50
 47/47 0s 3ms/step - loss:
 0.4494 - mae: 0.3887 - val_loss: 0.4089 - val_mae: 0.3729

Epoch 8/50
 47/47 0s 4ms/step - loss:
 0.4296 - mae: 0.3732 - val_loss: 0.3830 - val_mae: 0.3585

Epoch 9/50
 47/47 0s 3ms/step - loss:
 0.4025 - mae: 0.3652 - val_loss: 0.3729 - val_mae: 0.3507

Epoch 10/50
 47/47 0s 6ms/step - loss:
 0.3508 - mae: 0.3437 - val_loss: 0.3712 - val_mae: 0.3451

Epoch 11/50
 47/47 0s 4ms/step - loss:
 0.3764 - mae: 0.3429 - val_loss: 0.3593 - val_mae: 0.3397


Epoch 12/50
 47/47 0s 3ms/step - loss:
 0.3672 - mae: 0.3407 - val_loss: 0.3559 - val_mae: 0.3355

Epoch 13/50
 47/47 0s 4ms/step - loss:
 0.3309 - mae: 0.3244 - val_loss: 0.3342 - val_mae: 0.3296

Epoch 14/50

```
47/47          0s 4ms/step - loss:  
0.3097 - mae: 0.3209 - val_loss: 0.3507 - val_mae: 0.3377  
Epoch 15/50  
47/47          0s 3ms/step - loss:  
0.3028 - mae: 0.3227 - val_loss: 0.3422 - val_mae: 0.3306  
Epoch 16/50  
47/47          0s 4ms/step - loss:  
0.2948 - mae: 0.3175 - val_loss: 0.3292 - val_mae: 0.3247  
Epoch 17/50  
47/47          0s 3ms/step - loss:  
0.3019 - mae: 0.3177 - val_loss: 0.3462 - val_mae: 0.3338  
Epoch 18/50  
47/47          0s 3ms/step - loss:  
0.3101 - mae: 0.3217 - val_loss: 0.3093 - val_mae: 0.3134  
Epoch 19/50  
47/47          0s 5ms/step - loss:  
0.2767 - mae: 0.3079 - val_loss: 0.3242 - val_mae: 0.3162  
Epoch 20/50  
47/47          0s 4ms/step - loss:  
0.2804 - mae: 0.2997 - val_loss: 0.2964 - val_mae: 0.3101  
Epoch 21/50  
47/47          0s 4ms/step - loss:  
0.2342 - mae: 0.2929 - val_loss: 0.3292 - val_mae: 0.3179  
Epoch 22/50  
47/47          0s 4ms/step - loss:  
0.2467 - mae: 0.2897 - val_loss: 0.2917 - val_mae: 0.3025  
Epoch 23/50  
47/47          0s 3ms/step - loss:  
0.2239 - mae: 0.2811 - val_loss: 0.2919 - val_mae: 0.3066  
Epoch 24/50  
47/47          0s 4ms/step - loss:  
0.2503 - mae: 0.2900 - val_loss: 0.2831 - val_mae: 0.2956  
Epoch 25/50  
47/47          0s 3ms/step - loss:  
0.2252 - mae: 0.2788 - val_loss: 0.2957 - val_mae: 0.3025  
Epoch 26/50  
47/47          0s 4ms/step - loss:  
0.2242 - mae: 0.2768 - val_loss: 0.2798 - val_mae: 0.2959  
Epoch 27/50  
47/47          0s 4ms/step - loss:  
0.2072 - mae: 0.2720 - val_loss: 0.2699 - val_mae: 0.2895  
Epoch 28/50  
47/47          0s 3ms/step - loss:  
0.1985 - mae: 0.2631 - val_loss: 0.2821 - val_mae: 0.2943  
Epoch 29/50  
47/47          0s 4ms/step - loss:  
0.2140 - mae: 0.2714 - val_loss: 0.2678 - val_mae: 0.2838  
Epoch 30/50
```

```
47/47          0s 3ms/step - loss:  
0.1952 - mae: 0.2611 - val_loss: 0.2744 - val_mae: 0.2861  
Epoch 31/50  
47/47          0s 4ms/step - loss:  
0.2044 - mae: 0.2669 - val_loss: 0.2658 - val_mae: 0.2816  
Epoch 32/50  
47/47          0s 3ms/step - loss:  
0.1888 - mae: 0.2545 - val_loss: 0.2634 - val_mae: 0.2826  
Epoch 33/50  
47/47          0s 4ms/step - loss:  
0.1935 - mae: 0.2542 - val_loss: 0.2565 - val_mae: 0.2803  
Epoch 34/50  
47/47          0s 3ms/step - loss:  
0.1930 - mae: 0.2575 - val_loss: 0.2626 - val_mae: 0.2842  
Epoch 35/50  
47/47          0s 3ms/step - loss:  
0.1817 - mae: 0.2495 - val_loss: 0.2570 - val_mae: 0.2848  
Epoch 36/50  
47/47          0s 4ms/step - loss:  
0.1787 - mae: 0.2504 - val_loss: 0.2584 - val_mae: 0.2765  
Epoch 37/50  
47/47          0s 4ms/step - loss:  
0.1797 - mae: 0.2509 - val_loss: 0.2621 - val_mae: 0.2856  
Epoch 38/50  
47/47          0s 5ms/step - loss:  
0.1809 - mae: 0.2511 - val_loss: 0.2611 - val_mae: 0.2800  
Epoch 38: early stopping  
Restoring model weights from the end of the best epoch: 33.
```



0.2.2 Evaluating Sparse Autoencoder

```
[103]: sparse_errors = reconstruction_errors(sparse_autoencoder, x_test_scaled)

# new threshold
normal_sparse_errors = reconstruction_errors(sparse_autoencoder, ↴
    ↪x_test_scaled[y_test == 0])
threshold_sparse = np.percentile(normal_sparse_errors, 90)
print(f"Sparse AE Threshold: {threshold_sparse:.6f}")

# Predictions
y_pred_sparse = (sparse_errors > threshold_sparse).astype(int)

print("\nClassification Report (With Sparsity):")
print(classification_report(y_test, y_pred_sparse))
print(f"ROC-AUC Score (Sparse AE): {roc_auc_score(y_test, sparse_errors):.4f}")

# Confusion Matrix
cm_sparse = confusion_matrix(y_test, y_pred_sparse)
sns.heatmap(cm_sparse, annot=True, fmt='d', cmap='Greens')
plt.title("Confusion Matrix (With Sparsity)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
plt.show()
```


Sparse AE Threshold: 0.259583

Classification Report (With Sparsity):

	precision	recall	f1-score	support
0.0	0.98	0.90	0.94	416
1.0	0.99	1.00	0.99	2919
accuracy			0.99	3335
macro avg	0.98	0.95	0.96	3335
weighted avg	0.98	0.99	0.98	3335

ROC-AUC Score (Sparse AE): 0.9622

Confusion Matrix (With Sparsity)



0.3 Comparison:

Model Type	Accuracy	ROC-AUC
Standard Autoencoder	98%	0.948
Sparse Autoencoder	99%	0.962

- Incorporating sparsity regularization significantly enhanced the model's ability to distinguish between normal and abnormal ECG signals by promoting compact and discriminative feature learning.