

Experiment_7

September 14, 2025

Name:	Tufan Kundu
Registration no:	24MDT0184
Course Name:	Deep Learning Lab
Course Code:	PMDS603P
Experiment:	7
Date:	11 September,2025

0.1 Question 1: First try to fit a simple ANN model with atleast 2 hidden layers and get the accuracy of your model. So that you can compare how the CNN models you are going to fit next will work. Include early stopping in your model as well

0.1.1 Importing necessary libraries

```
[1]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping, ReduceLROnPlateau
from sklearn.metrics import accuracy_score
import numpy as np

import warnings
warnings.filterwarnings('ignore')
```


0.1.2 Loading the MNIST dataset

```
[2]: from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
```

0.1.3 Displaying random images

```
[3]: class_names = [str(i) for i in range(10)]
indices = np.random.choice(len(x_train), size = 5 , replace = False)

plt.figure(figsize = (2*5,3))
for i, idx in enumerate(indices):
    ax = plt.subplot(1,5,i+1)
    img = x_train[idx]
    plt.imshow(img, cmap = 'gray')
    label = class_names[y_train[idx]]
    plt.title(label)
    plt.axis('off')
plt.show()
```



0.1.4 Making the ANN

```
[5]: model = Sequential()
model.add(Flatten(input_shape = (28,28)))
model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.3))
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation = 'softmax'))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None , 784)	0
dense (Dense)	(None , 256)	200,960
dropout (Dropout)	(None , 256)	0

dense_1 (Dense)	(None, 128)	32,896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1,290

Total params: 235,146 (918.54 KB)

Trainable params: 235,146 (918.54 KB)

Non-trainable params: 0 (0.00 B)

```
[ ]: estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
    ↪ verbose = 1, restore_best_weights=True)
model.compile(loss = 'sparse_categorical_crossentropy', optimizer = 'adam',
    ↪ metrics = ['accuracy'])

[7]: history = model.fit(x_train,y_train,batch_size=128, epochs = 200, verbose = 1,
    ↪ validation_data=(x_test,y_test), callbacks=[estop])
```

```
Epoch 1/200
469/469      3s 5ms/step -
accuracy: 0.8054 - loss: 0.6438 - val_accuracy: 0.9578 - val_loss: 0.1405
Epoch 2/200
469/469      2s 4ms/step -
accuracy: 0.9514 - loss: 0.1642 - val_accuracy: 0.9717 - val_loss: 0.0967
Epoch 3/200
469/469      2s 4ms/step -
accuracy: 0.9643 - loss: 0.1163 - val_accuracy: 0.9759 - val_loss: 0.0787
Epoch 4/200
469/469      2s 4ms/step -
accuracy: 0.9721 - loss: 0.0927 - val_accuracy: 0.9751 - val_loss: 0.0779
Epoch 5/200
469/469      2s 4ms/step -
accuracy: 0.9742 - loss: 0.0807 - val_accuracy: 0.9778 - val_loss: 0.0725
Epoch 6/200
469/469      2s 4ms/step -
accuracy: 0.9792 - loss: 0.0688 - val_accuracy: 0.9811 - val_loss: 0.0663
Epoch 7/200
469/469      2s 4ms/step -
accuracy: 0.9808 - loss: 0.0611 - val_accuracy: 0.9800 - val_loss: 0.0658
Epoch 8/200
469/469      2s 4ms/step -
accuracy: 0.9809 - loss: 0.0588 - val_accuracy: 0.9809 - val_loss: 0.0625
Epoch 9/200
```

```

469/469          2s 4ms/step -
accuracy: 0.9841 - loss: 0.0523 - val_accuracy: 0.9819 - val_loss: 0.0596
Epoch 10/200
469/469          2s 4ms/step -
accuracy: 0.9859 - loss: 0.0445 - val_accuracy: 0.9808 - val_loss: 0.0645
Epoch 11/200
469/469          2s 5ms/step -
accuracy: 0.9848 - loss: 0.0456 - val_accuracy: 0.9824 - val_loss: 0.0612
Epoch 12/200
469/469          2s 5ms/step -
accuracy: 0.9873 - loss: 0.0391 - val_accuracy: 0.9814 - val_loss: 0.0635
Epoch 13/200
469/469          2s 5ms/step -
accuracy: 0.9873 - loss: 0.0367 - val_accuracy: 0.9833 - val_loss: 0.0586
Epoch 14/200
469/469          2s 5ms/step -
accuracy: 0.9888 - loss: 0.0346 - val_accuracy: 0.9809 - val_loss: 0.0692
Epoch 15/200
469/469          2s 5ms/step -
accuracy: 0.9895 - loss: 0.0305 - val_accuracy: 0.9825 - val_loss: 0.0679
Epoch 16/200
469/469          2s 5ms/step -
accuracy: 0.9885 - loss: 0.0333 - val_accuracy: 0.9833 - val_loss: 0.0580
Epoch 17/200
469/469          2s 5ms/step -
accuracy: 0.9902 - loss: 0.0292 - val_accuracy: 0.9823 - val_loss: 0.0631
Epoch 18/200
469/469          2s 5ms/step -
accuracy: 0.9897 - loss: 0.0293 - val_accuracy: 0.9818 - val_loss: 0.0616
Epoch 19/200
469/469          2s 5ms/step -
accuracy: 0.9908 - loss: 0.0274 - val_accuracy: 0.9826 - val_loss: 0.0674
Epoch 20/200
469/469          2s 5ms/step -
accuracy: 0.9900 - loss: 0.0296 - val_accuracy: 0.9817 - val_loss: 0.0649
Epoch 21/200
469/469          2s 5ms/step -
accuracy: 0.9912 - loss: 0.0246 - val_accuracy: 0.9821 - val_loss: 0.0705
Epoch 21: early stopping
Restoring model weights from the end of the best epoch: 16.

```

```
[8]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```


313/313          1s 2ms/step -
accuracy: 0.9795 - loss: 0.0711

```

```
[9]: print(f"Validation Accuracy with ANN model: {val_accuracy*100:.4f}%")
```

```
Validation Accuracy with ANN model: 98.3300%
```

```
[10]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epochs")
plt.ylabel('Loss')
plt.legend()
plt.show()
```



0.2 Question 2: Now let us try to fit a CNN model to accomplish the same task and look at the improvements. Import necessary modules and functions first. Here you can see we are including the Conv2D and MaxPool2d layers from keras.layers.

```
[12]: from tensorflow import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Flatten, Conv2D, MaxPool2D
from keras.optimizers import Adam, RMSprop
from keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split

[ ]: (x_train,y_train),(x_test,y_test) = mnist.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
x_train = x_train.reshape(-1,28,28,1)
```

```
x_test = x_test.reshape(-1,28,28,1)
```

0.2.1 Creating the CNN model

```
[15]: model = Sequential()
model.add(Conv2D(26,5, strides = (1,1), activation = 'relu', padding='valid',
↳input_shape = (28,28,1)))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2), padding='valid'))
model.add(Conv2D(20,3, strides = (1,1), activation = 'relu', padding='valid'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='valid'))
model.add(Conv2D(10,3, strides = (1,1), activation = 'relu', padding='valid'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='valid'))
model.add(Flatten())
model.add(Dense(100, activation = 'relu'))
model.add(Dense(10, activation='softmax'))
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 24, 24, 26)	676
max_pooling2d (MaxPooling2D)	(None, 12, 12, 26)	0
conv2d_1 (Conv2D)	(None, 10, 10, 20)	4,700
max_pooling2d_1 (MaxPooling2D)	(None, 9, 9, 20)	0
conv2d_2 (Conv2D)	(None, 7, 7, 10)	1,810
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 10)	0
flatten_1 (Flatten)	(None, 360)	0
dense_3 (Dense)	(None, 100)	36,100
dense_4 (Dense)	(None, 10)	1,010

Total params: 44,296 (173.03 KB)

Trainable params: 44,296 (173.03 KB)

Non-trainable params: 0 (0.00 B)

```
[16]: estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
    ↪ verbose = 1, restore_best_weights=True)
model.compile(loss = 'sparse_categorical_crossentropy', optimizer = 'rmsprop',
    ↪ metrics = ['accuracy'])
```

```
[17]: history = model.fit(x_train,y_train,batch_size=128, epochs = 200, verbose = 1,
    ↪ validation_data=(x_test,y_test), callbacks=[estop])
```

Epoch 1/200

469/469 6s 12ms/step -

accuracy: 0.8025 - loss: 0.6023 - val_accuracy: 0.9761 - val_loss: 0.0748

Epoch 2/200

469/469 6s 13ms/step -

accuracy: 0.9697 - loss: 0.0947 - val_accuracy: 0.9678 - val_loss: 0.0981

Epoch 3/200

469/469 6s 14ms/step -

accuracy: 0.9818 - loss: 0.0587 - val_accuracy: 0.9877 - val_loss: 0.0358

Epoch 4/200

469/469 7s 15ms/step -

accuracy: 0.9856 - loss: 0.0443 - val_accuracy: 0.9886 - val_loss: 0.0380

Epoch 5/200

469/469 7s 15ms/step -

accuracy: 0.9883 - loss: 0.0376 - val_accuracy: 0.9828 - val_loss: 0.0541

Epoch 6/200

469/469 6s 14ms/step -

accuracy: 0.9904 - loss: 0.0305 - val_accuracy: 0.9905 - val_loss: 0.0334

Epoch 7/200

469/469 6s 14ms/step -

accuracy: 0.9905 - loss: 0.0292 - val_accuracy: 0.9907 - val_loss: 0.0310

Epoch 8/200

469/469 7s 14ms/step -

accuracy: 0.9932 - loss: 0.0233 - val_accuracy: 0.9863 - val_loss: 0.0450

Epoch 9/200

469/469 6s 13ms/step -

accuracy: 0.9933 - loss: 0.0216 - val_accuracy: 0.9922 - val_loss: 0.0277

Epoch 10/200

469/469 6s 14ms/step -

accuracy: 0.9936 - loss: 0.0203 - val_accuracy: 0.9916 - val_loss: 0.0297

Epoch 11/200

469/469 6s 13ms/step -

accuracy: 0.9941 - loss: 0.0177 - val_accuracy: 0.9917 - val_loss: 0.0261

Epoch 12/200

469/469 7s 14ms/step -

accuracy: 0.9956 - loss: 0.0141 - val_accuracy: 0.9920 - val_loss: 0.0287

Epoch 13/200

```

469/469          6s 13ms/step -
accuracy: 0.9955 - loss: 0.0146 - val_accuracy: 0.9925 - val_loss: 0.0309
Epoch 14/200
469/469          6s 13ms/step -
accuracy: 0.9964 - loss: 0.0120 - val_accuracy: 0.9908 - val_loss: 0.0356
Epoch 15/200
469/469          7s 15ms/step -
accuracy: 0.9963 - loss: 0.0122 - val_accuracy: 0.9893 - val_loss: 0.0424
Epoch 16/200
469/469          6s 13ms/step -
accuracy: 0.9970 - loss: 0.0097 - val_accuracy: 0.9916 - val_loss: 0.0351
Epoch 16: early stopping
Restoring model weights from the end of the best epoch: 11.

```

```
[18]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```


313/313          1s 3ms/step -
accuracy: 0.9898 - loss: 0.0320

```

```
[19]: print(f"Validation Accuracy with CNN model: {val_accuracy*100:.4f}%")
```

```
Validation Accuracy with CNN model: 99.1700%
```

```
[20]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epochs")
plt.ylabel('Loss')
plt.legend()
plt.show()
```

0.3 Question 3: Now you can try to fit a CNN model with CIFAR10 dataset we have seen in the previous lab with an appropriate model. See whether you get an improved accuracy for the model when you are using CNN models

```
[21]: (x_train,y_train),(x_test,y_test) = keras.datasets.cifar10.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
```

0.3.1 Displaying random images from the dataset

```
[22]: class_names = ["Airplane","Automobile","Bird","Cat","Deer",
"Dog","Frog","Horse","Ship","Truck"]
indices = np.random.choice(len(x_train), size = 5 , replace = False)
plt.figure(figsize = (2*5,3))
for i, idx in enumerate(indices):
    ax = plt.subplot(1,5,i+1)
    img = x_train[idx]
    plt.imshow(img)
    label = class_names[y_train[idx][0]]
    plt.title(label)
    plt.axis('off')
plt.show()
```



0.3.2 Building the CNN model

```
[24]: model = Sequential()
model.add(Conv2D(32,5, strides = (1,1), activation = 'relu', padding='same',
    ↪input_shape = x_train.shape[1:]))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2), padding='valid'))
model.add(Dropout(0.5))
model.add(Conv2D(32,3, strides = (1,1), activation = 'relu', padding='same'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='same'))
model.add(Dropout(0.3))
model.add(Conv2D(64,3, strides = (1,1), activation = 'relu', padding='same'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='valid'))
model.add(Dropout(0.2))
model.add(Flatten())
model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 32, 32, 32)	2,432
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_2 (Dropout)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 32)	9,248
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_3 (Dropout)	(None, 16, 16, 32)	0
conv2d_5 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 15, 15, 64)	0
dropout_4 (Dropout)	(None, 15, 15, 64)	0
flatten_2 (Flatten)	(None, 14400)	0
dense_5 (Dense)	(None, 256)	3,686,656

dropout_5 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32,896
dropout_6 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 10)	1,290

Total params: 3,751,018 (14.31 MB)

Trainable params: 3,751,018 (14.31 MB)

Non-trainable params: 0 (0.00 B)

```
[25]: estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
    ↪ verbose = 1, restore_best_weights=True)
model.compile(loss = 'sparse_categorical_crossentropy', optimizer = 'adam',
    ↪ metrics = ['accuracy'])
```

```
[26]: history = model.fit(x_train,y_train,batch_size=128, epochs = 200, verbose = 1,
    ↪ validation_data=(x_test,y_test), callbacks=[estop])
```

Epoch 1/200

391/391 42s 102ms/step -

accuracy: 0.2432 - loss: 2.0091 - val_accuracy: 0.4601 - val_loss: 1.5177

Epoch 2/200

391/391 43s 110ms/step -

accuracy: 0.4584 - loss: 1.4745 - val_accuracy: 0.5367 - val_loss: 1.2985

Epoch 3/200

391/391 36s 93ms/step -

accuracy: 0.5290 - loss: 1.3127 - val_accuracy: 0.5738 - val_loss: 1.1988

Epoch 4/200

391/391 36s 92ms/step -

accuracy: 0.5711 - loss: 1.2005 - val_accuracy: 0.5973 - val_loss: 1.1224

Epoch 5/200

391/391 62s 158ms/step -

accuracy: 0.6028 - loss: 1.1177 - val_accuracy: 0.5804 - val_loss: 1.1996

Epoch 6/200

391/391 70s 179ms/step -

accuracy: 0.6270 - loss: 1.0588 - val_accuracy: 0.6431 - val_loss: 1.0033

Epoch 7/200

391/391 61s 155ms/step -

accuracy: 0.6525 - loss: 0.9932 - val_accuracy: 0.6644 - val_loss: 0.9693

Epoch 8/200

391/391 46s 117ms/step -
 accuracy: 0.6660 - loss: 0.9487 - val_accuracy: 0.6826 - val_loss: 0.8955
 Epoch 9/200
 391/391 47s 121ms/step -
 accuracy: 0.6765 - loss: 0.9172 - val_accuracy: 0.6811 - val_loss: 0.9246
 Epoch 10/200
 391/391 41s 105ms/step -
 accuracy: 0.6872 - loss: 0.8855 - val_accuracy: 0.6827 - val_loss: 0.9187
 Epoch 11/200
 391/391 45s 114ms/step -
 accuracy: 0.6980 - loss: 0.8572 - val_accuracy: 0.6887 - val_loss: 0.8899
 Epoch 12/200
 391/391 60s 154ms/step -
 accuracy: 0.7095 - loss: 0.8182 - val_accuracy: 0.6826 - val_loss: 0.9365
 Epoch 13/200
 391/391 58s 149ms/step -
 accuracy: 0.7153 - loss: 0.8114 - val_accuracy: 0.7134 - val_loss: 0.8171
 Epoch 14/200
 391/391 55s 141ms/step -
 accuracy: 0.7240 - loss: 0.7793 - val_accuracy: 0.6892 - val_loss: 0.9158
 Epoch 15/200
 391/391 47s 121ms/step -
 accuracy: 0.7265 - loss: 0.7747 - val_accuracy: 0.6833 - val_loss: 0.9396
 Epoch 16/200
 391/391 43s 110ms/step -
 accuracy: 0.7339 - loss: 0.7522 - val_accuracy: 0.7189 - val_loss: 0.8089
 Epoch 17/200
 391/391 43s 109ms/step -
 accuracy: 0.7401 - loss: 0.7380 - val_accuracy: 0.7099 - val_loss: 0.8359
 Epoch 18/200
 391/391 43s 111ms/step -
 accuracy: 0.7448 - loss: 0.7270 - val_accuracy: 0.7197 - val_loss: 0.8131
 Epoch 19/200
 391/391 45s 116ms/step -
 accuracy: 0.7515 - loss: 0.6976 - val_accuracy: 0.7342 - val_loss: 0.7767
 Epoch 20/200
 391/391 52s 133ms/step -
 accuracy: 0.7533 - loss: 0.7006 - val_accuracy: 0.6721 - val_loss: 1.0006
 Epoch 21/200
 391/391 45s 116ms/step -
 accuracy: 0.7540 - loss: 0.6898 - val_accuracy: 0.7278 - val_loss: 0.7985
 Epoch 22/200
 391/391 51s 129ms/step -
 accuracy: 0.7617 - loss: 0.6690 - val_accuracy: 0.7322 - val_loss: 0.7815
 Epoch 23/200
 391/391 46s 119ms/step -
 accuracy: 0.7739 - loss: 0.6405 - val_accuracy: 0.7308 - val_loss: 0.7874
 Epoch 24/200

```
391/391          52s 133ms/step -  
accuracy: 0.7718 - loss: 0.6396 - val_accuracy: 0.7250 - val_loss: 0.8232  
Epoch 24: early stopping  
Restoring model weights from the end of the best epoch: 19.
```

```
[27]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```
313/313          4s 13ms/step -  
accuracy: 0.7305 - loss: 0.7857
```

```
[28]: print(f"Validation Accuracy of Cifar10 dataset with CNN model:␣  
      ↪{val_accuracy*100:.4f}%")
```

Validation Accuracy of Cifar10 dataset with CNN model: 73.4200%

```
[29]: plt.figure(figsize=(12,4))  
plt.subplot(1,2,1)  
plt.plot(history.history['accuracy'], label='Train Acc')  
plt.plot(history.history['val_accuracy'], label='Val Acc')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.subplot(1,2,2)  
plt.plot(history.history['loss'], label='Train Loss')  
plt.plot(history.history['val_loss'], label='Val Loss')  
plt.xlabel("Epochs")  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```

