

# experiment-8

September 22, 2025

Name:	Tufan Kundu
Registration no:	24MDT0184
Course Name:	Deep Learning Lab
Course Code:	PMDS603P
Experiment:	8
Date:	18 September,2025

**0.1 Question 1:** First try to fit a basic CNN model to accomplish the same task. Import necessary modules and classes first. Load the CIFAR 10 dataset and normalize the data and reshape. Now you can check the shape of xtrain and xtest etc which you are loading. we need to provide  $32 \times 32 \times 3$  images to the CNN. Next compile the model and fit the model and check the performance. Also include early stopping in your model. Since we have a classification problem we have to use the softmax activation in the final layer.

## 0.1.1 Importing the necessary libraries

```
[1]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout, Conv2D, MaxPool2D
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping
import numpy as np

import warnings
warnings.filterwarnings('ignore')
```

```
2025-09-22 07:01:41.466626: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
```

```
E0000 00:00:1758524501.799542      19 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1758524501.892278      19 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
```

### 0.1.2 Loading and preprocessing the dataset

```
[2]: (x_train,y_train),(x_test,y_test) = cifar10.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071      2s
0us/step
```

### 0.1.3 Displaying random images from the dataset

```
[3]: class_names = ["Airplane","Automobile","Bird","Cat","Deer",
"Dog","Frog","Horse","Ship","Truck"]

indices = np.random.choice(len(x_train), size = 5, replace = False)
plt.figure(figsize = (2*5,3))
for i,idx in enumerate(indices):
    ax = plt.subplot(1,5,i+1)
    img = x_train[idx]
    plt.imshow(img)
    label = class_names[y_train[idx][0]]
    plt.title(label)
    plt.axis('off')
plt.show()
```



### 0.1.4 Building the CNN

```
[4]: model = Sequential()
model.add(Conv2D(32,5, strides = (1,1), activation = 'relu', padding='same',
    ↪input_shape = x_train.shape[1:]))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2), padding='valid'))
model.add(Dropout(0.5))
model.add(Conv2D(32,3, strides = (1,1), activation = 'relu', padding='same'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='same'))
model.add(Dropout(0.3))
model.add(Conv2D(64,3, strides = (1,1), activation = 'relu', padding='same'))
model.add(MaxPool2D(pool_size=(2,2), strides=(1,1), padding='valid'))
model.add(Dropout(0.2))
model.add(Flatten())
model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
model.summary()
```

```
I0000 00:00:1758524526.043863      19 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 13942 MB memory: -> device:
0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5
I0000 00:00:1758524526.044545      19 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:1 with 13942 MB memory: -> device:
1, name: Tesla T4, pci bus id: 0000:00:05.0, compute capability: 7.5
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	2,432
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_1 (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 64)	0

dropout_2 (Dropout)	(None, 15, 15, 64)	0
flatten (Flatten)	(None, 14400)	0
dense (Dense)	(None, 256)	3,686,656
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32,896
dropout_4 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1,290

Total params: 3,751,018 (14.31 MB)

Trainable params: 3,751,018 (14.31 MB)

Non-trainable params: 0 (0.00 B)

```
[5]: estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
    ↪ verbose = 1, restore_best_weights=True)
model.compile(loss = 'sparse_categorical_crossentropy', optimizer = 'adam',
    ↪ metrics = ['accuracy'])
```

```
[6]: history = model.fit(x_train,y_train,batch_size=128, epochs = 200, verbose = 1,
    ↪ validation_data=(x_test,y_test), callbacks=[estop])
```

Epoch 1/200

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1758524532.675782 60 service.cc:148] XLA service 0x79ade000e340 initialized for platform CUDA (this does not guarantee that XLA will be used).

Devices:

I0000 00:00:1758524532.677284 60 service.cc:156] StreamExecutor device (0): Tesla T4, Compute Capability 7.5

I0000 00:00:1758524532.677305 60 service.cc:156] StreamExecutor device (1): Tesla T4, Compute Capability 7.5

I0000 00:00:1758524533.089177 60 cuda\_dnn.cc:529] Loaded cuDNN version 90300

17/391 3s 11ms/step - accuracy: 0.0992 - loss: 2.5474

I0000 00:00:1758524538.780024 60 device\_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

```
391/391          18s 24ms/step -  
accuracy: 0.1974 - loss: 2.1412 - val_accuracy: 0.4352 - val_loss: 1.5567  
Epoch 2/200  
391/391          3s 8ms/step -  
accuracy: 0.4373 - loss: 1.5330 - val_accuracy: 0.5071 - val_loss: 1.3733  
Epoch 3/200  
391/391          3s 8ms/step -  
accuracy: 0.5047 - loss: 1.3702 - val_accuracy: 0.5463 - val_loss: 1.2609  
Epoch 4/200  
391/391          3s 8ms/step -  
accuracy: 0.5592 - loss: 1.2411 - val_accuracy: 0.5457 - val_loss: 1.3063  
Epoch 5/200  
391/391          3s 8ms/step -  
accuracy: 0.5835 - loss: 1.1716 - val_accuracy: 0.5649 - val_loss: 1.2374  
Epoch 6/200  
391/391          3s 8ms/step -  
accuracy: 0.6104 - loss: 1.1073 - val_accuracy: 0.5987 - val_loss: 1.1134  
Epoch 7/200  
391/391          3s 8ms/step -  
accuracy: 0.6280 - loss: 1.0520 - val_accuracy: 0.6404 - val_loss: 1.0358  
Epoch 8/200  
391/391          3s 8ms/step -  
accuracy: 0.6495 - loss: 0.9930 - val_accuracy: 0.6511 - val_loss: 0.9789  
Epoch 9/200  
391/391          3s 8ms/step -  
accuracy: 0.6595 - loss: 0.9636 - val_accuracy: 0.6597 - val_loss: 0.9773  
Epoch 10/200  
391/391          3s 8ms/step -  
accuracy: 0.6672 - loss: 0.9446 - val_accuracy: 0.6658 - val_loss: 0.9494  
Epoch 11/200  
391/391          3s 8ms/step -  
accuracy: 0.6867 - loss: 0.8915 - val_accuracy: 0.6867 - val_loss: 0.8908  
Epoch 12/200  
391/391          3s 8ms/step -  
accuracy: 0.6922 - loss: 0.8735 - val_accuracy: 0.6938 - val_loss: 0.8763  
Epoch 13/200  
391/391          3s 8ms/step -  
accuracy: 0.6981 - loss: 0.8566 - val_accuracy: 0.6957 - val_loss: 0.8615  
Epoch 14/200  
391/391          3s 8ms/step -  
accuracy: 0.7149 - loss: 0.8138 - val_accuracy: 0.6922 - val_loss: 0.8784  
Epoch 15/200  
391/391          3s 8ms/step -  
accuracy: 0.7176 - loss: 0.8000 - val_accuracy: 0.7113 - val_loss: 0.8355  
Epoch 16/200
```

```

391/391          3s 8ms/step -
accuracy: 0.7230 - loss: 0.7870 - val_accuracy: 0.6967 - val_loss: 0.8705
Epoch 17/200
391/391          3s 8ms/step -
accuracy: 0.7265 - loss: 0.7724 - val_accuracy: 0.7134 - val_loss: 0.8180
Epoch 18/200
391/391          3s 8ms/step -
accuracy: 0.7335 - loss: 0.7496 - val_accuracy: 0.7195 - val_loss: 0.7988
Epoch 19/200
391/391          3s 8ms/step -
accuracy: 0.7381 - loss: 0.7398 - val_accuracy: 0.7240 - val_loss: 0.7961
Epoch 20/200
391/391          3s 8ms/step -
accuracy: 0.7428 - loss: 0.7283 - val_accuracy: 0.7221 - val_loss: 0.7998
Epoch 21/200
391/391          3s 8ms/step -
accuracy: 0.7511 - loss: 0.7032 - val_accuracy: 0.7244 - val_loss: 0.7977
Epoch 22/200
391/391          3s 8ms/step -
accuracy: 0.7545 - loss: 0.6965 - val_accuracy: 0.7075 - val_loss: 0.8427
Epoch 23/200
391/391          3s 8ms/step -
accuracy: 0.7563 - loss: 0.6949 - val_accuracy: 0.7232 - val_loss: 0.7966
Epoch 24/200
391/391          3s 8ms/step -
accuracy: 0.7643 - loss: 0.6695 - val_accuracy: 0.7100 - val_loss: 0.8499
Epoch 24: early stopping
Restoring model weights from the end of the best epoch: 19.

```

```
[7]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```

313/313          1s 2ms/step -
accuracy: 0.7251 - loss: 0.7901


```

```
[8]: print(f"Validation Accuracy of Cifar10 dataset with CNN model:␣
      ↳{val_accuracy*100:.4f}%")
```

```
Validation Accuracy of Cifar10 dataset with CNN model: 72.4000%
```

```
[9]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
```

```
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epochs")
plt.ylabel('Loss')
plt.legend()
plt.show()
```



**0.2 Question 2:** Next we will see how you can define a new model incorporating the pre-trained model VGG16 in your model as its part. For that first you have to import the respective model details along with the necessary modules and classes.

### 0.2.1 Importing the necessary libraries

```
[10]: import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Resizing, Dropout
from keras.applications import VGG16
from keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split
from keras.layers import GlobalAveragePooling2D
```

### 0.2.2 Loading and preprocessing of the dataset

```
[11]: (x_train,y_train),(x_test,y_test) = cifar10.load_data()
x_train = x_train/255.0
x_test = x_test/255.0
```

```
[12]: base_model = VGG16(include_top = False, weights = 'imagenet', input_shape = (224,224,3))
for layer in base_model.layers:
    layer.trainable = False
```

```

## Building the model
model = Sequential()
model.add(Resizing(224,224, input_shape = (32,32,3)))
model.add(base_model)
model.add(GlobalAveragePooling2D())
model.add(Dense(100, activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(50, activation = 'relu'))
model.add(Dropout(0.3))
model.add(Dense(10, activation = 'softmax'))

model.summary()
model.compile(
    loss = 'sparse_categorical_crossentropy',
    optimizer = 'adam',
    metrics = ['accuracy']
)

estop = EarlyStopping(monitor = 'val_loss', min_delta= 1e-4, patience= 5,
↳ verbose = 1, restore_best_weights=True)
history = model.fit(x_train,y_train,batch_size=128, epochs = 200, verbose = 1,
↳ validation_data=(x_test,y_test), callbacks=[estop])

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
 58889256/58889256                      0s  
 0us/step

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
resizing ( <a href="#">Resizing</a> )	( <a href="#">None</a> , 224, 224, 3)	0
vgg16 ( <a href="#">Functional</a> )	( <a href="#">None</a> , 7, 7, 512)	14,714,688
global_average_pooling2d ( <a href="#">GlobalAveragePooling2D</a> )	( <a href="#">None</a> , 512)	0
dense_3 ( <a href="#">Dense</a> )	( <a href="#">None</a> , 100)	51,300
dropout_5 ( <a href="#">Dropout</a> )	( <a href="#">None</a> , 100)	0
dense_4 ( <a href="#">Dense</a> )	( <a href="#">None</a> , 50)	5,050
dropout_6 ( <a href="#">Dropout</a> )	( <a href="#">None</a> , 50)	0



dense\_5 (Dense)

(None, 10)

510

Total params: 14,771,548 (56.35 MB)

Trainable params: 56,860 (222.11 KB)

Non-trainable params: 14,714,688 (56.13 MB)

Epoch 1/200

2025-09-22 07:03:59.900674: E

external/local\_xla/xla/service/slow\_operation\_alarm.cc:65] Trying algorithm eng11{k2=1,k3=0} for conv (f32[128,64,224,224]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,64,224,224]{3,2,1,0}, f32[64,64,3,3]{3,2,1,0}, f32[64]{0}), window={size=3x3 pad=1\_1x1\_1}, dim\_labels=bf01\_oi01->bf01, custom\_call\_target="\_\_cudnn\$convBiasActivationForward", backend\_config={"cudnn\_c onv\_backend\_config":{"activation\_mode":"kRelu","conv\_result\_scale":1,"leakyrelu\_alpha":0,"side\_input\_scale":0},"force\_earliest\_schedule":false,"operation\_queue\_id":"0","wait\_on\_operation\_queues":[]}} is taking a while...

2025-09-22 07:04:00.167501: E

external/local\_xla/xla/service/slow\_operation\_alarm.cc:133] The operation took 1.267004147s

Trying algorithm eng11{k2=1,k3=0} for conv (f32[128,64,224,224]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,64,224,224]{3,2,1,0}, f32[64,64,3,3]{3,2,1,0}, f32[64]{0}), window={size=3x3 pad=1\_1x1\_1}, dim\_labels=bf01\_oi01->bf01, custom\_call\_target="\_\_cudnn\$convBiasActivationForward", backend\_config={"cudnn\_c onv\_backend\_config":{"activation\_mode":"kRelu","conv\_result\_scale":1,"leakyrelu\_alpha":0,"side\_input\_scale":0},"force\_earliest\_schedule":false,"operation\_queue\_id":"0","wait\_on\_operation\_queues":[]}} is taking a while...

2025-09-22 07:04:10.300850: E

external/local\_xla/xla/service/slow\_operation\_alarm.cc:65] Trying algorithm eng36{k2=3,k3=0} for conv (f32[128,128,112,112]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,128,112,112]{3,2,1,0}, f32[128,128,3,3]{3,2,1,0}, f32[128]{0}), window={size=3x3 pad=1\_1x1\_1}, dim\_labels=bf01\_oi01->bf01, custom\_call\_target="\_\_cudnn\$convBiasActivationForward", backend\_config={"cudnn\_c onv\_backend\_config":{"activation\_mode":"kRelu","conv\_result\_scale":1,"leakyrelu\_alpha":0,"side\_input\_scale":0},"force\_earliest\_schedule":false,"operation\_queue\_id":"0","wait\_on\_operation\_queues":[]}} is taking a while...

2025-09-22 07:04:10.988668: E

external/local\_xla/xla/service/slow\_operation\_alarm.cc:133] The operation took 1.687989021s

Trying algorithm eng36{k2=3,k3=0} for conv (f32[128,128,112,112]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,128,112,112]{3,2,1,0}, f32[128,128,3,3]{3,2,1,0}, f32[128]{0}), window={size=3x3 pad=1\_1x1\_1}, dim\_labels=bf01\_oi01->bf01, custom\_call\_target="\_\_cudnn\$convBiasActivationForward", backend\_config={"cudnn\_c

onv\_backend\_config":{"activation\_mode":"kRelu","conv\_result\_scale":1,"leakyrelu\_alpha":0,"side\_input\_scale":0},"force\_earliest\_schedule":false,"operation\_queue\_id":"0","wait\_on\_operation\_queues":[]} is taking a while...

391/391 359s 822ms/step -  
accuracy: 0.1941 - loss: 2.1650 - val\_accuracy: 0.4281 - val\_loss: 1.6345  
Epoch 2/200  
391/391 286s 731ms/step -  
accuracy: 0.3755 - loss: 1.6891 - val\_accuracy: 0.4951 - val\_loss: 1.4140  
Epoch 3/200  
391/391 284s 727ms/step -  
accuracy: 0.4337 - loss: 1.5387 - val\_accuracy: 0.5288 - val\_loss: 1.3201  
Epoch 4/200  
391/391 284s 726ms/step -  
accuracy: 0.4705 - loss: 1.4653 - val\_accuracy: 0.5440 - val\_loss: 1.2820  
Epoch 5/200  
391/391 286s 731ms/step -  
accuracy: 0.4826 - loss: 1.4295 - val\_accuracy: 0.5642 - val\_loss: 1.2317  
Epoch 6/200  
391/391 285s 728ms/step -  
accuracy: 0.5031 - loss: 1.3861 - val\_accuracy: 0.5725 - val\_loss: 1.2089  
Epoch 7/200  
391/391 284s 726ms/step -  
accuracy: 0.5128 - loss: 1.3609 - val\_accuracy: 0.5762 - val\_loss: 1.1942  
Epoch 8/200  
391/391 284s 727ms/step -  
accuracy: 0.5156 - loss: 1.3495 - val\_accuracy: 0.5866 - val\_loss: 1.1713  
Epoch 9/200  
391/391 285s 730ms/step -  
accuracy: 0.5256 - loss: 1.3329 - val\_accuracy: 0.5889 - val\_loss: 1.1565  
Epoch 10/200  
391/391 284s 727ms/step -  
accuracy: 0.5339 - loss: 1.3103 - val\_accuracy: 0.5987 - val\_loss: 1.1308  
Epoch 11/200  
391/391 284s 727ms/step -  
accuracy: 0.5375 - loss: 1.3088 - val\_accuracy: 0.6039 - val\_loss: 1.1226  
Epoch 12/200  
391/391 284s 727ms/step -  
accuracy: 0.5435 - loss: 1.2854 - val\_accuracy: 0.6087 - val\_loss: 1.1319  
Epoch 13/200  
391/391 284s 727ms/step -  
accuracy: 0.5457 - loss: 1.2880 - val\_accuracy: 0.6098 - val\_loss: 1.1193  
Epoch 14/200  
391/391 284s 727ms/step -  
accuracy: 0.5492 - loss: 1.2716 - val\_accuracy: 0.6082 - val\_loss: 1.1163  
Epoch 15/200  
391/391 284s 727ms/step -  
accuracy: 0.5489 - loss: 1.2674 - val\_accuracy: 0.6166 - val\_loss: 1.0942

Epoch 16/200  
391/391 285s 729ms/step -  
accuracy: 0.5596 - loss: 1.2488 - val\_accuracy: 0.6235 - val\_loss: 1.0898  
Epoch 17/200  
391/391 285s 729ms/step -  
accuracy: 0.5587 - loss: 1.2463 - val\_accuracy: 0.6172 - val\_loss: 1.1146  
Epoch 18/200  
391/391 285s 730ms/step -  
accuracy: 0.5645 - loss: 1.2449 - val\_accuracy: 0.6239 - val\_loss: 1.0686  
Epoch 19/200  
391/391 285s 729ms/step -  
accuracy: 0.5654 - loss: 1.2334 - val\_accuracy: 0.6188 - val\_loss: 1.0937  
Epoch 20/200  
391/391 285s 729ms/step -  
accuracy: 0.5656 - loss: 1.2420 - val\_accuracy: 0.6291 - val\_loss: 1.0641  
Epoch 21/200  
391/391 285s 729ms/step -  
accuracy: 0.5717 - loss: 1.2233 - val\_accuracy: 0.6211 - val\_loss: 1.0839  
Epoch 22/200  
391/391 287s 735ms/step -  
accuracy: 0.5653 - loss: 1.2345 - val\_accuracy: 0.6279 - val\_loss: 1.0642  
Epoch 23/200  
391/391 285s 730ms/step -  
accuracy: 0.5699 - loss: 1.2209 - val\_accuracy: 0.6235 - val\_loss: 1.0740  
Epoch 24/200  
391/391 284s 727ms/step -  
accuracy: 0.5702 - loss: 1.2209 - val\_accuracy: 0.6323 - val\_loss: 1.0532  
Epoch 25/200  
391/391 284s 727ms/step -  
accuracy: 0.5755 - loss: 1.2150 - val\_accuracy: 0.6375 - val\_loss: 1.0556  
Epoch 26/200  
391/391 284s 727ms/step -  
accuracy: 0.5780 - loss: 1.2083 - val\_accuracy: 0.6209 - val\_loss: 1.0729  
Epoch 27/200  
391/391 284s 727ms/step -  
accuracy: 0.5727 - loss: 1.2148 - val\_accuracy: 0.6356 - val\_loss: 1.0585  
Epoch 28/200  
391/391 284s 727ms/step -  
accuracy: 0.5758 - loss: 1.2060 - val\_accuracy: 0.6363 - val\_loss: 1.0438  
Epoch 29/200  
391/391 284s 727ms/step -  
accuracy: 0.5808 - loss: 1.1903 - val\_accuracy: 0.6377 - val\_loss: 1.0486  
Epoch 30/200  
391/391 285s 728ms/step -  
accuracy: 0.5827 - loss: 1.1939 - val\_accuracy: 0.6358 - val\_loss: 1.0545  
Epoch 31/200  
391/391 285s 729ms/step -  
accuracy: 0.5841 - loss: 1.1966 - val\_accuracy: 0.6379 - val\_loss: 1.0461

Epoch 32/200  
 391/391 285s 729ms/step -  
 accuracy: 0.5817 - loss: 1.2005 - val\_accuracy: 0.6398 - val\_loss: 1.0341  
 Epoch 33/200  
 391/391 285s 728ms/step -  
 accuracy: 0.5861 - loss: 1.1897 - val\_accuracy: 0.6439 - val\_loss: 1.0354  
 Epoch 34/200  
 391/391 285s 729ms/step -  
 accuracy: 0.5830 - loss: 1.1858 - val\_accuracy: 0.6321 - val\_loss: 1.0636  
 Epoch 35/200  
 391/391 285s 730ms/step -  
 accuracy: 0.5806 - loss: 1.2009 - val\_accuracy: 0.6395 - val\_loss: 1.0410  
 Epoch 36/200  
 391/391 287s 733ms/step -  
 accuracy: 0.5901 - loss: 1.1758 - val\_accuracy: 0.6483 - val\_loss: 1.0253  
 Epoch 37/200  
 391/391 285s 729ms/step -  
 accuracy: 0.5869 - loss: 1.1792 - val\_accuracy: 0.6478 - val\_loss: 1.0335  
 Epoch 38/200  
 391/391 284s 727ms/step -  
 accuracy: 0.5905 - loss: 1.1744 - val\_accuracy: 0.6396 - val\_loss: 1.0356  
 Epoch 39/200  
 391/391 284s 727ms/step -  
 accuracy: 0.5869 - loss: 1.1850 - val\_accuracy: 0.6454 - val\_loss: 1.0205  
 Epoch 40/200  
 391/391 285s 729ms/step -  
 accuracy: 0.5890 - loss: 1.1850 - val\_accuracy: 0.6466 - val\_loss: 1.0345  
 Epoch 41/200  
 391/391 285s 730ms/step -  
 accuracy: 0.5962 - loss: 1.1602 - val\_accuracy: 0.6416 - val\_loss: 1.0266  
 Epoch 42/200  
 391/391 285s 730ms/step -  
 accuracy: 0.5937 - loss: 1.1743 - val\_accuracy: 0.6493 - val\_loss: 1.0205  
 Epoch 43/200  
 391/391 285s 730ms/step -  
 accuracy: 0.5897 - loss: 1.1756 - val\_accuracy: 0.6413 - val\_loss: 1.0251  
 Epoch 44/200  
 391/391 284s 727ms/step -  
 accuracy: 0.5935 - loss: 1.1702 - val\_accuracy: 0.6468 - val\_loss: 1.0196  
 Epoch 45/200  
 391/391 284s 727ms/step -  
 accuracy: 0.5928 - loss: 1.1625 - val\_accuracy: 0.6456 - val\_loss: 1.0331  
 Epoch 46/200  
 391/391 285s 729ms/step -  
 accuracy: 0.5928 - loss: 1.1719 - val\_accuracy: 0.6543 - val\_loss: 1.0122  
 Epoch 47/200  
 391/391 286s 731ms/step -  
 accuracy: 0.5940 - loss: 1.1685 - val\_accuracy: 0.6461 - val\_loss: 1.0156

Epoch 48/200  
391/391 284s 727ms/step -  
accuracy: 0.5970 - loss: 1.1505 - val\_accuracy: 0.6520 - val\_loss: 1.0174  
Epoch 49/200  
391/391 284s 727ms/step -  
accuracy: 0.6017 - loss: 1.1496 - val\_accuracy: 0.6428 - val\_loss: 1.0254  
Epoch 50/200  
391/391 284s 727ms/step -  
accuracy: 0.5991 - loss: 1.1509 - val\_accuracy: 0.6548 - val\_loss: 0.9992  
Epoch 51/200  
391/391 285s 730ms/step -  
accuracy: 0.5973 - loss: 1.1534 - val\_accuracy: 0.6526 - val\_loss: 1.0083  
Epoch 52/200  
391/391 285s 729ms/step -  
accuracy: 0.6000 - loss: 1.1559 - val\_accuracy: 0.6431 - val\_loss: 1.0277  
Epoch 53/200  
391/391 285s 729ms/step -  
accuracy: 0.5993 - loss: 1.1635 - val\_accuracy: 0.6437 - val\_loss: 1.0237  
Epoch 54/200  
391/391 287s 733ms/step -  
accuracy: 0.6047 - loss: 1.1439 - val\_accuracy: 0.6591 - val\_loss: 0.9959  
Epoch 55/200  
391/391 285s 729ms/step -  
accuracy: 0.6036 - loss: 1.1447 - val\_accuracy: 0.6589 - val\_loss: 0.9961  
Epoch 56/200  
391/391 285s 729ms/step -  
accuracy: 0.6051 - loss: 1.1465 - val\_accuracy: 0.6444 - val\_loss: 1.0154  
Epoch 57/200  
391/391 285s 729ms/step -  
accuracy: 0.5945 - loss: 1.1630 - val\_accuracy: 0.6589 - val\_loss: 0.9936  
Epoch 58/200  
391/391 285s 730ms/step -  
accuracy: 0.5931 - loss: 1.1541 - val\_accuracy: 0.6588 - val\_loss: 1.0005  
Epoch 59/200  
391/391 286s 732ms/step -  
accuracy: 0.6005 - loss: 1.1477 - val\_accuracy: 0.6532 - val\_loss: 1.0074  
Epoch 60/200  
391/391 284s 727ms/step -  
accuracy: 0.6020 - loss: 1.1494 - val\_accuracy: 0.6514 - val\_loss: 1.0351  
Epoch 61/200  
391/391 284s 726ms/step -  
accuracy: 0.6037 - loss: 1.1454 - val\_accuracy: 0.6525 - val\_loss: 0.9974  
Epoch 62/200  
391/391 284s 727ms/step -  
accuracy: 0.6030 - loss: 1.1469 - val\_accuracy: 0.6548 - val\_loss: 1.0089  
Epoch 62: early stopping  
Restoring model weights from the end of the best epoch: 57.


```
[13]: loss, val_accuracy = model.evaluate(x_test,y_test)
```

```
313/313          76s 200ms/step -  
accuracy: 0.6574 - loss: 0.9963
```

```
[14]: print(f"Validation Accuracy of Cifar10 dataset with CNN model:␣  
      ↪{val_accuracy*100:.4f}%")
```

```
Validation Accuracy of Cifar10 dataset with CNN model: 65.8900%
```

```
[15]: plt.figure(figsize=(12,4))  
plt.subplot(1,2,1)  
plt.plot(history.history['accuracy'], label='Train Acc')  
plt.plot(history.history['val_accuracy'], label='Val Acc')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.subplot(1,2,2)  
plt.plot(history.history['loss'], label='Train Loss')  
plt.plot(history.history['val_loss'], label='Val Loss')  
plt.xlabel("Epochs")  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



**0.3 Question 3: Try to use the pre-trained model ResNet50 and design a new model in similar way and report the results obtained.**

**0.3.1 Importing the necessary libraries**

```
[16]: import tensorflow as tf  
from tensorflow import keras  
from keras import Sequential  
from keras.layers import Dense, Resizing, Dropout, GlobalAveragePooling2D
```

```

from keras.applications import ResNet50
from keras.callbacks import EarlyStopping
from keras.datasets import cifar10
import matplotlib.pyplot as plt

```

### 0.3.2 Loading and preprocessing the dataset

```

[17]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train / 255.0
x_test  = x_test / 255.0

```

```

[18]: base_model = ResNet50(
        include_top=False,
        weights='imagenet',
        input_shape=(224, 224, 3)
    )

    for layer in base_model.layers:
        layer.trainable = False
    ## Building the model
    model = Sequential()
    model.add(Resizing(224, 224, input_shape=(32, 32, 3)))
    model.add(base_model)
    model.add(GlobalAveragePooling2D())
    model.add(Dense(100, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(50, activation='relu'))
    model.add(Dropout(0.3))
    model.add(Dense(10, activation='softmax'))

    model.summary()

    model.compile(
        loss='sparse_categorical_crossentropy',
        optimizer='adam',
        metrics=['accuracy']
    )

    estop = EarlyStopping(
        monitor='val_loss',
        min_delta=1e-4,
        patience=5,
        verbose=1,
        restore_best_weights=True
    )

    history = model.fit(

```

```

    x_train, y_train,
    batch_size=128,
    epochs=200,
    verbose=1,
    validation_data=(x_test, y_test),
    callbacks=[estop]
)

loss, val_accuracy = model.evaluate(x_test, y_test)
print(f"Validation Accuracy of CIFAR-10 with ResNet50 model: {val_accuracy*100:.
↵4f}%")

plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel("Epochs")
plt.ylabel('Loss')
plt.legend()

plt.show()

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5)  
94765736/94765736 0s

0us/step

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
resizing_1 (Resizing)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, 7, 7, 2048)	23,587,712
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 2048)	0
dense_6 (Dense)	(None, 100)	204,900



dropout_7 (Dropout)	(None, 100)	0
dense_7 (Dense)	(None, 50)	5,050
dropout_8 (Dropout)	(None, 50)	0
dense_8 (Dense)	(None, 10)	510

Total params: 23,798,172 (90.78 MB)

Trainable params: 210,460 (822.11 KB)

Non-trainable params: 23,587,712 (89.98 MB)

Epoch 1/200

391/391 205s 460ms/step -  
accuracy: 0.0983 - loss: 2.3250 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 2/200

391/391 161s 411ms/step -  
accuracy: 0.0999 - loss: 2.3027 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 3/200

391/391 161s 411ms/step -  
accuracy: 0.0963 - loss: 2.3027 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 4/200

391/391 161s 412ms/step -  
accuracy: 0.0992 - loss: 2.3027 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 5/200

391/391 161s 412ms/step -  
accuracy: 0.0975 - loss: 2.3027 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 6/200

391/391 161s 412ms/step -  
accuracy: 0.0972 - loss: 2.3027 - val\_accuracy: 0.1000 - val\_loss: 2.3026

Epoch 6: early stopping

Restoring model weights from the end of the best epoch: 1.

313/313 32s 87ms/step -

accuracy: 0.0995 - loss: 2.3026

Validation Accuracy of CIFAR-10 with ResNet50 model: 10.0000%

